

notebook_modeling

April 6, 2023

1 Predicting Dengue Cases

- Student name: *Aysu Erdemir*
- Student pace: *Flex*
- Scheduled project review date/time: *March, 2023*
- Instructor name: *Abhineet Kulkarni*

1.1 Modeling:

The dengue data with labels (1990-2008) was split into training and test sets using the first 80% of the data as train, and the final 20% for test. Additional dataset with climate features only (without the knowledge of true case counts)(2008-2013) was used to forecast upcoming case counts for the best performing models.

Several versions of machine learning models were built, tuned and validated to be able to forecast the time series data:

- **Negative Binomial Regression** (multiple regression used for count data following the negative binomial). This method was chosen specifically because total_cases could be described by a negative binomial distribution with a population variance that is much larger than the population mean.
- **Sarimax** (Seasonal Autoregressive Integrated Moving Average Exogenous model)- a generalization of an autoregressive moving average (ARMA) model which supports time series data with a seasonal component.
- **XGBoost (Extreme Gradient Boosting) Regression** Gradient-boosted decision tree algorithm used for regression predictive modeling.
- **LSTM (long short-term memory network)** A variety of recurrent neural networks (RNNs) that are capable of learning long-term dependencies, especially in sequence prediction problems.

1.2 Evaluation:

- Performance is evaluated according to the **Mean Absolute Error**.

- MAE is a popular metric to use as the error value is easily interpreted. This is because the value is on the same scale as the target you are predicting for.

1.2.1 Import Necessary packages:

```
[1]: # Import basic packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import seaborn as sns
%matplotlib inline
```

```
[2]: # Import packages for Negative Binomial regression
import statsmodels.api as sm
import scipy.stats as stats
from statsmodels.formula.api import ols
import statsmodels.formula.api as smf
from statsmodels.tools import eval_measures
```

```
[3]: # Import packages for ARIMA
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.stattools import acf, pacf
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.statespace.sarimax import SARIMAX

# !pip install pmdarima
from pmdarima.arima import auto_arima
from pmdarima.arima import ADFTTest
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting pmdarima

Downloading pmdarima-2.0.3-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux_2_28_x86_64.whl (1.9 MB)

1.9/1.9 MB

18.6 MB/s eta 0:00:00

Requirement already satisfied: numpy>=1.21.2 in

/usr/local/lib/python3.9/dist-packages (from pmdarima) (1.22.4)

Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (1.4.4)

Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (67.6.1)

Requirement already satisfied: Cython!=0.29.18,!0.29.31,>=0.29 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (0.29.34)

Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (1.2.2)

Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (1.1.1)
 Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (1.10.1)
 Requirement already satisfied: statsmodels>=0.13.2 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (0.13.5)
 Requirement already satisfied: urllib3 in /usr/local/lib/python3.9/dist-packages (from pmdarima) (1.26.15)
 Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.9/dist-packages (from pandas>=0.19->pmdarima) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.9/dist-packages (from pandas>=0.19->pmdarima) (2022.7.1)
 Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.9/dist-packages (from scikit-learn>=0.22->pmdarima) (3.1.0)
 Requirement already satisfied: patsy>=0.5.2 in /usr/local/lib/python3.9/dist-packages (from statsmodels>=0.13.2->pmdarima) (0.5.3)
 Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.9/dist-packages (from statsmodels>=0.13.2->pmdarima) (23.0)
 Requirement already satisfied: six in /usr/local/lib/python3.9/dist-packages (from patsy>=0.5.2->statsmodels>=0.13.2->pmdarima) (1.16.0)
 Installing collected packages: pmdarima
 Successfully installed pmdarima-2.0.3

```
[4]: # Import packages for XGBoost
import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import TimeSeriesSplit
```

```
[5]: # Import packages for LSTM
# directly download from keras without importing tensorflow, otherwise there
# are issues with using TimeseriesGenerator and model fitting

from sklearn.preprocessing import MinMaxScaler
from keras.preprocessing.sequence import TimeseriesGenerator
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers import Dropout
from keras.callbacks import EarlyStopping
import random
```

```
[6]: # Import packages for model evaluation
from sklearn.metrics import mean_absolute_error as MAE
from sklearn.metrics import mean_squared_error as MSE
```

1.3 Get the data ready for modeling:

```
[7]: # Import train_final and test_final
from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving train_final.csv to train_final.csv

Saving test_final.csv to test_final.csv

```
[8]: # Read the Data
train_final = pd.read_csv("train_final.csv").iloc[:, 1:] # drop the first_
↳unnamed column of repeated index that was read.
test_final = pd.read_csv("test_final.csv").iloc[:, 1:]
```

```
[9]: # change `week_start_date` to datetime
train_final["week_start_date"] = pd.to_datetime(train_final["week_start_date"])
test_final["week_start_date"] = pd.to_datetime(test_final["week_start_date"])

# set the index to `week_start_date` to datetime
train_final = train_final.set_index("week_start_date")
test_final = test_final.set_index("week_start_date")
```

```
[10]: train_final.info()
```

<class 'pandas.core.frame.DataFrame'>

DatetimeIndex: 925 entries, 1990-07-16 to 2008-04-22

Data columns (total 41 columns):

#	Column	Non-Null Count	Dtype
0	total_cases	925 non-null	int64
1	year	925 non-null	float64
2	weekofyear	925 non-null	float64
3	month	925 non-null	int64
4	fall	925 non-null	int64
5	spring	925 non-null	int64
6	summer	925 non-null	int64
7	winter	925 non-null	int64
8	station_avg_temp_c	925 non-null	float64
9	station_max_temp_c	925 non-null	float64
10	station_min_temp_c	925 non-null	float64
11	reanalysis_tdtr_k	925 non-null	float64
12	reanalysis_specific_humidity_g_per_kg	925 non-null	float64
13	reanalysis_precip_amt_kg_per_m2	925 non-null	float64
14	ndvi_ne	925 non-null	float64
15	ndvi_nw	925 non-null	float64
16	ndvi_se	925 non-null	float64

```

17  ndvi_sw                925 non-null    float64
18  ndvi_average           925 non-null    float64
19  ndvi_average_cat       925 non-null    object
20  grassy                 925 non-null    int64
21  soily                  925 non-null    int64
22  watery                 925 non-null    int64
23  station_avg_temp_c_shift  925 non-null    float64
24  station_max_temp_c_shift  925 non-null    float64
25  station_min_temp_c_shift  925 non-null    float64
26  reanalysis_tdtr_k_shift  925 non-null    float64
27  reanalysis_specific_humidity_g_per_kg_shift  925 non-null    float64
28  reanalysis_precip_amt_kg_per_m2_shift  925 non-null    float64
29  grassy_shift           925 non-null    float64
30  soily_shift            925 non-null    float64
31  watery_shift           925 non-null    float64
32  station_max_temp_c_shift_18  925 non-null    float64
33  station_min_temp_c_shift_18  925 non-null    float64
34  station_avg_temp_c_shift_18  925 non-null    float64
35  reanalysis_tdtr_k_shift_8  925 non-null    float64
36  reanalysis_specific_humidity_g_per_kg_shift_12  925 non-null    float64
37  reanalysis_precip_amt_kg_per_m2_shift_8  925 non-null    float64
38  grassy_shift_20        925 non-null    float64
39  soily_shift_20         925 non-null    float64
40  watery_shift_20        925 non-null    float64
dtypes: float64(31), int64(9), object(1)
memory usage: 303.5+ KB

```

```
[11]: test_final.info()
```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 260 entries, 2008-04-29 to 2013-04-23
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   year                                260 non-null    float64
1   weekofyear                          260 non-null    float64
2   month                              260 non-null    int64
3   fall                                260 non-null    int64
4   spring                              260 non-null    int64
5   summer                              260 non-null    int64
6   winter                              260 non-null    int64
7   station_avg_temp_c                  260 non-null    float64
8   station_max_temp_c                  260 non-null    float64
9   station_min_temp_c                  260 non-null    float64
10  reanalysis_tdtr_k                    260 non-null    float64
11  reanalysis_specific_humidity_g_per_kg  260 non-null    float64
12  reanalysis_precip_amt_kg_per_m2       260 non-null    float64
13  ndvi_ne                              260 non-null    float64

```

14	ndvi_nw	260 non-null	float64
15	ndvi_se	260 non-null	float64
16	ndvi_sw	260 non-null	float64
17	ndvi_average	260 non-null	float64
18	ndvi_average_cat	260 non-null	object
19	grassy	260 non-null	int64
20	soily	260 non-null	int64
21	watery	260 non-null	int64
22	station_avg_temp_c_shift	260 non-null	float64
23	station_max_temp_c_shift	260 non-null	float64
24	station_min_temp_c_shift	260 non-null	float64
25	reanalysis_precip_amt_kg_per_m2_shift	260 non-null	float64
26	reanalysis_specific_humidity_g_per_kg_shift	260 non-null	float64
27	reanalysis_tdtr_k_shift	260 non-null	float64
28	grassy_shift	260 non-null	float64
29	soily_shift	260 non-null	float64
30	watery_shift	260 non-null	float64
31	station_max_temp_c_shift_18	260 non-null	float64
32	station_min_temp_c_shift_18	260 non-null	float64
33	station_avg_temp_c_shift_18	260 non-null	float64
34	reanalysis_tdtr_k_shift_8	260 non-null	float64
35	reanalysis_specific_humidity_g_per_kg_shift_12	260 non-null	float64
36	reanalysis_precip_amt_kg_per_m2_shift_8	260 non-null	float64
37	grassy_shift_20	260 non-null	float64
38	soily_shift_20	260 non-null	float64
39	watery_shift_20	260 non-null	float64

dtypes: float64(31), int64(8), object(1)
memory usage: 83.3+ KB

Both data sets look in line with one another.

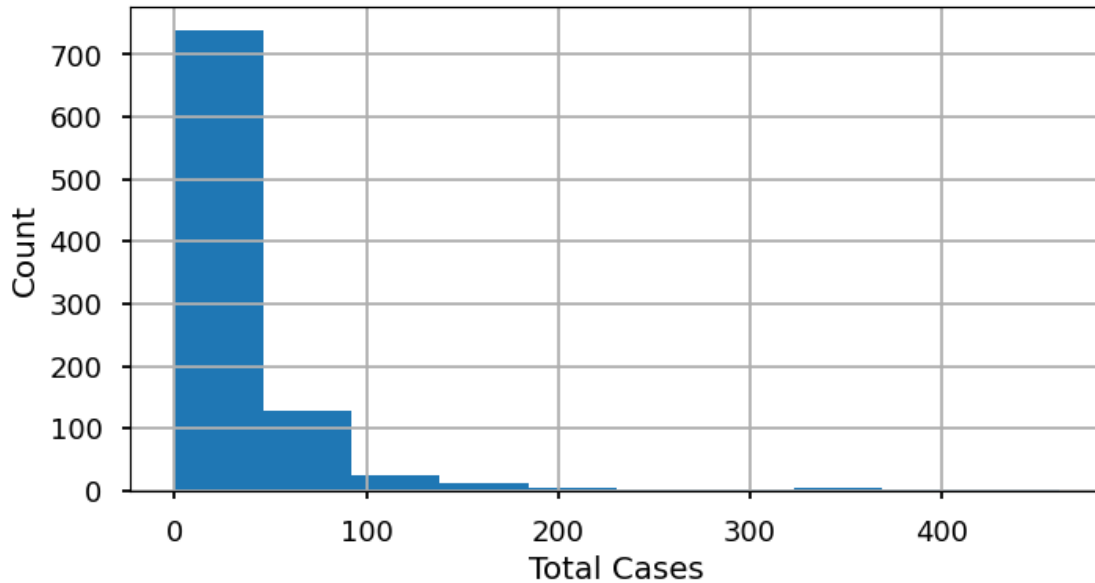
2 Negative Binomial Regression:

-
- Our target variable, total_cases is a non-negative integer, which means we're looking to make some count predictions. Standard regression techniques for this type of prediction would include Poisson regression and Negative binomial regression.
 - Poisson regression fits according to the assumption that the mean and variance of the population distribution are equal. When they aren't, specifically when the variance is much larger than the mean, the negative binomial approach is a better approach.

```
[12]: # Check the distribution of the target variable to see if it follows negative_
      ↪ binomial distribution.
with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=(8,4))
    train_final.total_cases.hist()
    ax.set_xlabel('Total Cases')
```

```
ax.set_ylabel('Count')
print('mean: ', train_final.total_cases.mean())
print('var :', train_final.total_cases.var())
```

```
mean: 34.52540540540541
var : 2661.2885854685856
```



```
[13]: # create a copy of the train_final for NBR
train_NBR = train_final.copy()
train_NBR.head()
```

```
[13]:
```

	total_cases	year	weekofyear	month	fall	spring	summer	\
week_start_date								
1990-07-16	2	1990.0	29.0	7	0	0	1	
1990-07-23	6	1990.0	30.0	7	0	0	1	
1990-07-30	17	1990.0	31.0	7	0	0	1	
1990-08-06	23	1990.0	32.0	8	0	0	1	
1990-08-13	13	1990.0	33.0	8	0	0	1	

	winter	station_avg_temp_c	station_max_temp_c	...	\
week_start_date					
1990-07-16	0	28.128571	32.8	...	
1990-07-23	0	28.114286	31.7	...	
1990-07-30	0	28.242857	34.4	...	
1990-08-06	0	28.200000	33.3	...	
1990-08-13	0	28.042857	32.8	...	

	watery_shift	station_max_temp_c_shift_18 \
week_start_date		
1990-07-16	0.0	32.990000
1990-07-23	0.0	32.872727
1990-07-30	0.0	32.866667
1990-08-06	0.0	32.776923
1990-08-13	0.0	32.892857

	station_min_temp_c_shift_18	station_avg_temp_c_shift_18 \
week_start_date		
1990-07-16	22.940000	27.584286
1990-07-23	22.827273	27.581818
1990-07-30	22.916667	27.627381
1990-08-06	22.907692	27.664835
1990-08-13	22.900000	27.706122

	reanalysis_tdtr_k_shift_8 \
week_start_date	
1990-07-16	2.169643
1990-07-23	2.151786
1990-07-30	2.150000
1990-08-06	2.032143
1990-08-13	2.092857

	reanalysis_specific_humidity_g_per_kg_shift_12 \
week_start_date	
1990-07-16	16.840286
1990-07-23	16.892857
1990-07-30	16.963214
1990-08-06	17.240595
1990-08-13	17.425714

	reanalysis_precip_amt_kg_per_m2_shift_8	grassy_shift_20 \
week_start_date		
1990-07-16	26.65000	0.300000
1990-07-23	28.85250	0.272727
1990-07-30	32.22750	0.333333
1990-08-06	36.01875	0.384615
1990-08-13	37.03250	0.428571

	soily_shift_20	watery_shift_20
week_start_date		
1990-07-16	0.700000	0.0
1990-07-23	0.727273	0.0
1990-07-30	0.666667	0.0
1990-08-06	0.615385	0.0
1990-08-13	0.571429	0.0

[5 rows x 41 columns]

2.0.1 Train test split:


- Splitting data into train/test sets is to evaluate under- and overfitting and help to choose the hyperparameters.

```
[14]: # Set the first 80% of the data to train, and remaining 20% to test:
train = train_NBR.head(750)
test = train_NBR.tail(train_NBR.shape[0] - 750)
```

```
[15]: # See the distribution of case counts per each successive year for presentation:
with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=(10,5))
    ax.plot(train_NBR['total_cases'])
    ax.set_xlabel('Date')
    ax.set_ylabel('Total Weekly Dengue Cases')
    fig.patch.set_alpha(0) # make the figure background transparent
    fig.savefig('total_cases_years.png', dpi=300)
    files.download("total_cases_years.png")
```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



```
[16]: # Plot test-train
with plt.style.context('seaborn-talk'):
```


```

fig, ax = plt.subplots(figsize=(10,5))
ax.plot(train['total_cases'], label='Train')
ax.plot(test['total_cases'], label='Test', color = 'limegreen')
ax.set_title('Train-Test Split')
ax.set_xlabel('Date')
ax.set_ylabel('Total Weekly Dengue Cases')
fig.patch.set_alpha(0) # make the figure background transparent
plt.legend();
fig.savefig('total_cases_test_train_split.png', dpi=300)
files.download("total_cases_test_train_split.png")

```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



2.1 Negative Binomial Regression Model #1

- A simple model with 4 original variables.

Create a function to:

- to get the best NBR model using the most optimum alpha that minimizes MAE.

```

[17]: def get_best_NBR_model(train, test, model_formula):

    # Create the aplha grid
    grid = np.linspace(0.000001, 2, 1000) # grid = np.arange(-8, -3, dtype=np.
    ↪float64)

```

```

best_alpha = []
best_MAE_score = 1000

# Find the best hyper parameter, alpha - specify the regularization
↪ distribution between L1 and L2.
for alpha in grid:
    model = smf.glm(formula=model_formula,
                    data=train,
                    family=sm.families.NegativeBinomial(alpha=alpha))

    results = model.fit()

    predictions_test = results.predict(test).astype(int)
    score_test = eval_measures.meanabs(predictions_test, test.total_cases)

    if score_test < best_MAE_score:
        best_alpha = alpha
        best_MAE_score = score_test

print('best alpha = ', best_alpha)
print('(best) test MAE score = ', best_MAE_score)

# refit on train dataset using best alpha
model = smf.glm(formula=model_formula,
                data=train,
                family=sm.families.NegativeBinomial(alpha=best_alpha))
fitted_model = model.fit()

predictions_train = fitted_model.predict(train).astype(int)
score_train = eval_measures.meanabs(predictions_train, train.total_cases)
print('train MAE score = ', score_train)

return fitted_model

```

```

[18]: # Create the model formula and run above function
model_formula = "total_cases ~ 1 + " \
                "station_avg_temp_c + " \
                "reanalysis_tdtr_k + " \
                "reanalysis_specific_humidity_g_per_kg + " \
                "reanalysis_precip_amt_kg_per_m2 " \

best_model = get_best_NBR_model(train, test, model_formula)
best_model.summary()

```

```

best alpha = 1e-06
(best) test MAE score = 22.617142857142856

```

```
train MAE score = 29.001333333333335
```

```
[18]: <class 'statsmodels.iolib.summary.Summary'>
      """
              Generalized Linear Model Regression Results
      =====
Dep. Variable:          total_cases      No. Observations:          750
Model:                  GLM             Df Residuals:              745
Model Family:          NegativeBinomial  Df Model:                  4
Link Function:          Log             Scale:                    1.0000
Method:                 IRLS            Log-Likelihood:            -17211.
Date:                  Thu, 06 Apr 2023  Deviance:                 30782.
Time:                  00:32:53          Pearson chi2:              5.22e+04
No. Iterations:        6                 Pseudo R-squ. (CS):       0.9669
Covariance Type:       nonrobust
      =====
      =====
                                coef      std err          z
P>|z|      [0.025      0.975]
-----
-----
Intercept                                -0.5271      0.146      -3.598
0.000      -0.814      -0.240
station_avg_temp_c                      0.1479      0.010      15.537
0.000       0.129       0.167
reanalysis_tdtr_k                       -0.3015      0.015     -20.405
0.000      -0.330      -0.273
reanalysis_specific_humidity_g_per_kg    0.0472      0.009       5.277
0.000       0.030       0.065
reanalysis_precip_amt_kg_per_m2          0.0014      0.000       9.855
0.000       0.001       0.002
      =====
      =====
      """
```

- p values for all variables are below 0.05 and statistically significant, meaning we can reject the null hypothesis that these variables do not correlate with total cases.

Create a function to:

- (1) plot the true total cases against test and train predictions.
- (2) display the Mean Absolute Error and Root Mean Square Error for both train and test. This helps to see the overall performance of the model on both train and test, and assess overfitting.

```
[19]: def forecast_graph(observed, predictions_train, predictions_test):

      with plt.style.context('seaborn-talk'):
```

```

fig, ax = plt.subplots(figsize=(14,6))
ax.plot(observed, label='observed')
ax.plot(predictions_train, label='train - fitted')
ax.plot(predictions_test, label='test - forecasted')
ax.set_title("Dengue Predicted Cases vs. Actual Cases")
ax.set_xlabel('Date')
ax.set_ylabel('Total Cases')
plt.legend();

```

```

[20]: def final_scores(y_train_true, y_train_pred, y_test_true, y_test_pred):
    MAE_train = MAE(y_train_true, y_train_pred)
    MAE_test = MAE(y_test_true, y_test_pred)

    RMSE_train = MSE(y_train_true, y_train_pred, squared=False) # Setting
    ↪ squared to False will return the RMSE.
    RMSE_test = MSE(y_test_true, y_test_pred, squared=False)


    print('MAE_train: %f' % MAE_train)
    print('MAE_test: %f' % MAE_test)
    print('-----')
    print('RMSE_train: %f' % RMSE_train)
    print('RMSE_test: %f' % RMSE_test)

```

```

[21]: forecast_graph(train_NBR.total_cases, best_model.predict(train), best_model.
    ↪ predict(test))

```



```

[22]: final_scores(train.total_cases, best_model.predict(train), test.total_cases,
    ↪ best_model.predict(test))

```

MAE_train: 29.173050

MAE_test: 22.943668

RMSE_train: 53.788695

RMSE_test: 30.244678

Summary:

- The model captures the basic seasonality, while missing all the individual outbreaks - peaks.
- Next let's see what happens when we add more variables.

2.2 Negative Binomial Regression Model #2

- Add more variables

```
[23]: model_formula = "total_cases ~ 1 + " \
        "station_avg_temp_c + " \
        "station_max_temp_c + " \
        "station_min_temp_c + " \
        "reanalysis_tdtr_k + " \
        "reanalysis_specific_humidity_g_per_kg + " \
        "reanalysis_precip_amt_kg_per_m2 + " \
        "month "

best_model = get_best_NBR_model(train, test, model_formula)
best_model.summary()
```

best alpha = 1e-06

(best) test MAE score = 21.017142857142858

train MAE score = 28.038666666666668

```
[23]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```


              Generalized Linear Model Regression Results
=====
Dep. Variable:          total_cases    No. Observations:          750
Model:                  GLM           Df Residuals:              742
Model Family:          NegativeBinomial  Df Model:                  7
Link Function:          Log            Scale:                   1.0000
Method:                 IRLS           Log-Likelihood:          -15364.
Date:                  Thu, 06 Apr 2023  Deviance:                27086.
Time:                  00:33:22         Pearson chi2:            4.10e+04
No. Iterations:         5               Pseudo R-squ. (CS):      0.9998
Covariance Type:        nonrobust
=====
=====
                                coef    std err          z
P>|z|    [0.025    0.975]
-----
```

```

-----
Intercept                                -0.5088      0.152      -3.337
0.001      -0.808      -0.210
station_avg_temp_c                       -0.0797      0.017      -4.633
0.000      -0.113      -0.046
station_max_temp_c                        0.2414      0.008      29.861
0.000      0.226      0.257
station_min_temp_c                       -0.0109      0.011      -1.020
0.308      -0.032      0.010
reanalysis_tdtr_k                         -0.3484      0.015     -22.924
0.000      -0.378      -0.319
reanalysis_specific_humidity_g_per_kg    -0.0725      0.010      -7.462
0.000      -0.092      -0.053
reanalysis_precip_amt_kg_per_m2          0.0008      0.000      4.951
0.000      0.000      0.001
month                                    0.1202      0.002     55.314
0.000      0.116      0.124
=====
=====
"""

```

```
[24]: forecast_graph(train_NBR.total_cases, best_model.predict(train), best_model.
      ↪predict(test))
```



```
[25]: final_scores(train.total_cases, best_model.predict(train), test.total_cases,
      ↪best_model.predict(test))
```

```
MAE_train: 28.226007
MAE_test: 21.302535
```

```
-----  
RMSE_train: 51.058028  
RMSE_test: 30.007795
```

Summary:

- The model with more variables is a little better as the MAE scores are slightly lower.
- It still captures the basic seasonality only, while missing all the individual outbreaks - peaks.
- There is also some asynchrony between the onset of each predicted peak.

2.3 Negative Binomial Regression Model #3

- Use time shifted variables

```
[26]: model_formula = "total_cases ~ 1 + " \  
      "station_max_temp_c_shift + " \  
      "station_min_temp_c_shift + " \  
      "station_avg_temp_c_shift + " \  
      "reanalysis_tdtr_k_shift + " \  
      "reanalysis_specific_humidity_g_per_kg_shift + " \  
      "reanalysis_precip_amt_kg_per_m2_shift + " \  
      "month + " \  
      "grassy_shift + " \  
      "watery_shift + " \  
      "soily_shift"  
  
best_model = get_best_NBR_model(train, test, model_formula)  
best_model.summary()
```

```
best alpha = 0.0020030010010010013  
(best) test MAE score = 20.405714285714286  
train MAE score = 27.653333333333332
```

```
[26]: <class 'statsmodels.iolib.summary.Summary'>  
      """
```

```
              Generalized Linear Model Regression Results  
=====
```

Dep. Variable:	total_cases	No. Observations:	750
Model:	GLM	Df Residuals:	740
Model Family:	NegativeBinomial	Df Model:	9
Link Function:	Log	Scale:	1.0000
Method:	IRLS	Log-Likelihood:	-13505.
Date:	Thu, 06 Apr 2023	Deviance:	23319.
Time:	00:33:56	Pearson chi2:	3.58e+04
No. Iterations:	6	Pseudo R-squ. (CS):	0.9998
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	z
--	------	---------	---


```

P>|z|      [0.025      0.975]
-----
-----
Intercept                    -0.6370      0.124      -5.126
0.000      -0.881      -0.393
station_max_temp_c_shift      0.1841      0.009      21.584
0.000      0.167      0.201
station_min_temp_c_shift      0.0655      0.011      5.815
0.000      0.043      0.088
station_avg_temp_c_shift      -0.0418      0.018      -2.274
0.023      -0.078      -0.006
reanalysis_tdtr_k_shift      -0.4561      0.016      -28.371
0.000      -0.488      -0.425
reanalysis_specific_humidity_g_per_kg_shift      -0.0909      0.010      -8.661
0.000      -0.111      -0.070
reanalysis_precip_amt_kg_per_m2_shift      0.0007      0.000      4.164
0.000      0.000      0.001
month                        0.0992      0.003      38.799
0.000      0.094      0.104
grassy_shift                 -0.3490      0.044      -7.863
0.000      -0.436      -0.262
watery_shift                 -0.2550      0.043      -5.985
0.000      -0.338      -0.171
soily_shift                  -0.0330      0.044      -0.752
0.452      -0.119      0.053
=====
=====
"""


```

- p values for all variables except soily_shift are below 0.05 and statistically significant, meaning we can reject the null hypothesis that these variables do not correlate with total cases.

```

[27]: forecast_graph(train_NBR.total_cases, best_model.predict(train), best_model.
      ↪ predict(test))

```



```
[28]: final_scores(train.total_cases, best_model.predict(train),
                  test.total_cases, best_model.predict(test))
```

MAE_train: 27.839518

MAE_test: 20.656355

RMSE_train: 50.823396

RMSE_test: 29.647056

Summary:

- This model is a little better as the MAE scores are slightly lower.
- However, it still captures the basic seasonality only, while missing all the individual outbreaks - peaks.
- Using time shifted variables did not help much.

2.4 Negative Binomial Regression Model #4

- Use time sifted variables with most highly correlated lagged means to account for sustained heat, humidity, precipitation and vegetation.

```
[29]: model_formula = "total_cases ~ 1 + " \
                      "station_max_temp_c_shift_18 + " \
                      "station_min_temp_c_shift_18 + " \
                      "station_avg_temp_c_shift_18 + " \
                      "reanalysis_tdtr_k_shift_8 + " \
                      "reanalysis_specific_humidity_g_per_kg_shift_12 + " \
                      "reanalysis_precip_amt_kg_per_m2_shift_8 + " \
                      "fall + " \
                      "spring + " \
                      "winter + " \
```

```

"summer +" \
"grassy_shift_20 +" \
"soily_shift_20 +" \
"watery_shift_20" \

best_model = get_best_NBR_model(train, test, model_formula)
best_model.summary()

```

```

best alpha = 1e-06
(best) test MAE score = 15.771428571428572
train MAE score = 20.158666666666665

```

```

[29]: <class 'statsmodels.iolib.summary.Summary'>
      """

```

```

                          Generalized Linear Model Regression Results
=====
Dep. Variable:            total_cases      No. Observations:            750
Model:                  GLM                Df Residuals:                738
Model Family:          NegativeBinomial    Df Model:                   11
Link Function:          Log                Scale:                     1.0000
Method:                 IRLS               Log-Likelihood:             -8939.3
Date:                  Thu, 06 Apr 2023    Deviance:                   14238.
Time:                  00:35:04            Pearson chi2:               1.84e+04
No. Iterations:        5                  Pseudo R-squ. (CS):        1.000
Covariance Type:       nonrobust
=====
=====

```


			coef	std err	z
P> z	[0.025	0.975]			
Intercept			-9.9703	0.246	-40.590
0.000	-10.452	-9.489			
station_max_temp_c_shift_18			1.4173	0.031	45.622
0.000	1.356	1.478			
station_min_temp_c_shift_18			2.2127	0.038	58.632
0.000	2.139	2.287			
station_avg_temp_c_shift_18			-2.4829	0.062	-40.011
0.000	-2.605	-2.361			
reanalysis_tdtr_k_shift_8			-0.5996	0.028	-21.653
0.000	-0.654	-0.545			
reanalysis_specific_humidity_g_per_kg_shift_12			-0.4644	0.016	-29.577
0.000	-0.495	-0.434			
reanalysis_precip_amt_kg_per_m2_shift_8			-0.0037	0.000	-8.484
0.000	-0.005	-0.003			
fall			-2.3035	0.074	-30.927
0.000	-2.449	-2.158			

spring			-2.5182	0.052	-48.305
0.000	-2.620	-2.416			
winter			-2.6845	0.067	-40.275
0.000	-2.815	-2.554			
summer			-2.4641	0.061	-40.504
0.000	-2.583	-2.345			
grassy_shift_20			-3.4037	0.102	-33.440
0.000	-3.603	-3.204			
soily_shift_20			-2.4105	0.106	-22.803
0.000	-2.618	-2.203			
watery_shift_20			-4.1560	0.087	-47.657
0.000	-4.327	-3.985			

```
=====
=====
"""
```

- Using the lagged variables p values for all 13 variables are below 0.05 and statistically significant.

```
[30]: forecast_graph(train_NBR.total_cases, best_model.predict(train), best_model.
      ↪ predict(test))
```



```
[31]: final_scores(train.total_cases, best_model.predict(train), test.total_cases,
      ↪ best_model.predict(test))
```

MAE_train: 20.243531

MAE_test: 15.921087

RMSE_train: 36.521933

RMSE_test: 25.074627

Summary:


- This model is significantly better as the MAE and RMSE scores are lower.
- This model captures some of the individual peaks-outbreak correctly and it generalizes to the test set better as well.

2.5 Refit on the whole dataset for feature importance and forecasting into future :

- Once we have done enough iterations and we are satisfied with the performance, we can retrain your model on the total labeled data.
- Splitting data into train/test sets is to evaluate under- and overfitting and help us choose the hyperparameters. Once this is achieved, it makes sense to get maximal performance before using your model in real applications.
- Let's see how the model predicts on the final test set, for which we do not have the true case counts available.

```
[32]: # refit on the whole data set to be able to project onto future and extract
      ↪ feature importances
model = smf.glm(formula=model_formula,
                 data=train_NBR,
                 family=sm.families.NegativeBinomial(alpha=1e-06))
fitted_model = model.fit()
```


```
[33]: forecast_graph(train_NBR.total_cases, fitted_model.predict(train_NBR),
      ↪ fitted_model.predict(test_final))
```



By refitting the model on the whole train set and projecting on to future, we see that there are two more moderate size outbreaks predicted by the end of year 2009 and 2012.

```
[34]: # Create a smoother version of the above graph:
total_cases_rolled = train_NBR.total_cases.rolling(window=6, min_periods = 1).
    ↪mean()
predicted_train_final_rolled = fitted_model.predict(train_NBR).
    ↪rolling(window=6, min_periods = 1).mean()
predicted_test_final_rolled = fitted_model.predict(test_final).
    ↪rolling(window=6, min_periods = 1).mean()

forecast_graph(total_cases_rolled, predicted_train_final_rolled,
    ↪predicted_test_final_rolled)
```



2.5.1 Extract Feature importance from the final model:

```
[35]: # Drop the model coefficient for graphing purposes:
coeff = fitted_model.params.drop('Intercept')


# Sort the coefficients:
coeff = coeff.abs().sort_values(ascending=False) # if you want to keep - or +:
    ↪coeff = coeff.iloc[(coeff.abs()*-1.0).argsort()]
# extract the index representing variable names
feature_names = pd.DataFrame(coeff).index

# Plot the coefficients using a TORNADO PLOT:
with plt.style.context('seaborn-talk'):
    base_color = sns.color_palette("husl", 9)[6]
    fig, ax = plt.subplots(figsize=(10, 5))
    sns.barplot(x=coeff.values, y=coeff.index, color = base_color, ax=ax,
    ↪orient='h')
```

```

ax.set_title("Relative Importance of Features \n for Predicting Dengue_
Cases \n", fontsize=16)
ax.set_xlabel("Feature importance", fontsize=14)
ax.set_ylabel("Feature", fontsize=14)
ax.set_yticklabels(labels=feature_names)
fig.tight_layout();
# fig.savefig('./images/TornadoPlot_Coefs.png', dpi=300)

```



- According to Negative Binomial Regression the most important features for predicting dengue cases are the NDVI variables followed by the seasons.

3 SARIMA

```
[36]: train_ARIMA = train_final.copy()
```

3.0.1 RE-SAMPLE the time series dataset from weekly to monthly:

- This step was necessary since hyperparameter search with `pm.auto_arima` did not work efficiently on weekly data.

```

[37]: train_ARIMA= train_ARIMA.resample('1M').mean()
train_ARIMA
# We are down to 214 rows.

```

```

[37]:
total_cases  year  weekofyear  month  fall  spring  summer \
week_start_date
1990-07-31      8.333333  1990.0      30.0    7.0    0.0    0.0    1.0

```

1990-08-31	21.250000	1990.0	33.5	8.0	0.0	0.0	1.0
1990-09-30	27.750000	1990.0	37.5	9.0	1.0	0.0	0.0
1990-10-31	42.400000	1990.0	42.0	10.0	1.0	0.0	0.0
1990-11-30	56.000000	1990.0	46.5	11.0	1.0	0.0	0.0
...
2007-12-31	14.750000	2007.0	50.5	12.0	0.0	0.0	0.0
2008-01-31	12.600000	2008.0	3.0	1.0	0.0	0.0	0.0
2008-02-29	5.250000	2008.0	7.5	2.0	0.0	0.0	0.0
2008-03-31	2.500000	2008.0	11.5	3.0	0.0	1.0	0.0
2008-04-30	3.000000	2008.0	15.5	4.0	0.0	1.0	0.0

	winter	station_avg_temp_c	station_max_temp_c	...	\
week_start_date				...	
1990-07-31	0.0	28.161905	32.966667	...	
1990-08-31	0.0	28.310714	32.900000	...	
1990-09-30	0.0	28.400000	33.075000	...	
1990-10-31	0.0	27.954286	32.980000	...	
1990-11-30	0.0	27.253571	32.225000	...	
...	
2007-12-31	1.0	25.442857	28.875000	...	
2008-01-31	1.0	24.780000	28.440000	...	
2008-02-29	1.0	24.664286	27.900000	...	
2008-03-31	0.0	25.171429	29.575000	...	
2008-04-30	0.0	25.900000	30.275000	...	

	watery_shift	station_max_temp_c_shift_18	\
week_start_date			
1990-07-31	0.00	32.909798	
1990-08-31	0.00	32.875570	
1990-09-30	0.00	33.029657	
1990-10-31	0.00	33.142222	
1990-11-30	0.00	32.968056	
...	
2007-12-31	0.75	32.277778	
2008-01-31	0.20	31.272222	
2008-02-29	0.25	30.031944	
2008-03-31	0.75	29.048611	
2008-04-30	0.75	28.756944	

	station_min_temp_c_shift_18	station_avg_temp_c_shift_18	\
week_start_date			
1990-07-31	22.894646	27.597828	
1990-08-31	22.913277	27.717010	
1990-09-30	23.124183	27.936006	
1990-10-31	23.448889	28.242063	
1990-11-30	23.409722	28.174603	
...	

2007-12-31	23.591667	27.821032
2008-01-31	23.035556	27.024603
2008-02-29	22.241667	26.101389
2008-03-31	21.706944	25.381151
2008-04-30	21.397222	25.061706

reanalysis_tdtr_k_shift_8 \

week_start_date	
1990-07-31	2.157143
1990-08-31	2.104464
1990-09-30	2.325893
1990-10-31	2.439286
1990-11-30	2.263393
...	...
2007-12-31	2.646429
2008-01-31	2.438214
2008-02-29	2.435268
2008-03-31	2.337054
2008-04-30	2.536161

reanalysis_specific_humidity_g_per_kg_shift_12 \

week_start_date	
1990-07-31	16.898786
1990-08-31	17.463720
1990-09-30	17.863095
1990-10-31	17.999024
1990-11-30	18.081935
...	...
2007-12-31	17.372321
2008-01-31	16.465048
2008-02-29	15.370804
2008-03-31	14.724107
2008-04-30	14.225060


reanalysis_precip_amt_kg_per_m2_shift_8 grassy_shift_20 \

week_start_date		
1990-07-31	29.243333	0.302020
1990-08-31	37.863438	0.397047
1990-09-30	41.957500	0.457473
1990-10-31	46.427500	0.590000
1990-11-30	78.978750	0.575000
...
2007-12-31	43.128750	0.050000
2008-01-31	39.525500	0.030000
2008-02-29	25.243125	0.000000
2008-03-31	13.421563	0.000000
2008-04-30	8.242500	0.000000

	soily_shift_20	watery_shift_20
week_start_date		
1990-07-31	0.697980	0.0000
1990-08-31	0.602953	0.0000
1990-09-30	0.542527	0.0000
1990-10-31	0.410000	0.0000
1990-11-30	0.425000	0.0000
...
2007-12-31	0.737500	0.2125
2008-01-31	0.690000	0.2800
2008-02-29	0.650000	0.3500
2008-03-31	0.575000	0.4250
2008-04-30	0.500000	0.5000

[214 rows x 40 columns]

```
[38]: train_ARIMA.total_cases.plot(figsize=(8,4));
      # we preserved the same shape
```



3.0.2 Plot the ACF (auto correlation function) and PACF (partial auto correlation function):

- Both the ACF and PACF start with a lag of 0, which is the correlation of the time series with itself and results in a correlation of 1.
- The partial autocorrelation function can be interpreted as a regression of the series against its past lags. It helps you come up with a possible order for the autoregressive term.

- To figure out the order of an AR model, we need to look at the PACF.
- To figure out the order of an MA model, we need to look at the ACF.


```
[39]: trainseasonal = train_ARIMA['total_cases']
```

```
[40]: fig, ax = plt.subplots(figsize=(16,3))
      plot_acf(trainseasonal, ax=ax, lags=40);

      fig, ax = plt.subplots(figsize=(16,3))
      plot_pacf(trainseasonal, ax=ax, lags=40);
```

/usr/local/lib/python3.9/dist-packages/statsmodels/graphics/tsaplots.py:348:
FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.


```
warnings.warn(
```



- There are several autocorrelations that are significantly non-zero. Therefore, the time series is non-random.
- High degree of autocorrelation between adjacent (lag = 1) and near-adjacent (lag = 2) observations in both ACF and PACF plots

```
[41]: # Plot seasonality and trends:
      # Since I'm working with monthly data, setting the frequency to 12 seems
      ↪ logical.
      result = seasonal_decompose(trainseasonal, period = 12, model='additive')
      fig = result.plot()
      fig.set_size_inches((10, 8))
```

```
fig.tight_layout()
plt.show()
```



- We can see the seasonality clearly, but there does not seem to be a strong trend in the data.
- Let's check for stationarity using a statistical test:

```
[42]: # Statistical test to see if the time series is stationary or not
adf_test = ADFTest(alpha = .05)
adf_test.should_diff(trainseasonal)
```

```
[42]: (0.01, False)
```


- Data is nonstationary! We need to use INtegrated (I) concept.

4 SARIMA #1 Baseline Model:

- Using only total cases as the predictor

```
[43]: # Test-train split
train = train_ARIMA['total_cases'].head(170)
test = train_ARIMA['total_cases'].tail(train_ARIMA.shape[0] - 170)

fig, ax = plt.subplots(figsize=(8,4))
ax.plot(train, label='train')
ax.plot(test, label='valid')
ax.set_title('Train-Validation Split')
plt.legend();
```



4.0.1 Parameter Search:

A seasonal ARIMA model is classified as an $ARIMA(p,d,q) \times (P,D,Q)$ model,

- **p** = number of autoregressive (AR) terms
- **d** = number of differences
- **q** = number of moving average (MA) terms
- **P** = number of seasonal autoregressive (SAR) terms
- **D** = number of seasonal differences
- **Q** = number of seasonal moving average (SMA) terms

Using the `auto_arima()` function from the `pmdarima` package, we can perform a parameter search for the optimal values of the model.

```
[44]: # Parameter Search
sarimax_best = auto_arima(train,
                           start_p=1, start_q=1, max_p=5, max_q=5,
                           d=1, D=1,
                           start_P=1, start_Q=1, max_P=5, max_Q=5,
                           m=12,
                           max_order=None,
                           error_action='ignore',
                           suppress_warnings=True,
                           trace=True,
                           stepwise=True)
```

Performing stepwise search to minimize aic

```
ARIMA(1,1,1)(1,1,1)[12]      : AIC=inf, Time=0.60 sec
ARIMA(0,1,0)(0,1,0)[12]      : AIC=1631.374, Time=0.03 sec
ARIMA(1,1,0)(1,1,0)[12]      : AIC=1565.034, Time=0.11 sec
ARIMA(0,1,1)(0,1,1)[12]      : AIC=inf, Time=0.47 sec
ARIMA(1,1,0)(0,1,0)[12]      : AIC=1608.520, Time=0.05 sec
ARIMA(1,1,0)(2,1,0)[12]      : AIC=1544.574, Time=0.30 sec
ARIMA(1,1,0)(3,1,0)[12]      : AIC=1545.150, Time=0.76 sec
ARIMA(1,1,0)(2,1,1)[12]      : AIC=inf, Time=0.93 sec
ARIMA(1,1,0)(1,1,1)[12]      : AIC=inf, Time=0.32 sec
ARIMA(1,1,0)(3,1,1)[12]      : AIC=inf, Time=4.11 sec
ARIMA(0,1,0)(2,1,0)[12]      : AIC=1562.275, Time=0.25 sec
ARIMA(2,1,0)(2,1,0)[12]      : AIC=1529.092, Time=0.49 sec
ARIMA(2,1,0)(1,1,0)[12]      : AIC=1547.536, Time=0.18 sec
ARIMA(2,1,0)(3,1,0)[12]      : AIC=1527.694, Time=1.25 sec
ARIMA(2,1,0)(4,1,0)[12]      : AIC=1520.050, Time=2.38 sec
ARIMA(2,1,0)(5,1,0)[12]      : AIC=1518.757, Time=2.84 sec
ARIMA(2,1,0)(5,1,1)[12]      : AIC=inf, Time=7.08 sec
ARIMA(2,1,0)(4,1,1)[12]      : AIC=inf, Time=4.77 sec
ARIMA(1,1,0)(5,1,0)[12]      : AIC=1532.452, Time=2.41 sec
ARIMA(3,1,0)(5,1,0)[12]      : AIC=1517.635, Time=2.99 sec
ARIMA(3,1,0)(4,1,0)[12]      : AIC=1518.524, Time=2.15 sec
ARIMA(3,1,0)(5,1,1)[12]      : AIC=inf, Time=9.71 sec
ARIMA(3,1,0)(4,1,1)[12]      : AIC=inf, Time=5.27 sec
ARIMA(4,1,0)(5,1,0)[12]      : AIC=1518.671, Time=3.95 sec
ARIMA(3,1,1)(5,1,0)[12]      : AIC=inf, Time=14.06 sec
ARIMA(2,1,1)(5,1,0)[12]      : AIC=inf, Time=23.59 sec
ARIMA(4,1,1)(5,1,0)[12]      : AIC=inf, Time=16.35 sec
ARIMA(3,1,0)(5,1,0)[12] intercept : AIC=1519.594, Time=5.73 sec
```

Best model: ARIMA(3,1,0)(5,1,0)[12]

Total fit time: 113.173 seconds

```
[45]: # Creating and fitting Final SARIMAX model
Final_model = SARIMAX(train.astype('int'),
```

```

order=sarimax_best.order,
seasonal_order=sarimax_best.seasonal_order,
enforce_invertibility=False,
enforce_stationarity=False)
Final_output = Final_model.fit()

```

```

[46]: # Displaying the model summary and diagnostics
display(Final_output.summary());
Final_output.plot_diagnostics(figsize=(15, 10));

```

```

<class 'statsmodels.iolib.summary.Summary'>
"""

```

SARIMAX Results

```

=====
Dep. Variable:                total_cases    No. Observations:          170
    ↳ 170
Model:                SARIMAX(3, 1, 0)x(5, 1, 0, 12)    Log Likelihood          -435.814
    ↳ -435.814
Date:                Thu, 06 Apr 2023    AIC          889.627
    ↳ 889.627
Time:                00:37:06    BIC          912.517
    ↳ 912.517
Sample:                07-31-1990    HQIC          898.873
    ↳ 898.873

```

- 08-31-2004

Covariance Type: opg

```

=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.1099      0.095      1.152      0.249      -0.077      0.297
ar.L2         -0.2342      0.116     -2.022      0.043      -0.461     -0.007
ar.L3         -0.1556      0.154     -1.013      0.311      -0.457      0.146
ar.S.L12       -0.6717      0.105     -6.413      0.000      -0.877     -0.466
ar.S.L24       -0.4578      0.139     -3.292      0.001      -0.730     -0.185
ar.S.L36       -0.3404      0.150     -2.266      0.023      -0.635     -0.046
ar.S.L48       -0.2625      0.176     -1.492      0.136      -0.607      0.082
ar.S.L60       -0.1015      0.141     -0.719      0.472      -0.378      0.175
sigma2        623.1333     58.891     10.581      0.000     507.710     738.557
=====

```

```

=====
Ljung-Box (L1) (Q):                0.05    Jarque-Bera (JB):                1246.
    ↳ 41
Prob(Q):                0.83    Prob(JB):                0.
    ↳ 00
Heteroskedasticity (H):            0.05    Skew:                -1.
    ↳ 88
Prob(H) (two-sided):            0.00    Kurtosis:                20.
    ↳ 44

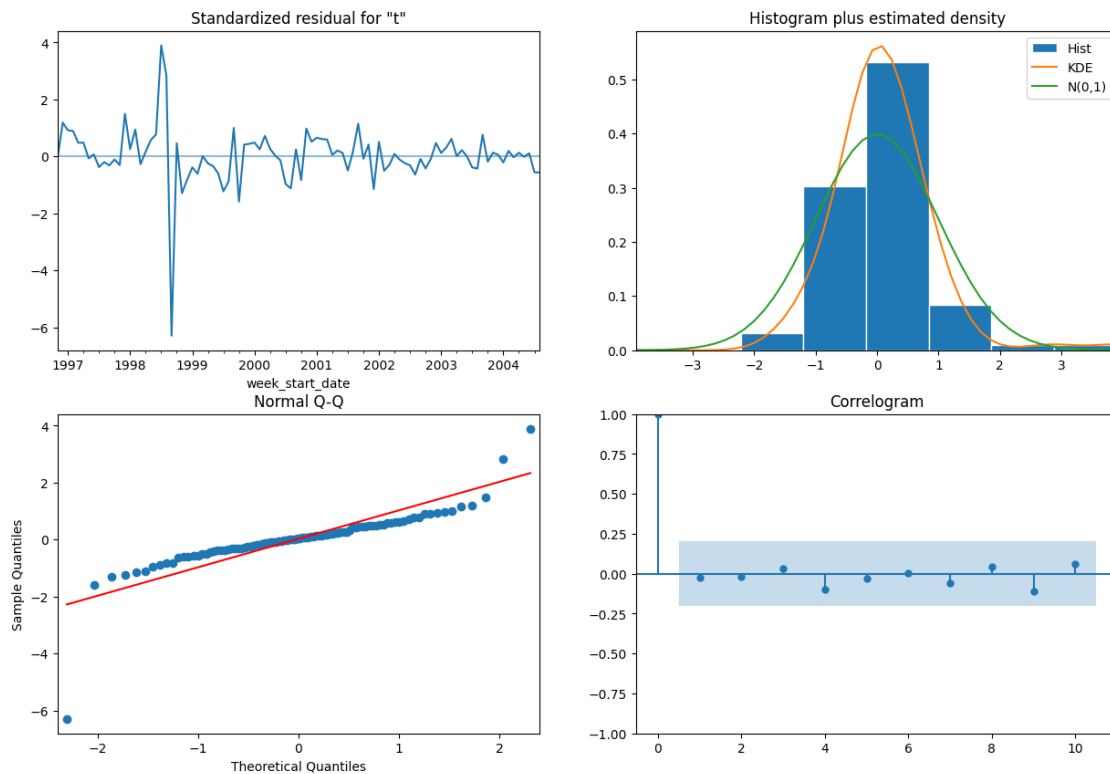
```

=====

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients
↳(complex-step).
```


"""



There does not seem to be any violation of assumptions for the model.

```
[47]: # Create predictions:
train_prediction = Final_output.predict(typ='levels') # sari_mod.
↳predict(start=train.index[0], end=train.index[-1])
test_prediction = Final_output.predict(start=test.index[0], end=test.
↳index[-1], typ='levels')
```

```
[48]: forecast_graph(train_ARIMA.total_cases, train_prediction, test_prediction)
```

```
[49]: final_scores(train, train_prediction, test, test_prediction)
```

```
MAE_train: 17.485334
```

```
MAE_test: 42.321119
```

```
-----
```

```
RMSE_train: 29.872551
```

```
RMSE_test: 51.515279
```

Summary:

- Basic Sarima overfits the data. It fits to the train data almost perfectly while performing very poorly for the test set.

5 SARIMAX #2 Full Multivariate Model:

- Using exogenous variables


```
[50]: train_ARIMA2 = train_ARIMA[['total_cases', 'month',
    ↪ 'station_max_temp_c_shift_18',
    ↪ 'station_min_temp_c_shift_18',
    ↪ 'station_avg_temp_c_shift_18',
    ↪ 'reanalysis_tdtr_k_shift_8',
    ↪ 'reanalysis_specific_humidity_g_per_kg_shift_12',
    ↪ 'reanalysis_precip_amt_kg_per_m2_shift_8',
    ↪ 'grassy_shift_20',
    ↪ 'soily_shift_20', 'watery_shift_20']]
```

```
[51]: train_ARIMA2 = train_ARIMA[['total_cases', 'month', 'station_max_temp_c',
    ↪ 'station_min_temp_c', 'station_avg_temp_c',
    ↪ 'reanalysis_tdtr_k',
```

```
'reanalysis_specific_humidity_g_per_kg',
'reanalysis_precip_amt_kg_per_m2', 'grassy',
'soily', 'watery']]
```

```
[52]: train = train_ARIMA2.head(170)
test = train_ARIMA2.tail(train_ARIMA2.shape[0] - 170)

fig, ax = plt.subplots(figsize=(8,4))
ax.plot(train.total_cases, label='train')
ax.plot(test.total_cases, label='test')
ax.set_title('Train-Test Split')
plt.legend();
```



```
[53]: X_train, y_train, = train.drop('total_cases', axis=1), train['total_cases']
X_test, y_test = test.drop('total_cases', axis=1), test['total_cases']
```

```
[54]: exog_varblds = ['month', 'station_max_temp_c',
                    'station_min_temp_c', 'station_avg_temp_c',
                    'reanalysis_tdtr_k',
                    'reanalysis_specific_humidity_g_per_kg',
                    'reanalysis_precip_amt_kg_per_m2', 'grassy',
                    'soily', 'watery']
```

```
[55]: endog_train = y_train
endog_test = y_test
```

```
[56]: list(range(0,10))
```

```
[56]: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

5.0.1 Reshape X_train and X_test into an array of exogenous regressors, shaped nobs x k:

```
[57]: exog_train = np.empty([len(X_train), len(exog_varblds)])  
      for i, var in zip(list(range(0,len(X_train))), exog_varblds):  
          exog_train[:,i] = np.array(X_train[var])
```

```
[58]: exog_test = np.empty([len(X_test), len(exog_varblds)])  
      for i, var in zip(list(range(0,len(X_test))), exog_varblds):  
          exog_test[:,i] = np.array(X_test[var])
```

```
[59]: exog_train.shape, exog_test.shape
```

```
[59]: ((170, 10), (44, 10))
```

```
[60]: # Parameter Search  
      sarimax_best = auto_arima(y = endog_train, # target  
                               X = exog_train, # external predictors  
                               start_p=1, start_q=1, max_p=5, max_q=5,  
                               d=1, D=1,  
                               start_P=1, start_Q=1, max_P=5, max_Q=5,  
                               m=12,  
                               max_order=None,  
                               error_action='ignore',  
                               suppress_warnings=True,  
                               trace=True,  
                               stepwise=True)
```

Performing stepwise search to minimize aic

ARIMA(1,1,1)(1,1,1)[12]	: AIC=inf, Time=4.19 sec
ARIMA(0,1,0)(0,1,0)[12]	: AIC=1648.917, Time=0.27 sec
ARIMA(1,1,0)(1,1,0)[12]	: AIC=1583.431, Time=2.09 sec
ARIMA(0,1,1)(0,1,1)[12]	: AIC=inf, Time=2.46 sec
ARIMA(1,1,0)(0,1,0)[12]	: AIC=1623.953, Time=0.65 sec
ARIMA(1,1,0)(2,1,0)[12]	: AIC=1562.388, Time=5.74 sec
ARIMA(1,1,0)(3,1,0)[12]	: AIC=1562.936, Time=9.27 sec
ARIMA(1,1,0)(2,1,1)[12]	: AIC=inf, Time=5.98 sec
ARIMA(1,1,0)(1,1,1)[12]	: AIC=inf, Time=2.57 sec
ARIMA(1,1,0)(3,1,1)[12]	: AIC=inf, Time=9.69 sec
ARIMA(0,1,0)(2,1,0)[12]	: AIC=1580.987, Time=4.16 sec
ARIMA(2,1,0)(2,1,0)[12]	: AIC=1542.558, Time=6.47 sec
ARIMA(2,1,0)(1,1,0)[12]	: AIC=1564.400, Time=2.52 sec
ARIMA(2,1,0)(3,1,0)[12]	: AIC=1540.180, Time=10.67 sec
ARIMA(2,1,0)(4,1,0)[12]	: AIC=1533.786, Time=16.80 sec
ARIMA(2,1,0)(5,1,0)[12]	: AIC=1534.384, Time=25.11 sec

```

ARIMA(2,1,0)(4,1,1)[12]      : AIC=inf, Time=28.83 sec
ARIMA(2,1,0)(3,1,1)[12]      : AIC=inf, Time=11.77 sec
ARIMA(2,1,0)(5,1,1)[12]      : AIC=inf, Time=26.39 sec
ARIMA(1,1,0)(4,1,0)[12]      : AIC=1552.500, Time=15.06 sec
ARIMA(3,1,0)(4,1,0)[12]      : AIC=1533.401, Time=19.13 sec
ARIMA(3,1,0)(3,1,0)[12]      : AIC=1539.780, Time=9.83 sec
ARIMA(3,1,0)(5,1,0)[12]      : AIC=1533.994, Time=25.37 sec
ARIMA(3,1,0)(4,1,1)[12]      : AIC=inf, Time=19.58 sec
ARIMA(3,1,0)(3,1,1)[12]      : AIC=inf, Time=11.78 sec
ARIMA(3,1,0)(5,1,1)[12]      : AIC=inf, Time=27.30 sec
ARIMA(4,1,0)(4,1,0)[12]      : AIC=1534.811, Time=18.66 sec
ARIMA(3,1,1)(4,1,0)[12]      : AIC=inf, Time=17.52 sec
ARIMA(2,1,1)(4,1,0)[12]      : AIC=inf, Time=19.34 sec
ARIMA(4,1,1)(4,1,0)[12]      : AIC=inf, Time=17.63 sec
ARIMA(3,1,0)(4,1,0)[12] intercept : AIC=1535.375, Time=18.83 sec

```

Best model: ARIMA(3,1,0)(4,1,0)[12]

Total fit time: 395.718 seconds

```

[61]: # Creating Final SARIMAX model
Final_model = SARIMAX(endog = endog_train,
                      exog = exog_train,
                      order=sarimax_best.order,
                      seasonal_order=sarimax_best.seasonal_order,
                      enforce_invertibility=False,
                      enforce_stationarity=False)
Final_output = Final_model.fit()

```

```

/usr/local/lib/python3.9/dist-packages/statsmodels/base/model.py:604:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to ")

```

```

[62]: # Displaying the model summary and diagnostics
display(Final_output.summary());
Final_output.plot_diagnostics(figsize=(15, 10));

```

```

<class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                SARIMAX Results
=====
Dep. Variable:                    total_cases    No. Observations:      170
    ↳ 170
Model:                SARIMAX(3, 1, 0)x(4, 1, 0, 12)    Log Likelihood      -488.362
    ↳ -488.362
Date:                Thu, 06 Apr 2023    AIC      1012.724
    ↳ 1012.724

```

Time: 00:43:54 BIC ┐
↪ 1060.666
Sample: 07-31-1990 HQIC ┐
↪ 1032.155


- 08-31-2004

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
x1	-7.875e-08	nan	nan	nan	nan	nan
x2	-2.0536	9.742	-0.211	0.833	-21.147	17.040
x3	-3.7804	12.941	-0.292	0.770	-29.144	21.584
x4	6.5039	19.119	0.340	0.734	-30.969	43.977
x5	11.4327	12.614	0.906	0.365	-13.290	36.156
x6	1.0381	11.378	0.091	0.927	-21.262	23.338
x7	0.0634	0.249	0.255	0.799	-0.424	0.551
x8	0.3277	10.211	0.032	0.974	-19.685	20.340
x9	0.7087	7.224	0.098	0.922	-13.451	14.868
x10	-1.0364	10.535	-0.098	0.922	-21.685	19.612
ar.L1	0.1991	0.092	2.168	0.030	0.019	0.379
ar.L2	-0.2849	0.100	-2.837	0.005	-0.482	-0.088
ar.L3	-0.1801	0.150	-1.204	0.229	-0.473	0.113
ar.S.L12	-0.7319	0.102	-7.199	0.000	-0.931	-0.533
ar.S.L24	-0.5634	0.126	-4.460	0.000	-0.811	-0.316
ar.S.L36	-0.3521	0.123	-2.870	0.004	-0.593	-0.112
ar.S.L48	-0.1959	0.087	-2.244	0.025	-0.367	-0.025
sigma2	587.8307	70.413	8.348	0.000	449.824	725.838
=====						
Ljung-Box (L1) (Q):			0.06	Jarque-Bera (JB):		1614.
↪41						
Prob(Q):			0.81	Prob(JB):		0.
↪00						
Heteroskedasticity (H):			0.08	Skew:		-2.
↪19						
Prob(H) (two-sided):			0.00	Kurtosis:		21.
↪61						
=====						

Warnings:


- [1] Covariance matrix calculated using the outer product of gradients┐
↪(complex-step).
 - [2] Covariance matrix is singular or near-singular, with condition number 8.
↪05e+27. Standard errors may be unstable.
- """



There does not seem to be any violation of assumptions for the model.

```
[63]: # Create Predictions:
train_prediction = Final_output.predict(typ='levels', exog=exog_train) #
    ↪ sari_mod.predict(start=train.index[0], end=train.index[-1])
test_prediction = Final_output.predict(start=test.index[0], end=test.
    ↪ index[-1], typ='levels', exog=exog_test)

[64]: forecast_graph(train_ARIMA.total_cases, train_prediction, test_prediction)
```



```
[65]: final_scores(train.total_cases, train_prediction, test.total_cases,
↳ test_prediction)
```

MAE_train: 17.613395

MAE_test: 19.492662

RMSE_train: 29.499109

RMSE_test: 32.196027

Summary:

- Multivariate Sarima performs better than the basic sarima. However it still does not capture the two peaks in the test dataset.
- While it performs very well for the train set, it does not generalize to unseen data.

6 XGB Regression

7 Model #1:

- using original variables


```
[66]: # Re-assign a new dataset called train_XGB
train_XGB = train_final.copy()
```

```
[67]: # Create a subset of the dataset to use in modeling:
train_XGB1 = train_XGB[['total_cases',
                        'month',
                        'station_avg_temp_c', 'station_min_temp_c',
```

```
'station_max_temp_c', 'reanalysis_tdtr_k',
'reanalysis_specific_humidity_g_per_kg',
'reanalysis_precip_amt_kg_per_m2',
'grassy', 'soily', 'watery']]
```

```
[68]: # Test-train split the dataset:
train = train_XGB1.head(750)
test = train_XGB1.tail(train_XGB1.shape[0] - 750)

fig, ax = plt.subplots(figsize=(8,4))
ax.plot(train.total_cases, label='train')
ax.plot(test.total_cases, label='test')
ax.set_title('Train-Validation Split')
plt.legend();
```



```
[69]: # Create train and test X and y
X_train, y_train, = train.drop('total_cases', axis=1), train['total_cases']
X_test, y_test = test.drop('total_cases', axis=1), test['total_cases']
```

```
[70]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
[70]: ((750, 10), (750,), (175, 10), (175,))
```

```
[71]: # time split the dataset into 5 folds to be used in cross validation
time_split = TimeSeriesSplit(n_splits=5)
[(el[0].shape, el[1].shape) for el in time_split.split(X_train)]
```



```
[71]: [((125,), (125,)),
      ((250,), (125,)),
      ((375,), (125,)),
      ((500,), (125,)),
      ((625,), (125,))]
```

```
[72]: # initiate the regressor
model = xgb.XGBRegressor()

# Specify the tunable parameters
parameters = {'objective':['reg:squarederror','reg:absoluteerror','reg:
↳squaredlogerror'], # 'count:poisson'
              'learning_rate':[0.01, 0.05, 0.1, 0.2], # default = 0.3, Lower
↳ratios avoid over-fitting.
              'max_depth': [2, 4, 6, 8], # default = 6, Lower
↳values avoid over-fitting.
              'min_child_weight': [1, 2, 3, 4], # default = 1, Larger
↳values avoid over-fitting.
              'gamma':[0.5, 1], # default = 0, Larger
↳values avoid over-fitting.
              'colsample_bytree':[0.5, 0.75], # default = 1, Lower
↳ratios avoid over-fitting.
              'subsample':[0.5, 0.75, 1]} # default = 1, Lower
↳ratios avoid over-fitting.

# Configure the GridSearchCV object to choose the best hyperparameters
# Using the neg_mean_squared_error metric to compare the results of 5-fold
↳cross-validation

xgb_grid = GridSearchCV(estimator = model,
                        cv = time_split,
                        param_grid = parameters,
                        scoring = 'neg_mean_squared_error',
↳#eval_metric='logloss'
                        verbose=0)


# Train the best model
xgb_grid.fit(X_train, y_train)

# Print best parameters and best score
print("Best parameters:", xgb_grid.best_params_)
print("Best Score (MAE): ", (xgb_grid.best_score_))
```

```
Best parameters: {'colsample_bytree': 0.75, 'gamma': 1, 'learning_rate': 0.2,
'max_depth': 4, 'min_child_weight': 1, 'objective': 'reg:absoluteerror',
'subsample': 0.75}
Best Score (MAE): -3116.6840740212792
```

```
[73]: # Create the train and test predictions as a DataFrame with index to use on the graph.
      # The best_estimator_field contains the best model trained by GridSearch.
      predicted_train = pd.DataFrame(xgb_grid.best_estimator_.predict(X_train),
      index= X_train.index)
      predicted_test = pd.DataFrame(xgb_grid.best_estimator_.predict(X_test), index=
      X_test.index)
```

```
[74]: forecast_graph(train_XGB.total_cases, predicted_train, predicted_test)
```



```
[75]: # Print the scores for both train and test
      final_scores(train.total_cases, xgb_grid.best_estimator_.predict(X_train),
      test.total_cases, xgb_grid.best_estimator_.predict(X_test))
```

MAE_train: 22.316101

MAE_test: 18.996592

RMSE_train: 54.488551

RMSE_test: 30.858615

Summary:

- The model only captures seasonality, and does not capture and of the individual peaks-outbreaks.
- Let's try another model with the lagged variables which gave the best results with NBR.

7.1 Model #2 using the lagged variables:

```
[76]: train_XGB2 = train_XGB[['total_cases', 'month',
                             'station_avg_temp_c_shift_18',
                             ↪ 'station_min_temp_c_shift_18',
                             'station_max_temp_c_shift_18',
                             ↪ 'reanalysis_tdtr_k_shift_8',
                             'reanalysis_specific_humidity_g_per_kg_shift_12',
                             'reanalysis_precip_amt_kg_per_m2_shift_8',
                             'grassy_shift_20',
                             'soily_shift_20',
                             'watery_shift_20']]

[77]: train = train_XGB2.head(750)
test = train_XGB2.tail(train_XGB2.shape[0] - 750)

[78]: X_train, y_train, = train.drop('total_cases', axis=1), train['total_cases']
X_test, y_test = test.drop('total_cases', axis=1), test['total_cases']

[79]: time_split = TimeSeriesSplit(n_splits=5)
[(el[0].shape, el[1].shape) for el in time_split.split(X_train)]

[79]: [((125,), (125,)),
      ((250,), (125,)),
      ((375,), (125,)),
      ((500,), (125,)),
      ((625,), (125,))]

[80]: model = xgb.XGBRegressor()

# Specify the tunable parameters
parameters = {'objective': ['reg:squarederror', 'reg:absoluteerror', 'reg:
↪ squaredlogerror'], # 'count:poisson'
              'learning_rate': [0.01, 0.05, 0.1, 0.2], # default = 0.3, Lower
↪ ratios avoid over-fitting.
              'max_depth': [2, 4, 6, 8], # default = 6, Lower
↪ values avoid over-fitting.
              'min_child_weight': [1, 2, 3, 4], # default = 1, Larger
↪ values avoid over-fitting.
              'gamma': [0.5, 1], # default = 0, Larger
↪ values avoid over-fitting.
              'colsample_bytree': [0.5, 0.75], # default = 1, Lower
↪ ratios avoid over-fitting.
              'subsample': [0.5, 0.75, 1]} # default = 1, Lower
↪ ratios avoid over-fitting.

# Configure the GridSearchCV object to choose the best hyperparameters
```

```

# using the neg_mean_squared_error metric to compare the results of 5-fold
↳cross-validation

xgb_grid = GridSearchCV(estimator = model,
                        cv = time_split,
                        param_grid = parameters,
                        scoring = 'neg_mean_squared_error',
↳#eval_metric='logloss'
                        verbose=0)

# Train the best model
xgb_grid.fit(X_train, y_train)

# Print best parameters and best score
print("Best parameters:", xgb_grid.best_params_)
print("Best Score (MAE): ", (xgb_grid.best_score_))

```

Best parameters: {'colsample_bytree': 0.75, 'gamma': 0.5, 'learning_rate': 0.05, 'max_depth': 2, 'min_child_weight': 3, 'objective': 'reg:squarederror', 'subsample': 0.5}

Best Score (MAE): -2608.85466896299

```


[81]: # Create the train and test predictions as a DataFrame with index to use on the
↳graph.
predicted_train = pd.DataFrame(xgb_grid.best_estimator_.predict(X_train),
↳columns= ['pred'], index= X_train.index)
predicted_test = pd.DataFrame(xgb_grid.best_estimator_.predict(X_test), columns=
↳['pred'], index= X_test.index)

```

```

[82]: forecast_graph(train_XGB.total_cases, predicted_train, predicted_test)

```



```
[83]: # Print the scores for both train and test
final_scores(train.total_cases, xgb_grid.best_estimator_.predict(X_train),
              test.total_cases, xgb_grid.best_estimator_.predict(X_test))
```

MAE_train: 14.660013

MAE_test: 17.724193

RMSE_train: 24.794945

RMSE_test: 26.477244

7.2 Best performing model so far:


Summary: * This model is by far the best model so far. It captures not only seasonality but also the majority of the individual peaks. * It performs equally well for train and test, with weekly case count errors of 14.6 and 17.7, respectively, and generalizes well to unseen data by detecting the two peaks present between years 2005-2008.

```
[84]: # re-create the graph with smooth lines
total_cases_rolled = train_XGB.total_cases.rolling(window=4, min_periods = 1).
    ↪mean()
predicted_train_rolled = predicted_train['pred'].rolling(window=4, min_periods=
    ↪1).mean()
predicted_test_rolled = predicted_test['pred'].rolling(window=4, min_periods =
    ↪1).mean()
```

```
[85]: with plt.style.context('seaborn-talk'):
fig, ax = plt.subplots(figsize=(12,6))
ax.plot(total_cases_rolled, label='observed')
ax.plot(predicted_train_rolled, label='train - fitted')
ax.plot(predicted_test_rolled, label='test - forecasted')
ax.set_title("Dengue Predicted Cases vs. Actual Cases")
ax.set_xlabel('Date')
ax.set_ylabel('Total Cases')
ax.xaxis.set_major_locator(mdates.YearLocator(2)) # Set the years on X axis
    ↪apart by 2
plt.legend()
fig.patch.set_alpha(0) # make the figure background transparent
# plt.tight_layout();
fig.savefig('XGB_Predict.png', dpi=300, bbox_inches='tight')
files.download("XGB_Predict.png")
```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



7.3 Refit on the whole dataset:

- Once we have done enough iterations and we are satisfied with the performance, we can retrain your model on the total labeled data to:
 - get maximal performance for forecasting into the future.
 - extract feature importances.

```
[86]: # re-create X,y train:
X_train_whole, y_train_whole = train_XGB2.drop('total_cases', axis=1),
↳train_XGB2['total_cases']
```

```
[87]: # get the same features out for test_final
X_test_final = test_final[['month',
                            'station_avg_temp_c_shift_18',
↳'station_min_temp_c_shift_18',
                            'station_max_temp_c_shift_18',
↳'reanalysis_tdtr_k_shift_8',
                            'reanalysis_specific_humidity_g_per_kg_shift_12',
                            'reanalysis_precip_amt_kg_per_m2_shift_8',
                            'grassy_shift_20',
                            'soily_shift_20',
                            'watery_shift_20']]
```

```
[88]: # Refit to the whole dataset using the best parameters:
model = xgb.XGBRegressor(objective='reg:squarederror',
                          learning_rate=0.05,
                          max_depth=2,
```

```

min_child_weight= 3,
subsample= 0.5,
colsample_bytree= 0.75,
gamma= 0.5)

```

#Creating the model on Training Data

```
final_model = model.fit(X_train_whole,y_train_whole)
```


[89]: *# Create predictions:*

```

predicted_train_final = pd.DataFrame(final_model.predict(X_train_whole),
    ↪ columns= ['pred'], index= X_train_whole.index)
predicted_test_final = pd.DataFrame(final_model.predict(X_test_final), columns=
    ↪ ['pred'], index= test_final.index)

```

[90]: `forecast_graph(train_XGB.total_cases, predicted_train_final,`
`↪ predicted_test_final)`



[91]: *# re-create the graph with smooth lines*

```

total_cases_rolled = train_XGB.total_cases.rolling(window=4, min_periods = 1).
    ↪ mean()
predicted_train_final_rolled = predicted_train_final['pred'].rolling(window=4,
    ↪ min_periods = 1).mean()
predicted_test_final_rolled = predicted_test_final['pred'].rolling(window=4,
    ↪ min_periods = 1).mean()

```

[92]: `with plt.style.context('seaborn-talk'):`
`fig, ax = plt.subplots(figsize=(14,6))`
`ax.plot(total_cases_rolled, label='observed')`
`ax.plot(predicted_train_final_rolled, label='train - fitted')`


```

ax.plot(predicted_test_final_rolled, label='test - forecasted')
ax.set_title("Dengue Predicted Cases vs. Actual Cases")
ax.set_xlabel('Date')
ax.set_ylabel('Total Cases')
ax.xaxis.set_major_locator(mdates.YearLocator(2)) # Set the years on X axis
↳ apart by 2
plt.legend()
fig.patch.set_alpha(0) # make the figure background transparent
# plt.tight_layout();
fig.savefig('XGB_Forecast.png', dpi=300, bbox_inches='tight')
files.download("XGB_Forecast.png")

```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



[93]: # Zoom on to the final predictions to find the corresponding the years for the
↳ peaks.


```

with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=(14,6))
    ax.plot(predicted_test_final_rolled, label='test - forecasted')
    ax.set_xlabel('Date')
    ax.set_ylabel('Predicted Total Cases')
    fig.patch.set_alpha(0) # make the figure background transparent
    # plt.tight_layout();
    fig.savefig('XGB_Forecast_Only.png', dpi=300, bbox_inches='tight')
    files.download("XGB_Forecast_Only.png")

```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



Best performing model predicts two more peaks by the end of years 2009 and 2012

Visualize feature importance:

```
[94]: # Calculate Feature Importance
feature_importances = final_model.feature_importances_
importance = pd.Series(feature_importances, index=X_train_whole.columns) #
    ↪ always positive value?
importance.sort_values()
```

```
[94]: month                                0.017704
station_avg_temp_c_shift_18                0.041669
watery_shift_20                            0.045804
reanalysis_precip_amt_kg_per_m2_shift_8    0.046078
reanalysis_tdtr_k_shift_8                  0.064036
grassy_shift_20                           0.064926
soily_shift_20                             0.108604
station_min_temp_c_shift_18                0.148184
station_max_temp_c_shift_18                0.192583
reanalysis_specific_humidity_g_per_kg_shift_12 0.270411
dtype: float32
```

```
[95]: # Visualise Feature Importance
with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=(8,5))
    importance.sort_values().plot.barh(ax=ax);
    ax.set_title("Relative Importance of Features \n for Predicting Dengue Cases_
    ↪ \n", fontsize=18)
    ax.set_xlabel('Feature importance')
```


```

ax.set_ylabel('Features')
ax.set_yticks([0,1,2,3,4,5,6,7,8,9])
ax.set_yticklabels(['Month', 'Avg Temp', 'Watery', 'Precipitation', 'Daily Temp_
Range', 'Grassy', 'Soily', 'Min Temp', 'Max Temp', 'Humidity'])
fig.patch.set_alpha(0) # make the figure background transparent
# plt.tight_layout()
fig.savefig('XGB_FeatureImportance.png', dpi=300, bbox_inches='tight')
files.download("XGB_FeatureImportance.png")

```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



- Sustained humidity and sustained maximum temperature are the two most important features in predicting dengue cases.

8 LSTM Neural Network

- You need an input shape of 3D tensor with shape (batch_size, timesteps, input_dim)

Scale and transform the data for neural network:

- Step 1: fit the scaler on the TRAINING data


- Step 2: use the scaler to transform the TRAINING data
- Step 3: use the transformed training data to fit the predictive model
- Step 4: use the scaler to transform the TEST data
- Step 5: predict using the trained model and the transformed TEST data.

```
[128]: train_LSTM = train_final.copy()
```

```
[129]: # Create a subset of the bigger dataset
train_LSTM1 = train_LSTM[['total_cases', 'month', 'station_max_temp_c',
    ↪ 'station_min_temp_c',
    'station_avg_temp_c', 'reanalysis_tdtr_k',
    'reanalysis_specific_humidity_g_per_kg',
    'reanalysis_precip_amt_kg_per_m2', 'grassy', 'soily',
    'watery']].copy()
```

```
[130]: # test train split
train = train_LSTM1.head(750)
test = train_LSTM1.tail(train_LSTM1.shape[0] - 750)

fig, ax = plt.subplots(figsize=(8,4))
ax.plot(train.total_cases, label='train')
ax.plot(test.total_cases, label='valid')
ax.set_title('Train-Validation Split')
plt.legend();
```



```
[131]: # Train and test X,y:
```

```
# the double brackets here are to keep the y in a dataframe format, otherwise it
↳will be pandas Series
X_train, y_train = train.drop('total_cases', axis=1).copy(),
↳train[['total_cases']].copy()
X_test, y_test = test.drop('total_cases', axis=1).copy(), test[['total_cases']].
↳copy()
```

```
[132]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[132]: ((750, 10), (175, 10), (750, 1), (175, 1))
```

```
[133]: # scale X and y using the sklearn MinMaxScaler model, so that their range will
↳be from 0 to 1.
Xscaler = MinMaxScaler(feature_range=(0, 1))
Xscaler.fit(X_train)
scaled_X_train = Xscaler.transform(X_train)
print(scaled_X_train.shape)

Yscaler = MinMaxScaler(feature_range=(0, 1))
Yscaler.fit(y_train)
scaled_y_train = Yscaler.transform(y_train)
print(scaled_y_train.shape)

# We need the shape of y to be (n, ), where n is the number of rows.
# remove the second dimension from y so the shape changes from (n,1) to (n,)
scaled_y_train = scaled_y_train.reshape(-1)
print(scaled_y_train.shape)

# print(type(scaled_X_train)) # Making sure they are numpy arrays
# print(type(scaled_y_train))
```

```
(750, 10)
```

```
(750, 1)
```

```
(750,)
```

8.0.1 Transform with TimeseriesGenerator:

- The TimeseriesGenerator transforms the separate X and y into a structure of samples ready to train deep learning models.
- The shape should be (batch_size,n_input,n_features)
- If batch size is equal to 3, the model will input the 3 sample videos and only after that 3 inputs, it will update the weights

```
[134]: # Create the train data
b_size = 32 #len(X_train)           # Number of timeseries samples in each batch
n_input = 12                        # how many samples/rows/timesteps to look
↳in the past in order to forecast the next sample
```

```

n_features= scaled_X_train.shape[1] # how many predictors/Xs/features we have,
    ↳to predict y

train_generator = TimeseriesGenerator(scaled_X_train, scaled_y_train,
    ↳length=n_input, batch_size=b_size)

# The shape should be (batch_size,n_input,n_features)
print(train_generator[0][0].shape)

```

(32, 12, 10)

```

[135]: # Create the test data
scaled_X_test = Xscaler.transform(X_test)
test_generator = TimeseriesGenerator(scaled_X_test, np.zeros(len(X_test)),
    ↳length=n_input, batch_size=b_size)
print(test_generator[0][0].shape)

```

(32, 12, 10)

8.1 LSTM Model # 1:

- Vanilla LSTM with a single hidden layer of LSTM units, and an output layer used to make a prediction.

```

[136]: model = Sequential()
model.add(LSTM(10, activation='relu', input_shape=(n_input, n_features)))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mae')
model.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
lstm_10 (LSTM)	(None, 10)	840
dense_7 (Dense)	(None, 1)	11
Total params: 851		
Trainable params: 851		
Non-trainable params: 0		


```

[137]: # Patience number of 10: the number of epochs to wait before early stop if no
    ↳progress on the validation set.
early_stop = EarlyStopping(monitor='loss', patience=10,
    ↳restore_best_weights=True)

```

```
[138]: # fit the model and plot the losses
def LSTM_fit_plotloss(train, test):
    random.seed(48)
    model.fit(train, epochs=100, verbose = 0, validation_data=test,
    ↪callbacks=[early_stop])
    loss_per_epoch = model.history.history['loss']
    val_loss_per_epoch = model.history.history['val_loss']
    plt.plot(range(len(loss_per_epoch)),loss_per_epoch, label = 'train_loss');
    plt.plot(range(len(val_loss_per_epoch)),val_loss_per_epoch, label =
    ↪'test_loss')
    plt.legend()
```

```
[139]: LSTM_fit_plotloss(train_generator, test_generator)
```



- Train and test loss converge after 60 epochs.
- Now the model is ready to use and we can make predictions on the train and test set.

```
[140]: # Make predictions on train and test
y_train_pred_scaled = model.predict(train_generator)
y_train_pred = Yscaler.inverse_transform(y_train_pred_scaled)
train_prediction = pd.DataFrame({'y_true':y_train['total_cases'].values[n_input:
    ↪],
```

```

        'y_pred':y_train_pred.ravel()},
        index=y_train.index[n_input:])

y_test_pred_scaled = model.predict(test_generator)
y_test_pred = Yscaler.inverse_transform(y_test_pred_scaled)
test_prediction = pd.DataFrame({'y_true':y_test['total_cases'].values[n_input:],
        'y_pred':y_test_pred.ravel()},
        index=y_test.index[n_input:])

```


24/24 [=====] - 0s 2ms/step

6/6 [=====] - 0s 3ms/step

```

[141]: forecast_graph(train_LSTM.total_cases, train_prediction.y_pred, test_prediction.
        ↪y_pred)

```



```

[142]: # Print the scores for both train and test
final_scores(train_prediction.y_true, train_prediction.y_pred,
        test_prediction.y_true, test_prediction.y_pred)

```

MAE_train: 22.701994

MAE_test: 19.718791

RMSE_train: 54.589900

RMSE_test: 33.485085

Summary:

- The model captures the basic seasonality, while missing all the individual outbreaks - peaks for not only test but also train.
- Let's make the model deeper , more complex by adding layers and neurons.

8.2 LSTM #2:

- A deeper model with more layers and neurons


```
[143]: # instantiate the model

model = Sequential()
model.add(LSTM(128, activation='relu', input_shape=(n_input, n_features),
    ↪return_sequences=True)) # returns a sequence of vectors of dimension 64
model.add(Dropout(0.2))
model.add(LSTM(64, activation='relu', return_sequences=True)) # returns a
    ↪sequence of vectors of dimension 32
model.add(Dropout(0.2))
model.add(LSTM(32)) # return a single vector of dimension 32
model.add(Dropout(0.2))
model.add(Dense(10))
model.add(Dropout(0.2))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mae')
model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
lstm_11 (LSTM)	(None, 12, 128)	71168
dropout_11 (Dropout)	(None, 12, 128)	0
lstm_12 (LSTM)	(None, 12, 64)	49408
dropout_12 (Dropout)	(None, 12, 64)	0
lstm_13 (LSTM)	(None, 32)	12416
dropout_13 (Dropout)	(None, 32)	0
dense_8 (Dense)	(None, 10)	330
dropout_14 (Dropout)	(None, 10)	0
dense_9 (Dense)	(None, 1)	11
Total params: 133,333		
Trainable params: 133,333		
Non-trainable params: 0		


```
[144]: LSTM_fit_plotloss(train_generator, test_generator)
```




```
[145]: y_train_pred_scaled = model.predict(train_generator)
y_train_pred = Yscaler.inverse_transform(y_train_pred_scaled)
train_prediction = pd.DataFrame({'y_true':y_train['total_cases'].values[n_input:
↪],
                                'y_pred':y_train_pred.ravel()},
                                index=y_train.index[n_input:])

y_test_pred_scaled = model.predict(test_generator)
y_test_pred = Yscaler.inverse_transform(y_test_pred_scaled)
test_prediction = pd.DataFrame({'y_true':y_test['total_cases'].values[n_input:],
                                'y_pred':y_test_pred.ravel()},
                                index=y_test.index[n_input:])
```

24/24 [=====] - 1s 10ms/step

6/6 [=====] - 0s 8ms/step

```
[146]: forecast_graph(train_LSTM1.total_cases, train_prediction.y_pred,
↪test_prediction.y_pred)
```



```
[147]: final_scores(train_prediction.y_true, train_prediction.y_pred, test_prediction.
        ↪ y_true, test_prediction.y_pred)
```

MAE_train: 22.583060

MAE_test: 19.037010

RMSE_train: 52.273007

RMSE_test: 31.049264

Summary:

- Increasing complexity improved the MAE and RMSE scores very slightly and it was not enough to be able to detect the individual peaks /outbreaks for train and test.
- Next, let's use lagged variables to see if that would make a difference like it did for NBR and XGBoost.

8.3 LSTM #3

- Lagged variables and a deeper model


```
[148]: train_LSTM3 = train_LSTM[['total_cases', 'month',
                                'station_max_temp_c_shift_18'],
                                ↪ 'station_min_temp_c_shift_18',
                                'station_avg_temp_c_shift_18'],
                                ↪ 'reanalysis_tdtr_k_shift_8', 'reanalysis_specific_humidity_g_per_kg_shift_12',
                                'reanalysis_precip_amt_kg_per_m2_shift_8'],
                                ↪ 'soily_shift_20', 'grassy_shift_20',
                                'watery_shift_20']].copy()

train_data = train_LSTM3.head(750)
```

```
test_data = train_LSTM3.tail(train_LSTM3.shape[0] - 750)

fig, ax = plt.subplots(figsize=(8,4))
ax.plot(train_data.total_cases, label='train')
ax.plot(test_data.total_cases, label='test')
ax.set_title('Train-Test Split');
plt.legend()

X_train, y_train = train_data.drop('total_cases', axis=1).copy(),
    ↪ train_data[['total_cases']].copy()
X_test, y_test = test_data.drop('total_cases', axis=1).copy(),
    ↪ test_data[['total_cases']].copy()
```



```
[149]: # scale X and y using the sklearn MinMaxScaler model, so that their range will
    ↪ be from 0 to 1.
Xscaler = MinMaxScaler(feature_range=(0, 1))
Xscaler.fit(X_train)
scaled_X_train = Xscaler.transform(X_train)
print(scaled_X_train.shape)

Yscaler = MinMaxScaler(feature_range=(0, 1))
Yscaler.fit(y_train)
scaled_y_train = Yscaler.transform(y_train)
print(scaled_y_train.shape)

# We need the shape of y to be (n, ), where n is the number of rows.
```

```
# remove the second dimension from y so the shape changes from (n,1) to (n,)
scaled_y_train = scaled_y_train.reshape(-1)
print(scaled_y_train.shape)
```

```
(750, 10)
(750, 1)
(750,)
```

```
[150]: b_size = 32                                # len(X_train) Number of timeseries samples
        ↪ in each batch
n_input = 12                                     # how many samples/rows/timesteps to look
        ↪ in the past in order to forecast the next sample
n_features= scaled_X_train.shape[1] # how many predictors/Xs/features we have
        ↪ to predict y

train_generator = TimeseriesGenerator(scaled_X_train, scaled_y_train,
        ↪ length=n_input, batch_size=b_size)

# The shape should be (batch_size,n_input,n_features)
print(train_generator[0][0].shape)
```

```
(32, 12, 10)
```

```
[151]: # create the validation data
scaled_X_test = Xscaler.transform(X_test)
test_generator = TimeseriesGenerator(scaled_X_test, np.zeros(len(X_test)),
        ↪ length=n_input, batch_size=b_size)
print(test_generator[0][0].shape)
```

```
(32, 12, 10)
```

```
[152]: model = Sequential()
model.add(LSTM(512, activation='relu', input_shape=(n_input, n_features),
        ↪ return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(256, activation='relu', return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(128, activation='relu', return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(64, activation='relu', return_sequences=True)) # returns a
        ↪ sequence of vectors of dimension 64
model.add(Dropout(0.2))
model.add(LSTM(32)) # return a single vector of dimension 32
model.add(Dropout(0.2))
model.add(Dense(32))
model.add(Dropout(0.2))
model.add(Dense(10))
```

```

model.add(Dropout(0.2))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mae')
model.summary()

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
lstm_14 (LSTM)	(None, 12, 512)	1071104
dropout_15 (Dropout)	(None, 12, 512)	0
lstm_15 (LSTM)	(None, 12, 256)	787456
dropout_16 (Dropout)	(None, 12, 256)	0
lstm_16 (LSTM)	(None, 12, 128)	197120
dropout_17 (Dropout)	(None, 12, 128)	0
lstm_17 (LSTM)	(None, 12, 64)	49408
dropout_18 (Dropout)	(None, 12, 64)	0
lstm_18 (LSTM)	(None, 32)	12416
dropout_19 (Dropout)	(None, 32)	0
dense_10 (Dense)	(None, 32)	1056
dropout_20 (Dropout)	(None, 32)	0
dense_11 (Dense)	(None, 10)	330
dropout_21 (Dropout)	(None, 10)	0
dense_12 (Dense)	(None, 1)	11

```


=====
Total params: 2,118,901
Trainable params: 2,118,901
Non-trainable params: 0
=====

```

```

[153]: # fit the model and plot the losses
LSTM_fit_plotloss(train_generator, test_generator)

```




```
[154]: # create predictions
y_train_pred_scaled = model.predict(train_generator)
y_train_pred = Yscaler.inverse_transform(y_train_pred_scaled)
train_prediction = pd.DataFrame({'y_true':y_train['total_cases'].values[n_input:
↪],
                                'y_pred':y_train_pred.ravel()},
                                index=y_train.index[n_input:])

y_test_pred_scaled = model.predict(test_generator)
y_test_pred = Yscaler.inverse_transform(y_test_pred_scaled)
test_prediction = pd.DataFrame({'y_true':y_test['total_cases'].values[n_input:],
↪ 'y_pred':y_test_pred.ravel()},
                                index=y_test.index[n_input:])
```

24/24 [=====] - 2s 55ms/step

6/6 [=====] - 0s 51ms/step

```
[155]: forecast_graph(train_LSTM.total_cases, train_prediction['y_pred'],
↪ test_prediction['y_pred'])
```



```
[156]: final_scores(train_prediction.y_true, train_prediction.y_pred, test_prediction.
      ↪ y_true, test_prediction.y_pred)
```

MAE_train: 21.902427

MAE_test: 20.202962

RMSE_train: 51.791441

RMSE_test: 34.619751

Summary:

- Using the best performing lagged variables with increased neural network complexity improved the model very slightly but it is still unable to detect individual peaks/outbreaks.
- We probably need a much more complex model since we are not able to detect even the peaks of the train set.

8.4 Conclusions / Recommendations:

-
- Dengue cases rely on climate variables, but the relationship is complex.
 - Further models should take into consideration cumulative computations of climate features over a period rather than isolated numbers.
 - Climate change and global warming may make dengue outbreaks and similar mosquito born illnesses more deadly in the future.
 - Knowing the next outbreak would help countries to allocate more resources to the health care system for timely intervention.

8.5 Limitations, Improvements, Next Steps

- More recent data needs to be collected to achieve more accurate predictions.
- Since the relationship between dengue and climate is complex:
 - Nonlinear relationships need to be taken into account with more complex models.
 - More meaningful and complex climate related features need to be engineered.

8.5.1 Export as PDF:

```
[125]: # Packages required for using nbconvert PDF
# ! apt-get install texlive texlive-xetex texlive-latex-extra pandoc
# ! pip install pypandoc
# ! pip install nbconvert

[126]: # First you need to download a copy of the ipynb notebook and upload it back to
↳ the drive, it is placed under /content/
# ! jupyter nbconvert --to pdf /content/notebook_modeling.ipynb
```