

notebook_preprocessing

April 4, 2023

1 Predicting Dengue Cases

- Student name: *Aysu Erdemir*
- Student pace: *Flex*
- Scheduled project review date/time: *March, 2023*
- Instructor name: *Abhineet Kulkarni*

1.1 Overview:

Dengue fever is a mosquito-borne disease that occurs in tropical and sub-tropical parts of the world. In mild cases, symptoms are similar to the flu: fever, rash, and muscle and joint pain. In severe cases, dengue fever can cause severe bleeding, low blood pressure, and even death.

Because it is carried by mosquitoes, the transmission dynamics of dengue are related to climate variables such as temperature and precipitation; however the relationship to climate is known to be complex. The way the disease spreads and causes endemics has significant public health implications worldwide.

- CDC is interested in predicting local epidemics of dengue fever so that they can take necessary precautions and efforts before the next spike. They want to know if we can predict the number of dengue fever cases reported each week in San Juan, Puerto Rico.
- My goal is to build several machine learning models to forecast the upcoming weekly dengue cases as accurately as possible.

1.2 Business and Data Understanding

- The data was obtained from [DrivenData](#). The data set included weekly dengue case counts along with environmental data collected by various U.S. Federal Government agencies—from the Centers for Disease Control and Prevention to the National Oceanic and Atmospheric Administration in the U.S. Department of Commerce.
- The full dataset included cases from year 1990 to 2008. The data from 2008-2013 included only features without case counts.
- In this project I will be focusing on data on **Puerto Rico** only. The relevant variables/features included in the dataset are:

Target Feature: * `total_cases` - Weekly total dengue cases.

Predictive Features:

Date Indicators:

- `week_start_date` - Date given in yyyy-mm-dd format.

NOAA's GHCN daily climate data weather station measurements:

- `station_max_temp_c` - Maximum temperature
- `station_min_temp_c` - Minimum temperature
- `station_avg_temp_c` - Average temperature
- `station_precip_mm` - Total precipitation
- `station_diur_temp_rng_c` - Diurnal temperature range

PERSIANN satellite precipitation measurements (0.25x0.25 degree scale):

- `precipitation_amt_mm` - Total precipitation

NOAA's NCEP Climate Forecast System Reanalysis measurements (0.5x0.5 degree scale):

- `reanalysis_sat_precip_amt_mm` - Total precipitation
- `reanalysis_dew_point_temp_k` - Mean dew point temperature
- `reanalysis_air_temp_k` - Mean air temperature
- `reanalysis_relative_humidity_percent` - Mean relative humidity
- `reanalysis_specific_humidity_g_per_kg` - Mean specific humidity
- `reanalysis_precip_amt_kg_per_m2` - Total precipitation
- `reanalysis_max_air_temp_k` - Maximum air temperature
- `reanalysis_min_air_temp_k` - Minimum air temperature
- `reanalysis_avg_temp_k` - Average air temperature
- `reanalysis_tdtr_k` - Diurnal temperature range

Satellite vegetation -greenness - Normalized difference vegetation index (NDVI) - NOAA's CDR Normalized Difference Vegetation Index (0.5x0.5 degree scale) measurements:

- `ndvi_se` - Pixel southeast of city centroid
- `ndvi_sw` - Pixel southwest of city centroid
- `ndvi_ne` - Pixel northeast of city centroid
- `ndvi_nw` - Pixel northwest of city centroid

For example, when you have negative values, it's highly likely that it's water. On the other hand, if you have an NDVI value close to +1, there's a high possibility that it's dense green leaves. But when NDVI is close to zero, there are likely no green leaves and it could even be an urbanized area.

1.3 Preprocessing:

1.3.1 Null Replacement:

- Null values for the climate features - except the four ndvi features - were imputed with **interpolation** since the missing data points are sparse.
- Null values for the four ndvi features were imputed using **k-Nearest Neighbors - KNN** since there were bigger chunks of missing values.

1.3.2 Feature Engineering:

- **Create month and seasons:** Created new variables representing the month and seasons.

- Create **average_ndvi** and its **categorical** version: Created a new feature representing the average NDVI values using the four different locations. Then created a categorical version of **average_ndvi** to represent watery, soily, sparse_grassy areas.
- Create **shifts** and **rolled averages** for the main climate variables: Research seems to indicate that past sustained heat, precipitation or humidity impacts dengue cases more profoundly than the climate situation right at the time of cases.
 - Shifted the variables by 2 weeks to account for the mosquito to reach adulthood and the incubation period of the virus until someone tests positive.
 - Create rolled means with a range of lags to see the variable with the highest correlation. The lag with the highest correlation was kept in the final dataset.

Some initial thoughts:

Precipitation: - Mosquitos thrive wet climates, the wetter the better! - A rise in **accumulated rainfall** was shown to result in an increase in the number and quality of breeding sites. - Elevated relative risk of dengue was observed when the weekly average rainfall was more than 150 mm at **lagged weeks 12 to 20**. - However, **above a certain rainfall level**, suitable mosquito breeding sites can be exposed to flooding, so the population is likely to decrease in such cases.

Humidity: - Humidity generates conditions that are favorable to adult mosquitoes increasing the life of the mosquito.

- Humidity range of **60% - 90%** is the optimum moisture for growth and development of the *Aedes aegypti* mosquito.

Temperature: - Mosquitoes more likely to transmit dengue virus in hot weather. Mosquitoes function best at 80 degrees F. **Higher temperatures (73-85°F)** results in more rapid viral growth and higher levels of virus. - The relative risk of dengue fever increases when the weekly average temperature is high **at lagged weeks 5 to 18**. - Under **fluctuating temperatures**, the mosquitoes show lower levels of virus in their salivary glands. - **Extreme heat waves** can negatively impact mosquito life as very high temperatures reduce adult lifespan and egg survival, resulting in reduced *Aedes* population and lower risk of dengue transmission.

NDVI index: - The relationship between dengue epidemic and greenness indexes is not clear or consistent. - Some studies indicate increased density of vegetation to provide suitable habitat for the immature mosquitoes. - Others indicated that low vegetation cover areas with increased dengue incidence rates. This inconsistency may be explained by regional differences.

```
[1]: # Import required packages

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import missingno
from sklearn.preprocessing import StandardScaler
from sklearn.impute import KNNImputer
```

```
# check package versions when necessary:
# pd.__version__
```

2 Data Exploration:

```
[2]: from google.colab import files
      uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving dengue_labels_train.csv to dengue_labels_train.csv
 Saving dengue_features_train.csv to dengue_features_train.csv
 Saving dengue_features_test.csv to dengue_features_test.csv

```
[3]: # Read the Data
      train_features = pd.read_csv("dengue_features_train.csv")
      train_labels = pd.read_csv("dengue_labels_train.csv")
      test_features = pd.read_csv("dengue_features_test.csv")
```

```
[4]: train_features.head()
```

```
[4]:   year  weekofyear  week_start_date  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  1990           18        4/30/90  0.122600  0.103725  0.198483  0.177617
1  1990           19        5/7/90  0.169900  0.142175  0.162357  0.155486
2  1990           20        5/14/90  0.032250  0.172967  0.157200  0.170843
3  1990           21        5/21/90  0.128633  0.245067  0.227557  0.235886
4  1990           22        5/28/90  0.196200  0.262200  0.251200  0.247340

      precipitation_amt_mm  reanalysis_air_temp_k  reanalysis_avg_temp_k  ...  \
0                12.42        297.572857        297.742857  ...
1                22.82        298.211429        298.442857  ...
2                34.54        298.781429        298.878571  ...
3                15.36        298.987143        299.228571  ...
4                 7.52        299.518571        299.664286  ...

      reanalysis_precip_amt_kg_per_m2  reanalysis_relative_humidity_percent  \
0                32.00                        73.365714
1                17.94                        77.368571
2                26.10                        82.052857
3                13.90                        80.337143
4                12.20                        80.460000

      reanalysis_sat_precip_amt_mm  reanalysis_specific_humidity_g_per_kg  \
0                12.42                        14.012857
1                22.82                        15.372857
2                34.54                        16.848571
3                15.36                        16.672857
```

4 7.52 17.210000

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c	\
0	2.628571	25.442857	6.900000	
1	2.371429	26.714286	6.371429	
2	2.300000	26.714286	6.485714	
3	2.428571	27.471429	6.771429	
4	3.014286	28.942857	9.371429	

	station_max_temp_c	station_min_temp_c	station_precip_mm
0	29.4	20.0	16.0
1	31.7	22.2	8.6
2	32.2	22.8	41.4
3	33.3	23.3	4.0
4	35.0	23.9	5.8

[5 rows x 23 columns]

```
[5]: train_labels.head()
```

```
[5]:   year  weekofyear  total_cases
0  1990           18            4
1  1990           19            5
2  1990           20            4
3  1990           21            3
4  1990           22            6
```

```
[6]: test_features.head()
```

```
[6]:   year  weekofyear  week_start_date  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  2008           18      4/29/08 -0.0189 -0.018900  0.102729  0.091200
1  2008           19      5/6/08 -0.0180 -0.012400  0.082043  0.072314
2  2008           20      5/13/08 -0.0015         NaN  0.151083  0.091529
3  2008           21      5/20/08         NaN -0.019867  0.124329  0.125686
4  2008           22      5/27/08  0.0568  0.039833  0.062267  0.075914
```

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	...	\
0	78.60	298.492857	298.550000	...	
1	12.56	298.475714	298.557143	...	
2	3.66	299.455714	299.357143	...	
3	0.00	299.690000	299.728571	...	
4	0.76	299.780000	299.671429	...	

	reanalysis_precip_amt_kg_per_m2	reanalysis_relative_humidity_percent	\
0	25.37	78.781429	
1	21.83	78.230000	
2	4.12	78.270000	

3	2.20	73.015714
4	4.36	74.084286

	reanalysis_sat_precip_amt_mm	reanalysis_specific_humidity_g_per_kg \
0	78.60	15.918571
1	12.56	15.791429
2	3.66	16.674286
3	0.00	15.775714
4	0.76	16.137143

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c \
0	3.128571	26.528571	7.057143
1	2.571429	26.071429	5.557143
2	4.428571	27.928571	7.785714
3	4.342857	28.057143	6.271429
4	3.542857	27.614286	7.085714

	station_max_temp_c	station_min_temp_c	station_precip_mm
0	33.3	21.7	75.2
1	30.0	22.2	34.3
2	32.8	22.8	3.0
3	33.3	24.4	0.3
4	33.3	23.3	84.1

[5 rows x 23 columns]

```
[7]: # Merge the features dataset with the labels dataset (total cases) to form
      ↪ train.
train = pd.merge(train_features, train_labels, on=["year", "weekofyear"])
train.head()
```

	year	weekofyear	week_start_date	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw \
0	1990	18	4/30/90	0.122600	0.103725	0.198483	0.177617
1	1990	19	5/7/90	0.169900	0.142175	0.162357	0.155486
2	1990	20	5/14/90	0.032250	0.172967	0.157200	0.170843
3	1990	21	5/21/90	0.128633	0.245067	0.227557	0.235886
4	1990	22	5/28/90	0.196200	0.262200	0.251200	0.247340

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k ... \
0	12.42	297.572857	297.742857 ...
1	22.82	298.211429	298.442857 ...
2	34.54	298.781429	298.878571 ...
3	15.36	298.987143	299.228571 ...
4	7.52	299.518571	299.664286 ...

	reanalysis_relative_humidity_percent	reanalysis_sat_precip_amt_mm \
0	73.365714	12.42

1	77.368571	22.82
2	82.052857	34.54
3	80.337143	15.36
4	80.460000	7.52

	reanalysis_specific_humidity_g_per_kg	reanalysis_tdtr_k \
0	14.012857	2.628571
1	15.372857	2.371429
2	16.848571	2.300000
3	16.672857	2.428571
4	17.210000	3.014286

	station_avg_temp_c	station_diur_temp_rng_c	station_max_temp_c \
0	25.442857	6.900000	29.4
1	26.714286	6.371429	31.7
2	26.714286	6.485714	32.2
3	27.471429	6.771429	33.3
4	28.942857	9.371429	35.0

	station_min_temp_c	station_precip_mm	total_cases
0	20.0	16.0	4
1	22.2	8.6	5
2	22.8	41.4	4
3	23.3	4.0	3
4	23.9	5.8	6

[5 rows x 24 columns]

```
[8]: # check row and column numbers
print(train_features.shape)
print(train_labels.shape)
print(train.shape)
```

(936, 23)

(936, 3)

(936, 24)

```
[9]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 936 entries, 0 to 935
```

```
Data columns (total 24 columns):
```

#	Column	Non-Null Count	Dtype
0	year	936 non-null	int64
1	weekofyear	936 non-null	int64
2	week_start_date	936 non-null	object
3	ndvi_ne	745 non-null	float64

```

4   ndvi_nw                887 non-null    float64
5   ndvi_se                917 non-null    float64
6   ndvi_sw                917 non-null    float64
7   precipitation_amt_mm   927 non-null    float64
8   reanalysis_air_temp_k  930 non-null    float64
9   reanalysis_avg_temp_k  930 non-null    float64
10  reanalysis_dew_point_temp_k  930 non-null    float64
11  reanalysis_max_air_temp_k  930 non-null    float64
12  reanalysis_min_air_temp_k  930 non-null    float64
13  reanalysis_precip_amt_kg_per_m2  930 non-null    float64
14  reanalysis_relative_humidity_percent  930 non-null    float64
15  reanalysis_sat_precip_amt_mm  927 non-null    float64
16  reanalysis_specific_humidity_g_per_kg  930 non-null    float64
17  reanalysis_tdtr_k      930 non-null    float64
18  station_avg_temp_c     930 non-null    float64
19  station_diur_temp_rng_c  930 non-null    float64
20  station_max_temp_c     930 non-null    float64
21  station_min_temp_c     930 non-null    float64
22  station_precip_mm      930 non-null    float64
23  total_cases            936 non-null    int64
dtypes: float64(20), int64(3), object(1)
memory usage: 182.8+ KB

```

```
[10]: train.describe()
```

```

[10]:
   count      year  weekofyear  ndvi_ne  ndvi_nw  ndvi_se  \
count    936.000000  936.000000  745.000000  887.000000  917.000000
mean    1998.826923   26.503205   0.057925   0.067469   0.177655
std         5.212076   15.021909   0.107153   0.092479   0.057166
min     1990.000000    1.000000  -0.406250  -0.456100  -0.015533
25%     1994.000000   13.750000   0.004500   0.016425   0.139283
50%     1999.000000   26.500000   0.057700   0.068075   0.177186
75%     2003.000000   39.250000   0.111100   0.115200   0.212557
max     2008.000000   53.000000   0.493400   0.437100   0.393129

   ndvi_sw  precipitation_amt_mm  reanalysis_air_temp_k  \
count    917.000000          927.000000          930.000000
mean      0.165956          35.470809          299.163653
std      0.056073          44.606137           1.236429
min     -0.063457           0.000000          295.938571
25%      0.129157           0.000000          298.195000
50%      0.165971          20.800000          299.254286
75%      0.202771          52.180000          300.132857
max      0.381420          390.600000          302.200000

   reanalysis_avg_temp_k  reanalysis_dew_point_temp_k  ...  \
count                930.000000                930.000000  ...

```


mean	299.276920	295.109519	...
std	1.218637	1.569943	...
min	296.114286	289.642857	...
25%	298.300000	293.847857	...
50%	299.378571	295.464286	...
75%	300.228571	296.418929	...
max	302.164286	297.795714	...

	reanalysis_relative_humidity_percent	reanalysis_sat_precip_amt_mm	\
count	930.000000	927.000000	
mean	78.568181	35.470809	
std	3.389488	44.606137	
min	66.735714	0.000000	
25%	76.246071	0.000000	
50%	78.667857	20.800000	
75%	80.963214	52.180000	
max	87.575714	390.600000	

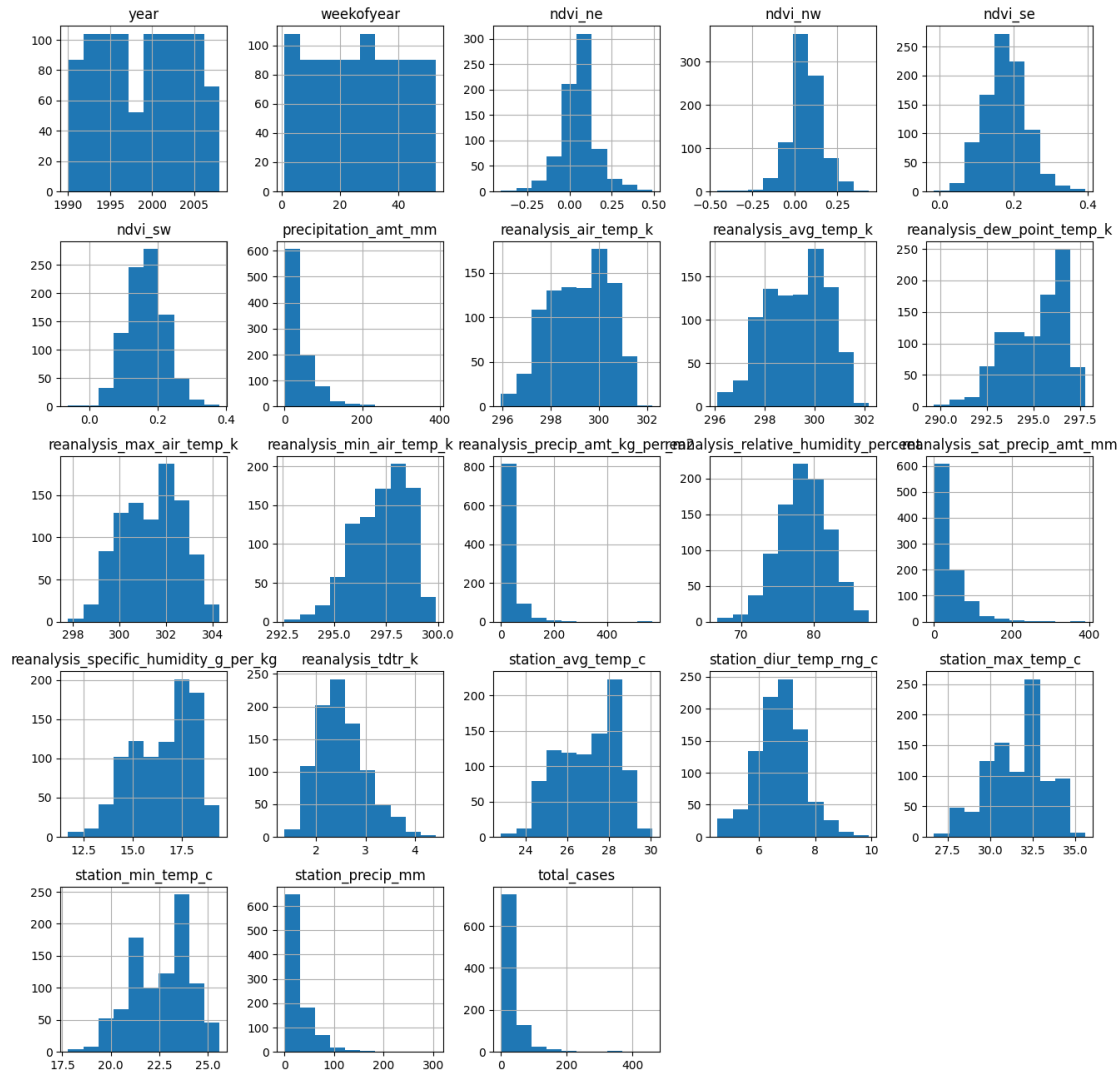
	reanalysis_specific_humidity_g_per_kg	reanalysis_tdtr_k	\
count	930.000000	930.000000	
mean	16.552409	2.516267	
std	1.560923	0.498892	
min	11.715714	1.357143	
25%	15.236429	2.157143	
50%	16.845714	2.457143	
75%	17.858571	2.800000	
max	19.440000	4.428571	

	station_avg_temp_c	station_diur_temp_rng_c	station_max_temp_c	\
count	930.000000	930.000000	930.000000	
mean	27.006528	6.757373	31.607957	
std	1.415473	0.835993	1.717297	
min	22.842857	4.528571	26.700000	
25%	25.842857	6.200000	30.600000	
50%	27.228571	6.757143	31.700000	
75%	28.185714	7.285714	32.800000	
max	30.071429	9.914286	35.600000	

	station_min_temp_c	station_precip_mm	total_cases
count	930.000000	930.000000	936.000000
mean	22.600645	26.785484	34.180556
std	1.506277	29.325811	51.381372
min	17.800000	0.000000	0.000000
25%	21.700000	6.825000	9.000000
50%	22.800000	17.750000	19.000000
75%	23.900000	35.450000	37.000000
max	25.600000	305.900000	461.000000

[8 rows x 23 columns]

```
[11]: # Let's see the distribution of the data for each one of the variables.  
train.hist(figsize = (16,16));
```



```
[12]: # See the distribution of case counts per each successive year:  
with plt.style.context('seaborn-talk'):  
    fig, ax = plt.subplots(figsize=(11,7))  
  
    years = np.arange(1990,2008,1).astype(int)  
    for year in years:  
        sns.lineplot(data=train[train.year == year],
```


```

x="week_start_date", y='total_cases', ax = ax, label = 'year)
ax.get_xaxis().set_ticks([]) # hide week_start_date
ax.set_xlabel('Time')
ax.set_ylabel('Total Cases')
ax.set_title('Total Dengue Cases on a Yearly Basis')
fig.patch.set_alpha(0) # make the figure background transparent
plt.legend(bbox_to_anchor=(1.12, 1), loc="upper right");
fig.savefig('total_cases_years.png', dpi=300)
files.download("total_cases_years.png")

```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



2.0.1 Summary:

- The data is from 1990 through 2008 with peak outbreaks at certain years.
- Total weekly cases range from 0 to 461.
- There are many null values, especially in ndvi index values.
- Many of the temperature features coded more than once in celcius and fehrenheit using different data sources.
- All variables appear as numerical except for week_start_date which is an object / string.
- All variables appear as normally distributed except target variable and precipitation measures.

3 Check for null values:

```
[13]: train.isnull().sum()  
# There are many null values
```

```
[13]: year                                0  
      weekofyear                          0  
      week_start_date                     0  
      ndvi_ne                             191  
      ndvi_nw                             49  
      ndvi_se                             19  
      ndvi_sw                             19  
      precipitation_amt_mm                 9  
      reanalysis_air_temp_k                6  
      reanalysis_avg_temp_k                6  
      reanalysis_dew_point_temp_k          6  
      reanalysis_max_air_temp_k            6  
      reanalysis_min_air_temp_k            6  
      reanalysis_precip_amt_kg_per_m2      6  
      reanalysis_relative_humidity_percent 6  
      reanalysis_sat_precip_amt_mm         9  
      reanalysis_specific_humidity_g_per_kg 6  
      reanalysis_tdtr_k                    6  
      station_avg_temp_c                   6  
      station_diur_temp_rng_c              6  
      station_max_temp_c                   6  
      station_min_temp_c                   6  
      station_precip_mm                    6  
      total_cases                          0  
      dtype: int64
```

```
[14]: # Proportion of null values for each variable:  
      nulls = ((train.isnull().sum()*100) / len(train_features)).  
             ↪sort_values(ascending=False)  
      nulls[nulls > 0]
```

```
[14]: ndvi_ne                20.405983  
      ndvi_nw                5.235043  
      ndvi_se                2.029915  
      ndvi_sw                2.029915  
      precipitation_amt_mm    0.961538  
      reanalysis_sat_precip_amt_mm 0.961538  
      reanalysis_min_air_temp_k 0.641026  
      reanalysis_precip_amt_kg_per_m2 0.641026  
      station_avg_temp_c      0.641026  
      reanalysis_tdtr_k       0.641026  
      reanalysis_specific_humidity_g_per_kg 0.641026
```

```

reanalysis_relative_humidity_percent    0.641026
reanalysis_dew_point_temp_k            0.641026
station_min_temp_c                    0.641026
reanalysis_max_air_temp_k              0.641026
station_max_temp_c                    0.641026
reanalysis_avg_temp_k                  0.641026
reanalysis_air_temp_k                  0.641026
station_precip_mm                      0.641026
station_diur_temp_rng_c                0.641026
dtype: float64

```

3.0.1 Display missing values:


- **Missingno** library offers a very nice way to visualize the distribution of Null values.

```

[15]: # Display null values across all rows/columns to check for specific patterns
      ↪ for the absence of data:
missingno.matrix(train_features , figsize=(14, 7), fontsize = 10)

```


[15]: <Axes: >




```

[16]: missingno.bar(train_features, figsize = (14,7), fontsize = 12);

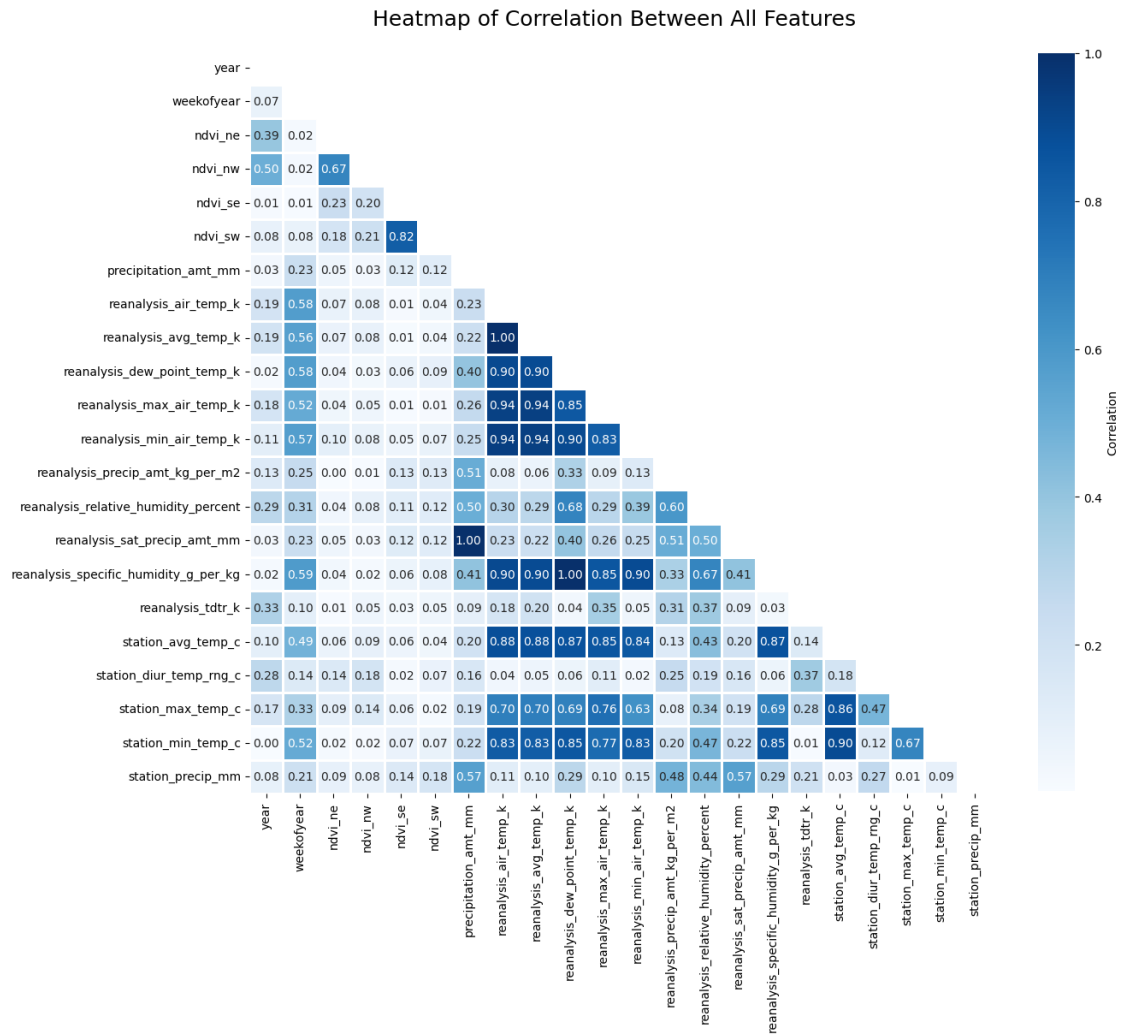
```



```
[17]: # Check the null matrix for the four variables with most null values to see if
      ↪ there is a pattern
missingno.matrix(train_features[['ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw']],
      ↪ figsize=(14, 7), fontsize = 12);
```



```
[18]: # Heat Map showing the correlation between all variables including the target
corr = train_features.corr().abs()
fig, ax = plt.subplots(figsize=(14,14))
matrix = np.triu(corr) # Getting the Lower Triangle of the correlation matrix
cbar_kws={"label": "Correlation", "shrink":0.8}
heatmap = sns.heatmap(data = corr, linewidths = 1, square= True,
    cmap='Blues', ax=ax, annot=True, mask=matrix, fmt= ".2f", cbar_kws=cbar_kws)
fig.suptitle('Heatmap of Correlation Between All Features', fontsize=18, y=.84,
    x = .43);
```



- There are strong correlations among the majority of the variables.
- Null values for most climate measures are scarce but ndvi indexes have null values in bigger chunks.

- All feature variables seem to follow more or less a normal distribution except for precipitation features.
- Precipitation follows a positive skew, where there are rare but extremely high values.

4 Feature Engineering: Null Replacement

4.0.1 Null replacement using interpolation and predictive modeling:

- We will replace the null values of all climate features except ndvi variables using **interpolation** since the missing data points are scarce.
- We will replace the null values of features for ndvi variables using **k-Nearest Neighbors** since there are bigger chunks of missing values.

```
[19]: # Null replacement with interpolation for the below variables:

train_features_interpolated = train_features

vars_to_interpolate = ['precipitation_amt_mm', 'reanalysis_air_temp_k',
                       'reanalysis_avg_temp_k', 'reanalysis_dew_point_temp_k',
                       'reanalysis_max_air_temp_k', 'reanalysis_min_air_temp_k',
                       'reanalysis_precip_amt_kg_per_m2',
                       'reanalysis_relative_humidity_percent', 'reanalysis_sat_precip_amt_mm',
                       'reanalysis_specific_humidity_g_per_kg', 'reanalysis_tdtr_k',
                       'station_avg_temp_c', 'station_diur_temp_rng_c', 'station_max_temp_c',
                       'station_min_temp_c', 'station_precip_mm']

for var in vars_to_interpolate:
    train_features_interpolated[var].interpolate(method='linear',
    ↪limit_direction='forward', inplace = True)
```

Null replacement using KNN neighbours for the remaining ndvi variables:

- The default distance measure is a Euclidean distance measure that is NaN aware.
- The number of neighbors is set to five by default and can be configured by the “n_neighbors” argument.
- `week_start_date` column was dropped and the data was scaled before imputing using KNN neighbours.
- After imputation, reversed the scaling to bring the values back to original scaling.

```
[20]: # Drop `week_start_date` since we cannot work with this variable.
train_features_interpolated = train_features_interpolated.
    ↪drop("week_start_date", axis = 1)
```

```
[21]: # Scale the data first using StandardScaler
scaler = StandardScaler()
train_features_scaled = pd.DataFrame(scaler.
    ↪fit_transform(train_features_interpolated), columns =
    ↪train_features_interpolated.columns)
```



```

# Define imputer
imputer = KNNImputer(n_neighbors=5)
# The fit imputer is applied to the dataset to create a copy of the dataset
↳with all missing values for each column replaced with an estimated value.
train_features_imputed= pd.DataFrame(imputer.
↳fit_transform(train_features_scaled), columns = train_features_scaled.
↳columns)
# inverse the Standard Scaling
train_features_full = pd.DataFrame(scaler.
↳inverse_transform(train_features_imputed), columns = train_features_imputed.
↳columns)
train_features_full.head()

```

```

[21]:      year  weekofyear  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  1990.0         18.0  0.122600  0.103725  0.198483  0.177617
1  1990.0         19.0  0.169900  0.142175  0.162357  0.155486
2  1990.0         20.0  0.032250  0.172967  0.157200  0.170843
3  1990.0         21.0  0.128633  0.245067  0.227557  0.235886
4  1990.0         22.0  0.196200  0.262200  0.251200  0.247340

      precipitation_amt_mm  reanalysis_air_temp_k  reanalysis_avg_temp_k  \
0                12.42          297.572857          297.742857
1                22.82          298.211429          298.442857
2                34.54          298.781429          298.878571
3                15.36          298.987143          299.228571
4                 7.52          299.518571          299.664286

      reanalysis_dew_point_temp_k  ...  reanalysis_precip_amt_kg_per_m2  \
0          292.414286  ...                32.00
1          293.951429  ...                17.94
2          295.434286  ...                26.10
3          295.310000  ...                13.90
4          295.821429  ...                12.20

      reanalysis_relative_humidity_percent  reanalysis_sat_precip_amt_mm  \
0                73.365714                12.42
1                77.368571                22.82
2                82.052857                34.54
3                80.337143                15.36
4                80.460000                 7.52

      reanalysis_specific_humidity_g_per_kg  reanalysis_tdtr_k  \
0                14.012857                2.628571
1                15.372857                2.371429
2                16.848571                2.300000
3                16.672857                2.428571
4                17.210000                3.014286

```

	station_avg_temp_c	station_diur_temp_rng_c	station_max_temp_c	\
0	25.442857	6.900000	29.4	
1	26.714286	6.371429	31.7	
2	26.714286	6.485714	32.2	
3	27.471429	6.771429	33.3	
4	28.942857	9.371429	35.0	

	station_min_temp_c	station_precip_mm
0	20.0	16.0
1	22.2	8.6
2	22.8	41.4
3	23.3	4.0
4	23.9	5.8

[5 rows x 22 columns]

```
[22]: # Making sure no null values remained
train_features_imputed.isna().sum().any() == 0
```

[22]: True

```
[23]: # Display the dataset for ndvi values before and after knn imputation
with plt.style.context('seaborn-talk'):
    fig, (ax1, ax2) = plt.subplots(nrows=2, ncols=1, figsize=(14,10))
    for var in ["ndvi_ne", "ndvi_nw", "ndvi_se", "ndvi_sw"]:
        train_features[150:300][var].plot.line(lw=1.2, ax = ax1)
        ax1.set_title('Vegetation Index over Time - Original Data')
        ax1.set_xlabel('Time')
        ax1.legend(loc='lower right')
    for var in ["ndvi_ne", "ndvi_nw", "ndvi_se", "ndvi_sw"]:
        train_features_full[150:300][var].plot.line(lw=1.2, ax = ax2)
        ax2.set_title('Vegetation Index over Time - Imputed Data')
        ax2.set_xlabel('Time')
        ax2.legend(loc='lower right')
    fig.tight_layout();
    fig.patch.set_alpha(0) # make the figure background transparent
    fig.savefig('KNN_ndvi.png', dpi=300, bbox_inches='tight')
    files.download("KNN_ndvi.png")
```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



```
[24]: # Merge the imputed dataset with the labels
train_full = pd.merge(train_features_full, train_labels, on=[ "year",
↪ "weekofyear" ])
train_full.head()
```

```
[24]:
```

	year	weekofyear	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	\
0	1990.0	18.0	0.122600	0.103725	0.198483	0.177617	
1	1990.0	19.0	0.169900	0.142175	0.162357	0.155486	
2	1990.0	20.0	0.032250	0.172967	0.157200	0.170843	
3	1990.0	21.0	0.128633	0.245067	0.227557	0.235886	
4	1990.0	22.0	0.196200	0.262200	0.251200	0.247340	

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	\
0	12.42	297.572857	297.742857	
1	22.82	298.211429	298.442857	
2	34.54	298.781429	298.878571	
3	15.36	298.987143	299.228571	
4	7.52	299.518571	299.664286	

	reanalysis_dew_point_temp_k	...	reanalysis_relative_humidity_percent	\
0	292.414286	...	73.365714	
1	293.951429	...	77.368571	

2	295.434286	...	82.052857
3	295.310000	...	80.337143
4	295.821429	...	80.460000

	reanalysis_sat_precip_amt_mm	reanalysis_specific_humidity_g_per_kg	\
0	12.42	14.012857	
1	22.82	15.372857	
2	34.54	16.848571	
3	15.36	16.672857	
4	7.52	17.210000	

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c	\
0	2.628571	25.442857	6.900000	
1	2.371429	26.714286	6.371429	
2	2.300000	26.714286	6.485714	
3	2.428571	27.471429	6.771429	
4	3.014286	28.942857	9.371429	

	station_max_temp_c	station_min_temp_c	station_precip_mm	total_cases
0	29.4	20.0	16.0	4
1	31.7	22.2	8.6	5
2	32.2	22.8	41.4	4
3	33.3	23.3	4.0	3
4	35.0	23.9	5.8	6

[5 rows x 23 columns]

4.1 Repeat all imputation steps for the test_features dataset:

```
[25]: # See how many null values present in the test dataset
test_features.isnull().sum()
```

```
[25]: year                0
weekofyear              0
week_start_date         0
ndvi_ne                 43
ndvi_nw                 11
ndvi_se                 1
ndvi_sw                 1
precipitation_amt_mm    2
reanalysis_air_temp_k   2
reanalysis_avg_temp_k   2
reanalysis_dew_point_temp_k 2
reanalysis_max_air_temp_k 2
reanalysis_min_air_temp_k 2
reanalysis_precip_amt_kg_per_m2 2
reanalysis_relative_humidity_percent 2
```

```

reanalysis_sat_precip_amt_mm          2
reanalysis_specific_humidity_g_per_kg  2
reanalysis_tdtr_k                      2
station_avg_temp_c                     2
station_diur_temp_rng_c                2
station_max_temp_c                     2
station_min_temp_c                     2
station_precip_mm                       2
dtype: int64

```

```

[26]: # Interpolation
test_features_interpolated = test_features
for var in vars_to_interpolate:
    test_features_interpolated[var].interpolate(method='linear',
    ↪ limit_direction='forward', inplace=True)

# Drop week_start_date:
test_features_interpolated = test_features_interpolated.drop("week_start_date",
    ↪ axis=1)

# Scale, imputer using KNN imputer, inverse scale
test_features_scaled = pd.DataFrame(scaler.
    ↪ fit_transform(test_features_interpolated), columns=
    ↪ test_features_interpolated.columns)
test_features_imputed = pd.DataFrame(imputer.
    ↪ fit_transform(test_features_scaled), columns=test_features_scaled.columns)
test_features_full = pd.DataFrame(scaler.
    ↪ inverse_transform(test_features_imputed), columns=test_features_imputed.
    ↪ columns)

# Making sure no null values remained
test_features_full.isna().sum().any() == 0

```

[26]: True

Full Imputed datasets are: * train_full * test_features_full

5 Feature Engineering: Feature Selection / Creation

```

[27]: train_featured = train_full.copy()

```

Let's create a new month variable and dummy coded season variables:

```

[28]: # Add `the week_start_date` column from the original dataset to the new dataset
train_featured['week_start_date'] = train_features['week_start_date']

```

```
[29]: # create a new month variable:
train_featured["week_start_date"] = pd.
    ↳to_datetime(train_featured["week_start_date"])
train_featured['month'] = train_featured['week_start_date'].dt.month
```

```
[30]: # create a new season variable:
seasons = ["winter", "winter", "spring", "spring", "spring",
           "summer", "summer", "summer", "fall", "fall", "fall", "winter"]

month_to_season = dict(zip(range(1,13), seasons))
month_to_season

train_featured['season'] = train_featured['month'].map(month_to_season)
```

```
[31]: # See the new variables:
train_featured[['week_start_date', 'month', 'season']].sample(5)
```

```
[31]:      week_start_date  month  season
400      1998-01-08        1  winter
208      1994-04-30        4  spring
167      1993-07-16        7  summer
363      1997-04-23        4  spring
852      2006-09-17        9   fall
```

```
[32]: # Get the season dummy coded
season_features = pd.get_dummies(train_featured['season'])
train_featured = pd.concat([train_featured, season_features], axis = 1)
train_featured.head()
```

```
[32]:      year  weekofyear  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  1990.0         18.0  0.122600  0.103725  0.198483  0.177617
1  1990.0         19.0  0.169900  0.142175  0.162357  0.155486
2  1990.0         20.0  0.032250  0.172967  0.157200  0.170843
3  1990.0         21.0  0.128633  0.245067  0.227557  0.235886
4  1990.0         22.0  0.196200  0.262200  0.251200  0.247340

      precipitation_amt_mm  reanalysis_air_temp_k  reanalysis_avg_temp_k  \
0              12.42          297.572857          297.742857
1              22.82          298.211429          298.442857
2              34.54          298.781429          298.878571
3              15.36          298.987143          299.228571
4               7.52          299.518571          299.664286


      reanalysis_dew_point_temp_k  ...  station_min_temp_c  station_precip_mm  \
0              292.414286  ...              20.0             16.0
1              293.951429  ...              22.2              8.6
2              295.434286  ...              22.8             41.4
```

3		295.310000	...			23.3		4.0
4		295.821429	...			23.9		5.8

	total_cases	week_start_date	month	season	fall	spring	summer	winter
0	4	1990-04-30	4	spring	0	1	0	0
1	5	1990-05-07	5	spring	0	1	0	0
2	4	1990-05-14	5	spring	0	1	0	0
3	3	1990-05-21	5	spring	0	1	0	0
4	6	1990-05-28	5	spring	0	1	0	0

[5 rows x 30 columns]

```
[33]: # Seasonality: See the distribution of case counts on a monthly basis:
with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=(10,4))
    sns.lineplot(data=train_featured, x="month", y='total_cases', ax = ax)
    # Maximum number of cases are usually are seen in the fall.
```




6 Feature elimination / selection:

```
[34]: # Show the correlated pairs starting with strongest correlations:
dataCorr = train_featured.drop('total_cases',axis =1).corr().abs()
dataCorr = dataCorr.mask(np.triu(np.ones(dataCorr.shape)).astype(bool)) #
    ↳convert upper triangle of values to NaN to remove repeated values from the
    ↳table
dataCorr = dataCorr.stack().reset_index().sort_values(0, ascending=False) #0 is
    ↳the column automatically generated by the stacking
dataCorr = dataCorr[(dataCorr[0]>.8) & (dataCorr[0]<1)]
```


207	station_avg_temp_c	0.898506
113	reanalysis_avg_temp_k	0.896420
116	reanalysis_min_air_temp_k	0.896376
44	reanalysis_avg_temp_k	0.895373
143	reanalysis_air_temp_k	0.880871
144	reanalysis_avg_temp_k	0.879118
151	reanalysis_specific_humidity_g_per_kg	0.869982
145	reanalysis_dew_point_temp_k	0.868837
188	station_avg_temp_c	0.865240
115	reanalysis_max_air_temp_k	0.853629
146	reanalysis_max_air_temp_k	0.852831
199	reanalysis_dew_point_temp_k	0.850479
205	reanalysis_specific_humidity_g_per_kg	0.849573
54	reanalysis_dew_point_temp_k	0.847654
147	reanalysis_min_air_temp_k	0.841300
197	reanalysis_air_temp_k	0.833158
201	reanalysis_min_air_temp_k	0.829792
65	reanalysis_max_air_temp_k	0.828665
198	reanalysis_avg_temp_k	0.827497
14	ndvi_se	0.820109

[35]: *# Show how strongly the features are correlated with the target variable - total cases:*

```
fig, ax = plt.subplots(figsize=(14,7))
train_featured.corr()['total_cases'].drop('total_cases').
    sort_values(ascending=False).plot.barh(ax=ax);
```



- Many of the temperature data are strongly correlated with one another.
- However, none of the features seem to have a strong relationship with the target variable - total_cases. Total_cases seems to only have weak correlations with other variables.
- We need to engineer some new features hoping they would have stronger relationship with

total_cases.

6.0.1 Select the best average temperature variable:

- station_avg_temp_c has the strongest correlation


```
[36]: # see the correlations among all average temperature variables:
train_featured[['total_cases', 'station_avg_temp_c', 'reanalysis_air_temp_k', 'reanalysis_avg_tem
↪corr()
```

```
[36]:
```

	total_cases	station_avg_temp_c	reanalysis_air_temp_k	\
total_cases	1.000000	0.194823	0.180311	
station_avg_temp_c	0.194823	1.000000	0.880871	
reanalysis_air_temp_k	0.180311	0.880871	1.000000	
reanalysis_avg_temp_k	0.173670	0.879118	0.997507	

	reanalysis_avg_temp_k
total_cases	0.173670
station_avg_temp_c	0.879118
reanalysis_air_temp_k	0.997507
reanalysis_avg_temp_k	1.000000

```
[37]: with plt.style.context('seaborn-talk'):
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=(14,4))
sns.regplot(data=train_featured, x="station_avg_temp_c", y='total_cases',
↪ax = ax1, label = 'station_avg_temp_c')
sns.regplot(data=train_featured, x="reanalysis_air_temp_k",
↪y='total_cases', ax = ax2, label = 'reanalysis_air_temp_k')
```



6.0.2 Select the best daily temperature change variable:

- reanalysis_tdtr_k has the strongest correlation
- Let's also replace the single outlier with a better value.


```
[38]: train_featured[['total_cases', 'station_diur_temp_rng_c', 'reanalysis_tdtr_k']].
      ↪corr()
```

```
[38]:
```

	total_cases	station_diur_temp_rng_c	\
total_cases	1.000000	0.035303	
station_diur_temp_rng_c	0.035303	1.000000	
reanalysis_tdtr_k	-0.067623	0.372414	

	reanalysis_tdtr_k
total_cases	-0.067623
station_diur_temp_rng_c	0.372414
reanalysis_tdtr_k	1.000000

```
[39]: with plt.style.context('seaborn-talk'):
      fig, ax = plt.subplots(figsize=(10,4))
      sns.regplot(data=train_featured, x="station_diur_temp_rng_c",
      ↪y='total_cases', ax = ax, label = 'station_diur_temp_rng_c')
      sns.regplot(data=train_featured, x="reanalysis_tdtr_k", y='total_cases', ax=
      ↪ax, label = 'reanalysis_tdtr_k')
```



```
[40]: # check out the outlier for 'reanalysis_tdtr_k':
      train_featured[train_featured['reanalysis_tdtr_k'] ==
      ↪train_featured['reanalysis_tdtr_k'].max()]
```

```
[40]:
```

	year	weekofyear	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	\
799	2005.0	36.0	0.0022	-0.0271	0.205029	0.220233	

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	\
799	23.3	301.465714	301.514286	

```

reanalysis_dew_point_temp_k ... station_min_temp_c station_precip_mm \
799          296.642857 ...          24.4          8.9

total_cases week_start_date month season fall spring summer winter
799          131      2005-09-10     9   fall     1     0     0     0

[1 rows x 30 columns]

```

```

[41]: # replace the outlier with the previous value in the series
train_featured = train_featured.replace(train_featured['reanalysis_tdtr_k'].
    ↳max(), method='ffill')
train_featured[799:800]

```

```

[41]:      year weekofyear ndvi_ne ndvi_nw ndvi_se ndvi_sw \
799  2005.0         36.0  0.0022 -0.0271  0.205029  0.220233

precipitation_amt_mm reanalysis_air_temp_k reanalysis_avg_temp_k \
799          23.3          301.465714          301.514286

reanalysis_dew_point_temp_k ... station_min_temp_c station_precip_mm \
799          296.642857 ...          24.4          8.9

total_cases week_start_date month season fall spring summer winter
799          131      2005-09-10     9   fall     1     0     0     0

[1 rows x 30 columns]

```

```

[42]: # check the correlations again:
train_featured[['total_cases', 'station_diur_temp_rng_c', 'reanalysis_tdtr_k']].
    ↳corr()

```

```

[42]:      total_cases station_diur_temp_rng_c \
total_cases          1.000000          0.035303
station_diur_temp_rng_c  0.035303          1.000000
reanalysis_tdtr_k      -0.073160          0.374047

reanalysis_tdtr_k
total_cases          -0.073160
station_diur_temp_rng_c  0.374047
reanalysis_tdtr_k          1.000000


```

```

[43]: with plt.style.context('seaborn-talk'):
fig, ax = plt.subplots(figsize=(10,4))
sns.regplot(data=train_featured, x="station_diur_temp_rng_c",
    ↳y='total_cases', ax = ax, label = 'station_diur_temp_rng_c')

```

```
sns.regplot(data=train_featured, x="reanalysis_tdtr_k", y='total_cases', ax=
↳ ax, label = 'reanalysis_tdtr_k')
```



6.0.3 Select the best humidity variable:

- reanalysis_specific_humidity_g_per_kg has the strongest correlation


```
[44]: train_featured[['total_cases', 'reanalysis_relative_humidity_percent',
↳
↳ 'reanalysis_specific_humidity_g_per_kg', 'reanalysis_dew_point_temp_k']].
↳ corr()
```

```
[44]:
```

	total_cases	\
total_cases	1.000000	
reanalysis_relative_humidity_percent	0.144404	
reanalysis_specific_humidity_g_per_kg	0.206942	
reanalysis_dew_point_temp_k	0.202807	
	reanalysis_relative_humidity_percent	\
total_cases	0.144404	
reanalysis_relative_humidity_percent	1.000000	
reanalysis_specific_humidity_g_per_kg	0.673010	
reanalysis_dew_point_temp_k	0.678116	
	reanalysis_specific_humidity_g_per_kg	\
total_cases	0.206942	
reanalysis_relative_humidity_percent	0.673010	
reanalysis_specific_humidity_g_per_kg	1.000000	
reanalysis_dew_point_temp_k	0.998533	

	reanalysis_dew_point_temp_k
total_cases	0.202807
reanalysis_relative_humidity_percent	0.678116
reanalysis_specific_humidity_g_per_kg	0.998533
reanalysis_dew_point_temp_k	1.000000

```
[45]: with plt.style.context('seaborn-talk'):
fig, (ax1, ax2, ax3) = plt.subplots(ncols = 3, nrows = 1, figsize=(16,4))
sns.regplot(data=train_featured, x="reanalysis_relative_humidity_percent",
↪y='total_cases', ax = ax1, label = 'reanalysis_relative_humidity_percent')
sns.regplot(data=train_featured, x="reanalysis_specific_humidity_g_per_kg",
↪y='total_cases', ax = ax2, label = 'reanalysis_specific_humidity_g_per_kg')
sns.regplot(data=train_featured, x="reanalysis_dew_point_temp_k",
↪y='total_cases', ax = ax3, label = 'reanalysis_dew_point_temp_k')
```



6.0.4 Select the best precipitation variable:

- reanalysis_precip_amt_kg_per_m2 has the strongest correlation

```
[46]: train_featured[['total_cases', 'reanalysis_sat_precip_amt_mm', 'station_precip_mm',
↪'reanalysis_precip_amt_kg_per_m2', 'precipitation_amt_mm']].
↪corr()
```

```
[46]:
```

	total_cases	reanalysis_sat_precip_amt_mm	\
total_cases	1.000000	0.060296	
reanalysis_sat_precip_amt_mm	0.060296	1.000000	
station_precip_mm	0.051883	0.566660	
reanalysis_precip_amt_kg_per_m2	0.106939	0.508274	
precipitation_amt_mm	0.060296	1.000000	

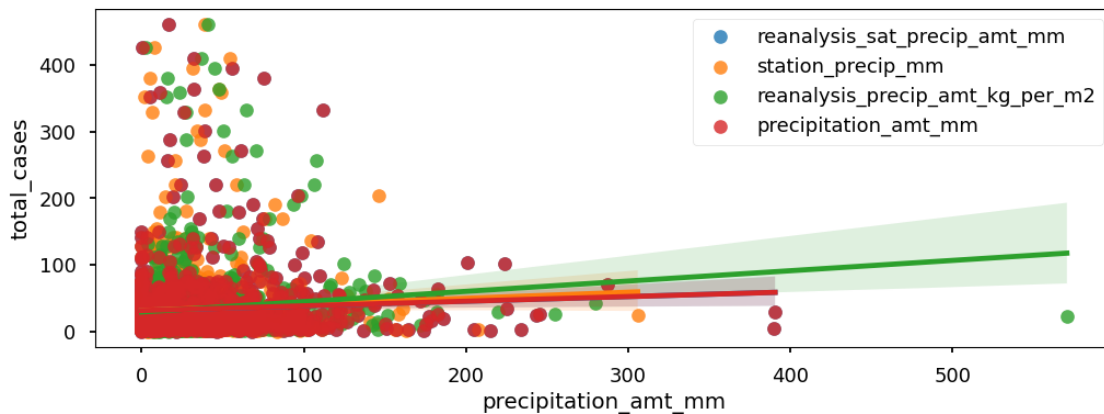
	station_precip_mm	\
total_cases	0.051883	
reanalysis_sat_precip_amt_mm	0.566660	

station_precip_mm	1.000000
reanalysis_precip_amt_kg_per_m2	0.477984
precipitation_amt_mm	0.566660

	reanalysis_precip_amt_kg_per_m2 \
total_cases	0.106939
reanalysis_sat_precip_amt_mm	0.508274
station_precip_mm	0.477984
reanalysis_precip_amt_kg_per_m2	1.000000
precipitation_amt_mm	0.508274

	precipitation_amt_mm
total_cases	0.060296
reanalysis_sat_precip_amt_mm	1.000000
station_precip_mm	0.566660
reanalysis_precip_amt_kg_per_m2	0.508274
precipitation_amt_mm	1.000000

```
[47]: with plt.style.context('seaborn-talk'):
fig, ax = plt.subplots(figsize=(12,4))
for var in ["reanalysis_sat_precip_amt_mm", "station_precip_mm",
↪ "reanalysis_precip_amt_kg_per_m2", "precipitation_amt_mm"]:
    sns.regplot(data=train_featured, x=var, y='total_cases', ax = ax, label =
↪ var)
plt.legend()
```



6.1 Summary - feature selection:

Let's focus on these variables below since they provide the highest correlations. Let's keep all the temperature variables same scale (celcius) for interpretability except the diurnal which gave better correlation in kelvin.

Using NOAA's NCEP Climate Forecast System Reanalysis measurements (0.5x0.5 degree scale):

Temperature variables: station_avg_temp_c, station_min_temp_c, station_max_temp_c, reanalysis_tdtr_k (Diurnal temperature range)

Humidity: reanalysis_specific_humidity_g_per_kg

Precipitation: reanalysis_precip_amt_kg_per_m2 (Total precipitation)


Vegetation: ndvi_ne, ndvi_nw, ndvi_se, ndvi_sw

```
[48]: train_featured.columns
```

```
[48]: Index(['year', 'weekofyear', 'ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw',  
        'precipitation_amt_mm', 'reanalysis_air_temp_k',  
        'reanalysis_avg_temp_k', 'reanalysis_dew_point_temp_k',  
        'reanalysis_max_air_temp_k', 'reanalysis_min_air_temp_k',  
        'reanalysis_precip_amt_kg_per_m2',  
        'reanalysis_relative_humidity_percent', 'reanalysis_sat_precip_amt_mm',  
        'reanalysis_specific_humidity_g_per_kg', 'reanalysis_tdtr_k',  
        'station_avg_temp_c', 'station_diur_temp_rng_c', 'station_max_temp_c',  
        'station_min_temp_c', 'station_precip_mm', 'total_cases',  
        'week_start_date', 'month', 'season', 'fall', 'spring', 'summer',  
        'winter'],  
        dtype='object')
```

```
[49]: train_featured = train_featured[['total_cases', 'year', 'weekofyear',  
        ↪ 'week_start_date',  
        'month', 'fall', 'spring', 'summer', 'winter',  
        'station_avg_temp_c', 'station_max_temp_c',  
        'station_min_temp_c', 'reanalysis_tdtr_k',  
        'reanalysis_specific_humidity_g_per_kg',  
        'reanalysis_precip_amt_kg_per_m2',  
        'ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw']]
```

```
[50]: fig, ax = plt.subplots(figsize=(12,6))  
train_featured.corr()['total_cases'].drop('total_cases').  
        ↪ sort_values(ascending=False).plot.barh(ax=ax);
```

```
[51]: # Plot how the total case numbers differ based on each climate feature:
columns = ['station_avg_temp_c', 'station_max_temp_c',
           'station_min_temp_c', 'reanalysis_precip_amt_kg_per_m2',
           'reanalysis_tdtr_k', 'reanalysis_specific_humidity_g_per_kg',
           'ndvi_ne', 'ndvi_nw',
           'ndvi_se', 'ndvi_sw']

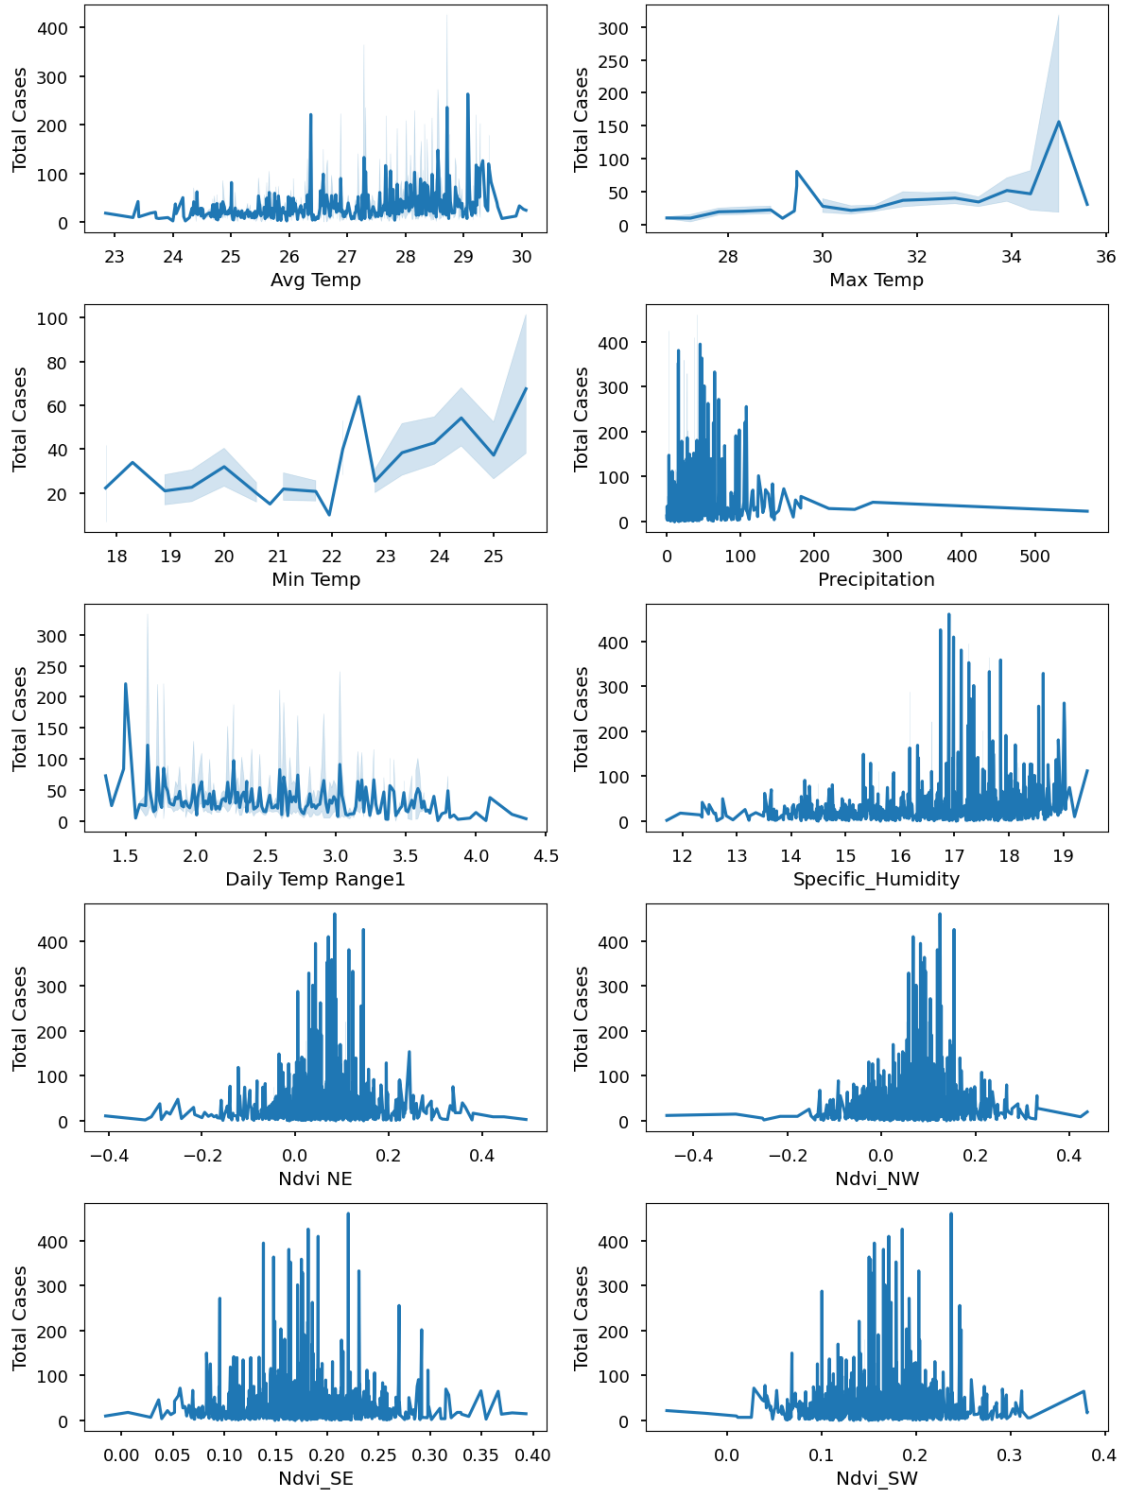
labels = ["Avg Temp", "Max Temp",
          "Min Temp", "Precipitation",
          "Daily Temp Range1", "Specific_Humidity",
          "Ndvi NE", "Ndvi_NW",
          "Ndvi_SE", "Ndvi_SW"]

nrows = 5
ncols = 2
with plt.style.context('seaborn-talk'):

    fig, ax_list = plt.subplots(nrows = nrows, ncols = ncols, figsize=(12,16))

    j=0
    for i in range(nrows):
        for u in range(ncols):
            sns.lineplot(data =train_featured, x=columns[j], y="total_cases",
            ↪ax = ax_list[i,u]) # need to use index for column because otherwise it does
            ↪not iterate.
            ax_list[i,u].set_xlabel(labels[j])
            ax_list[i,u].set_ylabel("Total Cases")
            j = j+1

    fig.tight_layout();
```



- For all the variables, the relationship seems stronger until case number 100 reaches, possible because there are fewer extra high dengue cases.
- Extremely high **average temp**, **maximum temp**, **temp range**, or **precipitation** seem to

impact dengue cases negatively.

- There is no clear linear relationship between ndvi variables and total cases.

6.2 Convert NDVI into Categorical variables:

- NDVI calculation range from -1 to 1. Negative values correspond to areas with water surfaces, manmade structures, rocks, clouds, snow. Bare soil usually falls within 0.1- 0.2 range. Plants will always have positive values between 0.2 and 1. Healthy, dense vegetation canopy should be above 0.5. Sparse vegetation will most likely fall within 0.2 to 0.5.
- Since there is no clear linear relationship between ndvi and total cases, let's create a categorical version of the variables. |
 - Let's first create a new feature representing the average NDVI values from the four different locations.
 - Then let's create a categorical version of the variable to represent **watery**, **soily**, **spare_grassy** areas.


```
[52]: # Create `ndvi_average` by taking the mean of all 4 coordinates.
train_featured['ndvi_average'] = (
    (train_featured['ndvi_ne']+train_featured['ndvi_nw']+
     train_featured['ndvi_se']+train_featured['ndvi_sw']))/4
```

```
[53]: # Let's check the distribution of values:
train_featured[['ndvi_ne','ndvi_nw','ndvi_se','ndvi_sw', 'ndvi_average']].
describe()
```

```
[53]:
```

	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	ndvi_average
count	936.000000	936.000000	936.000000	936.000000	936.000000
mean	0.057266	0.067853	0.177654	0.165855	0.117157
std	0.100001	0.090603	0.056694	0.055681	0.056231
min	-0.406250	-0.456100	-0.015533	-0.063457	-0.092565
25%	0.008050	0.018706	0.139862	0.129778	0.079570
50%	0.057667	0.068750	0.177171	0.165906	0.112724
75%	0.108288	0.115017	0.212336	0.202549	0.147122
max	0.493400	0.437100	0.393129	0.381420	0.342338

```
[54]: # Let's see the bottom and top 15th percentile for `ndvi_average`:
fig, ax = plt.subplots(figsize=(12,6))
sns.regplot(data = train_featured, x='ndvi_average', y="total_cases", ax = ax)
ax.axvline(x=train_featured['ndvi_average'].quantile(0.15), ymin=0, ymax=1,
           color='y', linestyle='--')
ax.axvline(x=train_featured['ndvi_average'].quantile(0.85), ymin=0, ymax=1,
           color='y', linestyle='--');
```




- Total cases seem to be low for low and high ndvi, and high for moderate ndvi

```
[55]: # Create a function to transform `ndvi_average` into a categorical variables.
# watery (dead plants) --> soily (unhealthy plants) --> grassy (healthy plants)
def get_ndvi_cat(x):
    if x < train_featured['ndvi_average'].quantile(0.15):
        return "watery"
    if x >= train_featured['ndvi_average'].quantile(0.15) and x <=
    ↪ train_featured['ndvi_average'].quantile(0.85):
        return "soily"
    else:
        return "grassy"
```

```
[56]: # Apply the transformation
train_featured["ndvi_average_cat"] = train_featured["ndvi_average"].
    ↪ apply(get_ndvi_cat)
```

```
[57]: # See the count values after transformation
train_featured['ndvi_average_cat'].value_counts().plot(kind='bar');
```



```
[58]: # Get dummy codes for 'ndvi_average_cat' and merge with the dataset:
ndvi_features = pd.get_dummies(train_featured['ndvi_average_cat'])
train_featured = pd.concat([train_featured, ndvi_features], axis = 1)
train_featured.head()
```

```
[58]:
```

	total_cases	year	weekofyear	week_start_date	month	fall	spring	\
0	4	1990.0	18.0	1990-04-30	4	0	1	
1	5	1990.0	19.0	1990-05-07	5	0	1	
2	4	1990.0	20.0	1990-05-14	5	0	1	
3	3	1990.0	21.0	1990-05-21	5	0	1	
4	6	1990.0	22.0	1990-05-28	5	0	1	

	summer	winter	station_avg_temp_c	...	reanalysis_precip_amt_kg_per_m2	\
0	0	0	25.442857	...	32.00	
1	0	0	26.714286	...	17.94	
2	0	0	26.714286	...	26.10	
3	0	0	27.471429	...	13.90	
4	0	0	28.942857	...	12.20	

	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	ndvi_average	ndvi_average_cat	\
--	---------	---------	---------	---------	--------------	------------------	---

0	0.122600	0.103725	0.198483	0.177617	0.150606	soily
1	0.169900	0.142175	0.162357	0.155486	0.157479	soily
2	0.032250	0.172967	0.157200	0.170843	0.133315	soily
3	0.128633	0.245067	0.227557	0.235886	0.209286	grassy
4	0.196200	0.262200	0.251200	0.247340	0.239235	grassy

	grassy	soily	watery
0	0	1	0
1	0	1	0
2	0	1	0
3	1	0	0
4	1	0	0

[5 rows x 24 columns]

```
[59]: # Let's see if correlations improved:
ndvi_data =
↳ train_featured[['total_cases', 'ndvi_average', 'grassy', 'soily', 'watery']]
ndvi_data.corr()['total_cases']
```

```
[59]: total_cases      1.000000
ndvi_average      0.052466
grassy           -0.043124
soily             0.102880
watery           -0.088839
Name: total_cases, dtype: float64
```

6.3 Create new shifted variables with rolled means:

- Research seems to indicate that past **sustained** heat, precipitation or humidity impacts dengue cases more profoundly than the climate situation right at the time of cases.
- I will be creating a series of rolled means for each of the chosen variable.
 - First shift the variables by 2 weeks to account for the growth of mosquito and the incubation period of the virus until testing positive.
 - Then create rolled means with a range of lags to find the variable with the highest correlation.

```
[60]: train_shifted = train_featured.copy()
```

```
[61]: # The variables to shift and roll:
shifted_varbls = ['station_avg_temp_c', 'station_max_temp_c',
↳ 'station_min_temp_c', 'reanalysis_tdtr_k',
                  'reanalysis_specific_humidity_g_per_kg',
↳ 'reanalysis_precip_amt_kg_per_m2',
                  'grassy', 'soily', 'watery']
```

```
[62]: # shift the variables two weeks ahead so that total cases would correspond with
      ↪ climate variables from two weeks ago:
      for var in shifted_varbls:
          train_shifted[f"{var}_shift"] = train_shifted[var].shift(2)
      # drop the two rows with NA
      train_shifted.dropna(axis=0, inplace=True)
      # reset the index
      train_shifted = train_shifted.reset_index(drop=True)
```

```
[63]: # Making sure it shifted corretly
      train_shifted[['year', 'weekofyear', 'week_start_date', 'station_avg_temp_c',
      ↪ 'station_avg_temp_c_shift', 'grassy', 'grassy_shift']].head()
```

```
[63]:
```

	year	weekofyear	week_start_date	station_avg_temp_c \
0	1990.0	20.0	1990-05-14	26.714286
1	1990.0	21.0	1990-05-21	27.471429
2	1990.0	22.0	1990-05-28	28.942857
3	1990.0	23.0	1990-06-04	28.114286
4	1990.0	24.0	1990-06-11	27.414286

	station_avg_temp_c_shift	grassy	grassy_shift
0	25.442857	0	0.0
1	26.714286	1	0.0
2	26.714286	1	0.0
3	27.471429	1	1.0
4	28.942857	0	1.0

```
[64]: # create another copy to get the rolled means
      train_rolled = train_shifted.copy()
```

```
[65]: # We will check the correlations between these variables and its shifted
      ↪ versions
      varbls_to_see_lags = ['total_cases',
                           'reanalysis_precip_amt_kg_per_m2',
                           'reanalysis_specific_humidity_g_per_kg',
                           'reanalysis_tdtr_k',
                           'station_avg_temp_c',
                           'station_max_temp_c',
                           'station_min_temp_c',
                           'grassy', 'soily', 'watery',
                           'reanalysis_precip_amt_kg_per_m2_shift',
                           'reanalysis_specific_humidity_g_per_kg_shift',
                           'reanalysis_tdtr_k_shift',
                           'station_avg_temp_c_shift',
                           'station_max_temp_c_shift',
                           'station_min_temp_c_shift',
                           'grassy_shift', 'soily_shift', 'watery_shift']
```

```
[66]: train_rolled = train_rolled[varbls_to_see_lags]
```

```
[67]: rolled_varbls = ['reanalysis_precip_amt_kg_per_m2_shift',  
                      'reanalysis_specific_humidity_g_per_kg_shift',  
                      'reanalysis_tdtr_k_shift',  
                      'station_avg_temp_c_shift',  
                      'station_max_temp_c_shift',  
                      'station_min_temp_c_shift',  
                      'grassy_shift', 'soily_shift', 'watery_shift']
```

```
[68]: # Create cumulative means for lags of 2 through 24:  
# The minimum lag is 2 weeks, maximum lag is 18 weeks (about 3-4 months):  
  
window = np.linspace(2,20,10).astype(int)  
min_periods = np.linspace(1,10,10).astype(int)  
  
for var in rolled_varbls:  
    for num,min in zip(window,min_periods):  
        train_rolled[f"{var}_{num}"] = train_rolled[var].rolling(window = num,  
↪min_periods = min).mean()
```

```
[69]: # Create 4 separate datasets for temp, humid, prec, ndvi variables with shifted,  
↪and rolled versions:  
temp_cols = [col for col in train_rolled.columns if 'temp' in col or 'tdtr' in_  
↪col]  
hum_cols = [col for col in train_rolled.columns if 'hum' in col]  
prec_cols = [col for col in train_rolled.columns if 'prec' in col]  
ndvi_cols = [col for col in train_rolled.columns if 'grassy' in col or 'soily'_  
↪in col or 'watery' in col]
```

```
[70]: # Add total_cases  
temp = train_rolled[temp_cols]  
temp['total_cases'] = train_rolled['total_cases']  
  
hum = train_rolled[hum_cols]  
hum['total_cases'] = train_rolled['total_cases']  
  
prec = train_rolled[prec_cols]  
prec['total_cases'] = train_rolled['total_cases']  
  
ndvi = train_rolled[ndvi_cols]  
ndvi['total_cases'] = train_rolled['total_cases']
```

<ipython-input-70-98b4a2a7251a>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
temp['total_cases'] = train_rolled['total_cases']  
<ipython-input-70-98b4a2a7251a>:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
hum['total_cases'] = train_rolled['total_cases']  
<ipython-input-70-98b4a2a7251a>:9: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
prec['total_cases'] = train_rolled['total_cases']  
<ipython-input-70-98b4a2a7251a>:12: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
ndvi['total_cases'] = train_rolled['total_cases']
```

```
[71]: fig, ax = plt.subplots(figsize=(12,12))  
temp.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).  
      plot.barh(ax=ax);
```



```
[72]: # create a function to see how lagged versions of a variable correlates with
      ↪ total cases.
def lag_graph(df, var):

    columns = [f"{var}", f"{var}_shift",
               f"{var}_shift_2", f"{var}_shift_4",
               f"{var}_shift_6", f"{var}_shift_8",
               f"{var}_shift_10", f"{var}_shift_12",
               f"{var}_shift_14", f"{var}_shift_16",
               f"{var}_shift_18", f"{var}_shift_20"]

    labels = ["original", "shift_2", "shift_2 + roll_2", "shift_2 + roll_4",
              ↪ "shift_2 + roll_6",
               "shift_2 + roll_8", "shift_2 + roll_10", "shift_2 + roll_12",
              ↪ "shift_2 + roll_14",
               "shift_2 + roll_16", "shift_2 + roll_18", "shift_2 + roll_20"]

    ncols = 2
```

```

nrows = 6
with plt.style.context('seaborn-talk'):

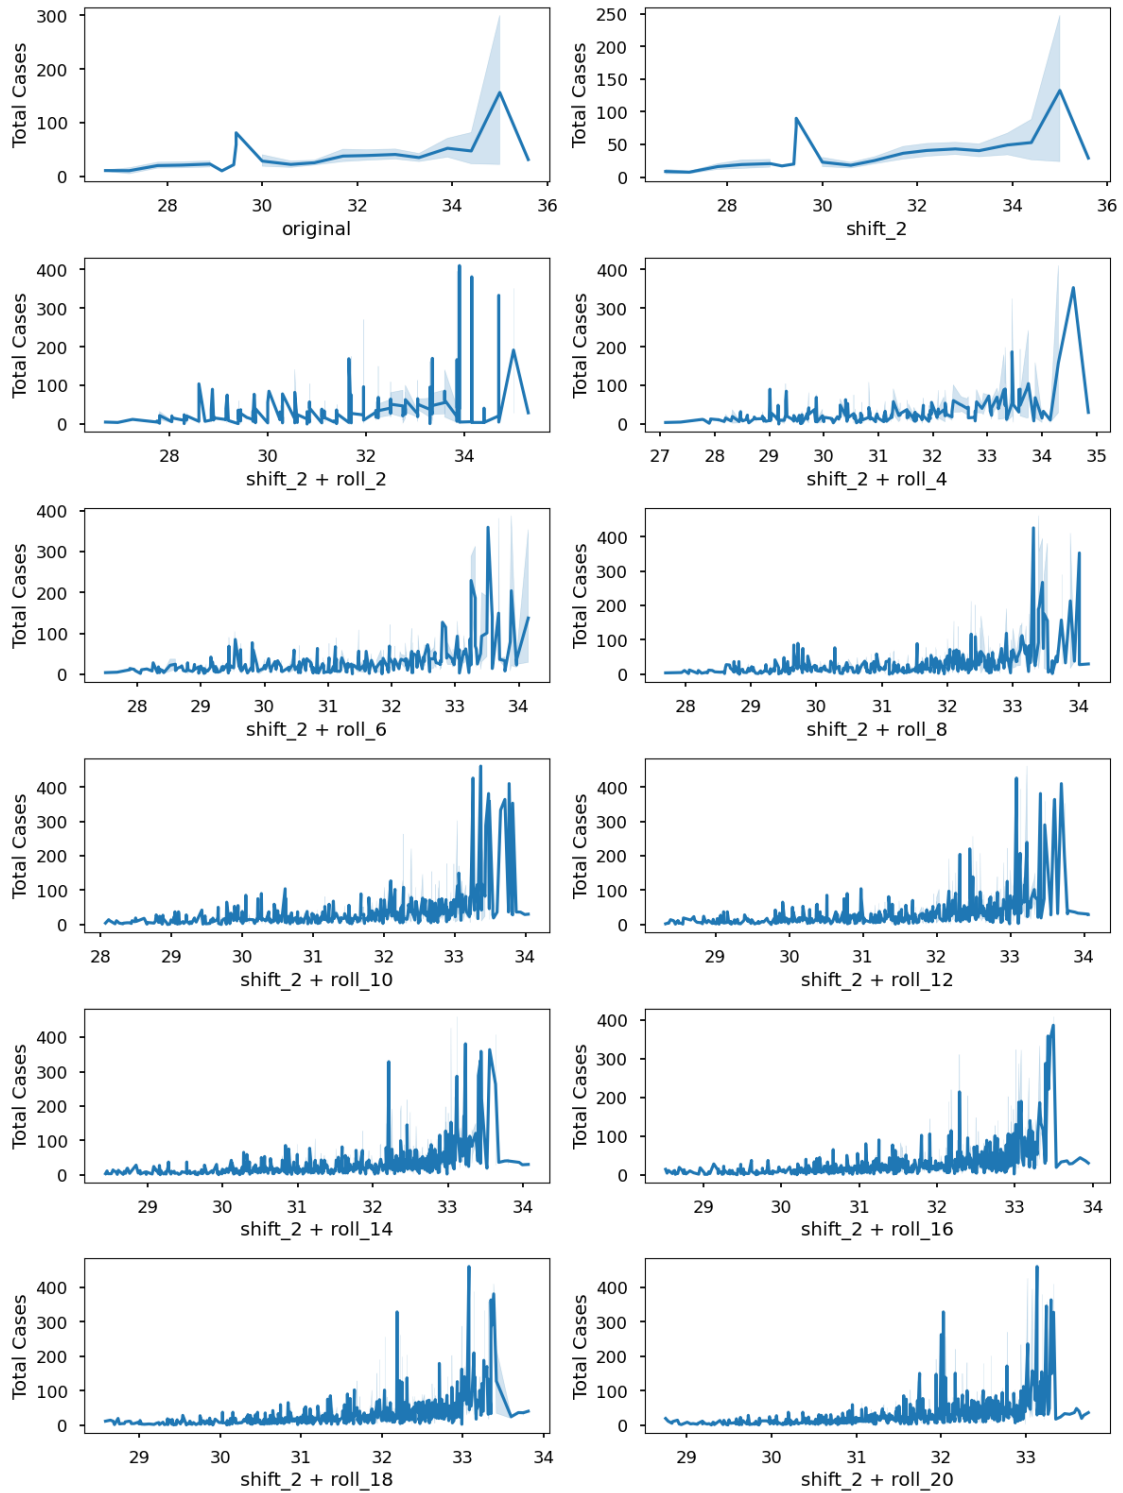
    fig, ax_list = plt.subplots(nrows = nrows, ncols = ncols,
↪figsize=(12,16))

    j=0
    for i in range(nrows):
        for u in range(ncols):
            sns.lineplot(data = df, x=columns[j], y="total_cases", ax =
↪ax_list[i,u]) # need to use index for column because otherwise it does not
↪iterate.
                ax_list[i,u].set_xlabel(labels[j])
                ax_list[i,u].set_ylabel("Total Cases")
                j = j+1

    fig.tight_layout();


```

```
[73]: lag_graph(train_rolled, "station_max_temp_c")
```




```
[74]: fig, ax = plt.subplots(figsize=(10,6))
```

```
hum.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).plot.  
↪barh(ax=ax);
```




```
[75]: lag_graph(train_rolled, "reanalysis_specific_humidity_g_per_kg")
```




```
[76]: fig, ax = plt.subplots(figsize=(10,6))
```

```
prec.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).  
    ↪plot.barh(ax=ax);
```



```
[77]: fig, ax = plt.subplots(figsize=(10,12))  
ndvi.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).  
    ↪plot.barh(ax=ax);
```



7 Based on above graphs I will be taking using variables:

- station_max_temp_c_shift_18,
- station_min_temp_c_shift_18,
- station_avg_temp_c_shift_18,
- reanalysis_tdtr_k_shift_8,
- reanalysis_specific_humidity_g_per_kg_shift_12,
- reanalysis_precip_amt_kg_per_m2_shift_8,
- grassy_shift_20,
- soily_shift_20,
- watery_shift_20


```
[78]: rolled_varblds_to_use = ['station_max_temp_c_shift_18',
                              'station_min_temp_c_shift_18',
                              'station_avg_temp_c_shift_18',
                              'reanalysis_tdtr_k_shift_8',
                              'reanalysis_specific_humidity_g_per_kg_shift_12',
                              'reanalysis_precip_amt_kg_per_m2_shift_8',
                              'grassy_shift_20',
                              'soily_shift_20', 'watery_shift_20']
```

```
[79]: # Add the rolled variables to the dataset
train_final = train_shifted.join(train_rolled[rolled_varblds_to_use])
train_final.head(20)
```

```
[79]:
```

	total_cases	year	weekofyear	week_start_date	month	fall	spring	\
0	4	1990.0	20.0	1990-05-14	5	0	1	
1	3	1990.0	21.0	1990-05-21	5	0	1	
2	6	1990.0	22.0	1990-05-28	5	0	1	
3	2	1990.0	23.0	1990-06-04	6	0	0	
4	4	1990.0	24.0	1990-06-11	6	0	0	
5	5	1990.0	25.0	1990-06-18	6	0	0	
6	10	1990.0	26.0	1990-06-25	6	0	0	
7	6	1990.0	27.0	1990-07-02	7	0	0	
8	8	1990.0	28.0	1990-07-09	7	0	0	
9	2	1990.0	29.0	1990-07-16	7	0	0	
10	6	1990.0	30.0	1990-07-23	7	0	0	
11	17	1990.0	31.0	1990-07-30	7	0	0	
12	23	1990.0	32.0	1990-08-06	8	0	0	
13	13	1990.0	33.0	1990-08-13	8	0	0	
14	21	1990.0	34.0	1990-08-20	8	0	0	
15	28	1990.0	35.0	1990-08-27	8	0	0	
16	24	1990.0	36.0	1990-09-03	9	1	0	
17	20	1990.0	37.0	1990-09-10	9	1	0	
18	40	1990.0	38.0	1990-09-17	9	1	0	
19	27	1990.0	39.0	1990-09-24	9	1	0	

	summer	winter	station_avg_temp_c	...	watery_shift	\
0	0	0	26.714286	...	0.0	
1	0	0	27.471429	...	0.0	
2	0	0	28.942857	...	0.0	
3	1	0	28.114286	...	0.0	
4	1	0	27.414286	...	0.0	
5	1	0	28.371429	...	0.0	
6	1	0	28.328571	...	0.0	
7	1	0	28.328571	...	0.0	
8	1	0	27.557143	...	0.0	
9	1	0	28.128571	...	0.0	
10	1	0	28.114286	...	0.0	

11	1	0	28.242857	...	0.0
12	1	0	28.200000	...	0.0
13	1	0	28.042857	...	0.0
14	1	0	28.342857	...	0.0
15	1	0	28.657143	...	0.0
16	0	0	28.328571	...	0.0
17	0	0	28.685714	...	0.0
18	0	0	28.242857	...	0.0
19	0	0	28.342857	...	0.0

	station_max_temp_c_shift_18	station_min_temp_c_shift_18	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	
5	NaN	NaN	
6	NaN	NaN	
7	NaN	NaN	
8	32.888889	22.777778	
9	32.990000	22.940000	
10	32.872727	22.827273	
11	32.866667	22.916667	
12	32.776923	22.907692	
13	32.892857	22.900000	
14	32.920000	22.926667	
15	32.912500	22.918750	
16	32.935294	22.941176	
17	32.894444	23.022222	
18	33.083333	23.205556	
19	33.205556	23.327778	

	station_avg_temp_c_shift_18	reanalysis_tdtr_k_shift_8	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	2.432143	
4	NaN	2.548571	
5	NaN	2.473810	
6	NaN	2.412245	
7	NaN	2.307143	
8	27.501587	2.214286	
9	27.584286	2.169643	
10	27.581818	2.151786	
11	27.627381	2.150000	
12	27.664835	2.032143	
13	27.706122	2.092857	

14	27.739048	2.128571
15	27.758036	2.164286
16	27.792437	2.239286
17	27.840476	2.225000
18	28.000794	2.389286
19	28.110317	2.450000

	reanalysis_specific_humidity_g_per_kg_shift_12 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
5	16.221667
6	16.366327
7	16.567679
8	16.703492
9	16.840286
10	16.892857
11	16.963214
12	17.240595
13	17.425714
14	17.517500
15	17.671071
16	17.751429
17	17.860000
18	17.887143
19	17.953810

	reanalysis_precip_amt_kg_per_m2_shift_8	grassy_shift_20	soily_shift_20 \
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	22.485000	NaN	NaN
4	20.428000	NaN	NaN
5	21.438333	NaN	NaN
6	23.890000	NaN	NaN
7	24.653750	NaN	NaN
8	25.342500	NaN	NaN
9	26.650000	0.300000	0.700000
10	28.852500	0.272727	0.727273
11	32.227500	0.333333	0.666667
12	36.018750	0.384615	0.615385
13	37.032500	0.428571	0.571429
14	34.707500	0.400000	0.600000
15	43.695000	0.375000	0.625000
16	42.243750	0.411765	0.588235

17	43.668750	0.444444	0.555556
18	42.003750	0.473684	0.526316
19	39.913750	0.500000	0.500000

	watery_shift_20
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
5	NaN
6	NaN
7	NaN
8	NaN
9	0.0
10	0.0
11	0.0
12	0.0
13	0.0
14	0.0
15	0.0
16	0.0
17	0.0
18	0.0
19	0.0

[20 rows x 42 columns]

```
[80]: # We are losing the first 21 rows
train_final = train_final.dropna().reset_index(drop=True)
```

```
[81]: train_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 925 entries, 0 to 924
Data columns (total 42 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   total_cases                          925 non-null    int64
1   year                                925 non-null    float64
2   weekofyear                          925 non-null    float64
3   week_start_date                     925 non-null    datetime64[ns]
4   month                               925 non-null    int64
5   fall                                925 non-null    uint8
6   spring                              925 non-null    uint8
7   summer                              925 non-null    uint8
```

```

8    winter                                925 non-null    uint8
9    station_avg_temp_c                   925 non-null    float64
10   station_max_temp_c                   925 non-null    float64
11   station_min_temp_c                   925 non-null    float64
12   reanalysis_tdtr_k                    925 non-null    float64
13   reanalysis_specific_humidity_g_per_kg 925 non-null    float64
14   reanalysis_precip_amt_kg_per_m2       925 non-null    float64
15   ndvi_ne                              925 non-null    float64
16   ndvi_nw                              925 non-null    float64
17   ndvi_se                              925 non-null    float64
18   ndvi_sw                              925 non-null    float64
19   ndvi_average                         925 non-null    float64
20   ndvi_average_cat                     925 non-null    object
21   grassy                              925 non-null    uint8
22   soily                                925 non-null    uint8
23   watery                                925 non-null    uint8
24   station_avg_temp_c_shift              925 non-null    float64
25   station_max_temp_c_shift              925 non-null    float64
26   station_min_temp_c_shift              925 non-null    float64
27   reanalysis_tdtr_k_shift               925 non-null    float64
28   reanalysis_specific_humidity_g_per_kg_shift 925 non-null    float64
29   reanalysis_precip_amt_kg_per_m2_shift 925 non-null    float64
30   grassy_shift                         925 non-null    float64
31   soily_shift                          925 non-null    float64
32   watery_shift                         925 non-null    float64
33   station_max_temp_c_shift_18           925 non-null    float64
34   station_min_temp_c_shift_18           925 non-null    float64
35   station_avg_temp_c_shift_18           925 non-null    float64
36   reanalysis_tdtr_k_shift_8             925 non-null    float64
37   reanalysis_specific_humidity_g_per_kg_shift_12 925 non-null    float64
38   reanalysis_precip_amt_kg_per_m2_shift_8 925 non-null    float64
39   grassy_shift_20                      925 non-null    float64
40   soily_shift_20                       925 non-null    float64
41   watery_shift_20                      925 non-null    float64

```

dtypes: datetime64[ns](1), float64(31), int64(2), object(1), uint8(7)

memory usage: 259.4+ KB

7.1 Repeat all steps for the final test set:

- First add the last 23 (21 +2 for shifting) rows of the train_full to test_full to not to lose data while transforming

```
[82]: # Add week_start_date to both datasets
test_features_full['week_start_date'] = test_features['week_start_date']
train_features_full['week_start_date'] = train_features['week_start_date']
```

```
[83]: test_features_long = pd.concat([train_features_full.tail(21),
    ↪ test_features_full], ignore_index=True)
```

```
test_features_long
```

```
[83]:
```

	year	weekofyear	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	\
0	2007.0	49.0	-0.03976	-0.042350	0.095600	0.089000	
1	2007.0	50.0	-0.13305	-0.045550	0.151440	0.143171	
2	2007.0	51.0	0.02945	-0.039000	0.173417	0.150171	
3	2007.0	52.0	0.01480	0.016300	0.207267	0.144578	
4	2008.0	1.0	0.00060	-0.309600	0.239814	0.195557	
..	
276	2013.0	13.0	-0.08740	-0.016183	0.156343	0.105186	
277	2013.0	14.0	-0.20325	-0.077833	0.204171	0.178914	
278	2013.0	15.0	-0.11760	-0.008200	0.192700	0.170429	
279	2013.0	16.0	0.08275	0.031200	0.135014	0.074857	
280	2013.0	17.0	-0.08730	-0.048667	0.129814	0.117671	

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	\
0	17.85	299.020000	299.021429	
1	31.30	298.900000	298.971429	
2	62.11	298.668571	298.757143	
3	0.00	298.602857	298.750000	
4	0.00	298.038571	298.121429	
..	
276	30.34	298.670000	298.885714	
277	6.55	298.035714	298.157143	
278	0.00	299.057143	299.328571	
279	0.00	298.912857	299.064286	
280	45.47	298.067143	298.042857	

	reanalysis_dew_point_temp_k	...	reanalysis_relative_humidity_percent	\
0	294.288571	...	75.368571	
1	294.774286	...	78.015714	
2	294.977143	...	80.178571	
3	293.928571	...	75.448571	
4	293.514286	...	76.148571	
..	
276	294.675714	...	78.780000	
277	294.628571	...	81.650000	
278	294.948571	...	78.285714	
279	294.678571	...	77.674286	
280	294.132857	...	79.045714	

	reanalysis_sat_precip_amt_mm	reanalysis_specific_humidity_g_per_kg	\
0	17.85	15.675714	
1	31.30	16.130000	
2	62.11	16.344286	
3	0.00	15.318571	
4	0.00	14.911429	

..
276	30.34	15.985714
277	6.55	15.881429
278	0.00	16.212857
279	0.00	15.965714
280	45.47	15.451429

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c	\
0	2.100000	25.842857	5.400000	
1	2.485714	25.771429	5.085714	
2	2.371429	25.071429	4.914286	
3	2.985714	25.085714	6.242857	
4	1.842857	25.400000	5.300000	
..	
276	3.314286	27.542857	7.942857	
277	2.828571	26.642857	6.642857	
278	3.171429	27.914286	8.114286	
279	3.042857	27.728571	6.942857	
280	2.342857	26.442857	6.742857	

	station_max_temp_c	station_min_temp_c	station_precip_mm	\
0	29.4	22.8	34.5	
1	28.9	22.2	30.2	
2	28.9	21.7	108.2	
3	28.3	21.1	16.8	
4	29.4	22.2	55.5	
..	
276	33.9	22.8	3.5	
277	33.3	22.8	17.6	
278	32.8	23.3	9.4	
279	31.7	23.9	22.9	
280	31.1	21.7	47.5	

	week_start_date
0	12/3/07
1	12/10/07
2	12/17/07
3	12/24/07
4	1/1/08
..	...
276	3/26/13
277	4/2/13
278	4/9/13
279	4/16/13
280	4/23/13

[281 rows x 23 columns]

```
[84]: # create a new month variable:

test_featured = test_features_long.copy()
test_featured["week_start_date"] = pd.
    ↳to_datetime(test_featured["week_start_date"])
test_featured['month'] = test_featured['week_start_date'].dt.month

# create a new season variable:
test_featured['season'] = test_featured['month'].map(month_to_season)
season_features = pd.get_dummies(test_featured['season'])
test_featured = pd.concat([test_featured, season_features], axis = 1)

test_featured = test_featured[['year', 'weekofyear', 'week_start_date',
                                'month', 'fall', 'spring', 'summer', 'winter',
                                'station_avg_temp_c', 'station_max_temp_c',
                                'station_min_temp_c', 'reanalysis_tdtr_k',
                                'reanalysis_specific_humidity_g_per_kg',
                                'reanalysis_precip_amt_kg_per_m2',
                                'ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw']]

test_featured['ndvi_average'] = ␣
    ↳(test_featured['ndvi_ne']+test_featured['ndvi_nw']+
        ␣
    ↳test_featured['ndvi_se']+test_featured['ndvi_sw'])/4

test_featured["ndvi_average_cat"] = test_featured["ndvi_average"].
    ↳apply(get_ndvi_cat)

ndvi_features = pd.get_dummies(test_featured['ndvi_average_cat'])
test_featured = pd.concat([test_featured, ndvi_features], axis = 1)

# shift by 2 weeks
test_shifted = test_featured.copy()

for var in shifted_varbls:
    test_shifted[f"{var}_shift"] = test_shifted[var].shift(2)

test_shifted.dropna(axis=0, inplace=True)
# test_shifted = test_shifted.reset_index(drop=True)

# get rolled means
test_rolled = test_shifted.copy()

varbls_to_see_lags = ['reanalysis_precip_amt_kg_per_m2',
                      'reanalysis_specific_humidity_g_per_kg',
                      'reanalysis_tdtr_k',
```



```

        'station_avg_temp_c',
        'station_max_temp_c',
        'station_min_temp_c',
        'grassy', 'soily', 'watery',
        'reanalysis_precip_amt_kg_per_m2_shift',
        'reanalysis_specific_humidity_g_per_kg_shift',
        'reanalysis_tdtr_k_shift',
        'station_avg_temp_c_shift',
        'station_max_temp_c_shift',
        'station_min_temp_c_shift',
        'grassy_shift', 'soily_shift', 'watery_shift']

test_rolled = test_rolled[varbls_to_see_lags]

for var in rolled_varbls:
    for num in window:
        test_rolled[f"{var}_{num}"] = test_rolled[var].rolling(num).mean()

rolled_varbls_to_use = ['station_avg_temp_c_shift',
                        'station_max_temp_c_shift',
                        'station_min_temp_c_shift',
                        'reanalysis_precip_amt_kg_per_m2_shift',
                        'reanalysis_specific_humidity_g_per_kg_shift',
                        'reanalysis_tdtr_k_shift',
                        'grassy_shift', 'soily_shift', 'watery_shift',
                        'station_max_temp_c_shift_18',
                        'station_min_temp_c_shift_18',
                        'station_avg_temp_c_shift_18',
                        'reanalysis_tdtr_k_shift_8',
                        'reanalysis_specific_humidity_g_per_kg_shift_12',
                        'reanalysis_precip_amt_kg_per_m2_shift_8',
                        'grassy_shift_20',
                        'soily_shift_20', 'watery_shift_20']

test_final = test_featured.join(test_rolled[rolled_varbls_to_use])
test_final = test_final.dropna().reset_index(drop=True)
test_final

```

```

[84]:
   year  weekofyear week_start_date  month  fall  spring  summer  winter  \
0  2008.0         18.0    2008-04-29      4     0       1       0       0
1  2008.0         19.0    2008-05-06      5     0       1       0       0
2  2008.0         20.0    2008-05-13      5     0       1       0       0
3  2008.0         21.0    2008-05-20      5     0       1       0       0
4  2008.0         22.0    2008-05-27      5     0       1       0       0
..    ...         ...             ...    ...    ...    ...    ...

```

255	2013.0	13.0	2013-03-26	3	0	1	0	0
256	2013.0	14.0	2013-04-02	4	0	1	0	0
257	2013.0	15.0	2013-04-09	4	0	1	0	0
258	2013.0	16.0	2013-04-16	4	0	1	0	0
259	2013.0	17.0	2013-04-23	4	0	1	0	0

	station_avg_temp_c	station_max_temp_c	...	watery_shift	\
0	26.528571	33.3	...	0.0	
1	26.071429	30.0	...	1.0	
2	27.928571	32.8	...	1.0	
3	28.057143	33.3	...	1.0	
4	27.614286	33.3	...	0.0	
..	
255	27.542857	33.9	...	0.0	
256	26.642857	33.3	...	1.0	
257	27.914286	32.8	...	1.0	
258	27.728571	31.7	...	1.0	
259	26.442857	31.1	...	1.0	

	station_max_temp_c_shift_18	station_min_temp_c_shift_18	\
0	28.816667	21.300000	
1	28.972222	21.388889	
2	29.250000	21.422222	
3	29.283333	21.422222	
4	29.533333	21.516667	
..	
255	30.766667	22.377778	
256	30.550000	22.194444	
257	30.611111	22.105556	
258	30.733333	22.016667	
259	30.827778	21.983333	

	station_avg_temp_c_shift_18	reanalysis_tdtr_k_shift_8	\
0	24.997619	2.678571	
1	25.107143	2.887500	
2	25.187302	2.951786	
3	25.224603	3.012500	
4	25.388889	3.082143	
..	
255	26.509524	2.723214	
256	26.352381	2.775000	
257	26.298413	2.891071	
258	26.227778	2.910714	
259	26.244444	2.971429	

	reanalysis_specific_humidity_g_per_kg_shift_12	\
0	14.065833	

1	14.171071
2	14.395476
3	14.514167
4	14.692143
..	...
255	14.662024
256	14.439048
257	14.439524
258	14.490000
259	14.648452

	reanalysis_precip_amt_kg_per_m2_shift_8	grassy_shift_20	soily_shift_20	\
0	9.28500	0.0	0.50	
1	8.74875	0.0	0.50	
2	11.10750	0.0	0.50	
3	12.66375	0.0	0.45	
4	13.06625	0.0	0.45	
..	
255	7.32500	0.0	0.60	
256	6.88750	0.0	0.55	
257	6.24375	0.0	0.55	
258	12.39375	0.0	0.50	
259	11.53125	0.0	0.45	

	watery_shift_20
0	0.50
1	0.50
2	0.50
3	0.55
4	0.55
..	...
255	0.40
256	0.45
257	0.45
258	0.50
259	0.55

[260 rows x 41 columns]

```
[85]: # Making sure test_final has the same length with older version
len(test_features_full) == len(test_final)
```

[85]: True

```
[86]: # Making sure test_final follows train_final corrretly in terms of date
print(train_final['week_start_date'])
print('-----')
```

```
print(test_final['week_start_date'])
```

```
0    1990-07-16
1    1990-07-23
2    1990-07-30
3    1990-08-06
4    1990-08-13
...
920   2008-03-25
921   2008-04-01
922   2008-04-08
923   2008-04-15
924   2008-04-22
Name: week_start_date, Length: 925, dtype: datetime64[ns]
-----
0    2008-04-29
1    2008-05-06
2    2008-05-13
3    2008-05-20
4    2008-05-27
...
255   2013-03-26
256   2013-04-02
257   2013-04-09
258   2013-04-16
259   2013-04-23
Name: week_start_date, Length: 260, dtype: datetime64[ns]
```

```
[87]: train_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 925 entries, 0 to 924
Data columns (total 42 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   total_cases                          925 non-null    int64
1   year                                925 non-null    float64
2   weekofyear                          925 non-null    float64
3   week_start_date                     925 non-null    datetime64[ns]
4   month                               925 non-null    int64
5   fall                                925 non-null    uint8
6   spring                              925 non-null    uint8
7   summer                              925 non-null    uint8
8   winter                              925 non-null    uint8
9   station_avg_temp_c                  925 non-null    float64
10  station_max_temp_c                  925 non-null    float64
11  station_min_temp_c                  925 non-null    float64
```

12	reanalysis_tdtr_k	925 non-null	float64
13	reanalysis_specific_humidity_g_per_kg	925 non-null	float64
14	reanalysis_precip_amt_kg_per_m2	925 non-null	float64
15	ndvi_ne	925 non-null	float64
16	ndvi_nw	925 non-null	float64
17	ndvi_se	925 non-null	float64
18	ndvi_sw	925 non-null	float64
19	ndvi_average	925 non-null	float64
20	ndvi_average_cat	925 non-null	object
21	grassy	925 non-null	uint8
22	soily	925 non-null	uint8
23	watery	925 non-null	uint8
24	station_avg_temp_c_shift	925 non-null	float64
25	station_max_temp_c_shift	925 non-null	float64
26	station_min_temp_c_shift	925 non-null	float64
27	reanalysis_tdtr_k_shift	925 non-null	float64
28	reanalysis_specific_humidity_g_per_kg_shift	925 non-null	float64
29	reanalysis_precip_amt_kg_per_m2_shift	925 non-null	float64
30	grassy_shift	925 non-null	float64
31	soily_shift	925 non-null	float64
32	watery_shift	925 non-null	float64
33	station_max_temp_c_shift_18	925 non-null	float64
34	station_min_temp_c_shift_18	925 non-null	float64
35	station_avg_temp_c_shift_18	925 non-null	float64
36	reanalysis_tdtr_k_shift_8	925 non-null	float64
37	reanalysis_specific_humidity_g_per_kg_shift_12	925 non-null	float64
38	reanalysis_precip_amt_kg_per_m2_shift_8	925 non-null	float64
39	grassy_shift_20	925 non-null	float64
40	soily_shift_20	925 non-null	float64
41	watery_shift_20	925 non-null	float64

dtypes: datetime64[ns](1), float64(31), int64(2), object(1), uint8(7)

memory usage: 259.4+ KB

```
[88]: test_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 260 entries, 0 to 259
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   year                                260 non-null    float64
1   weekofyear                          260 non-null    float64
2   week_start_date                     260 non-null
datetime64[ns]
3   month                              260 non-null    int64
4   fall                               260 non-null    uint8
5   spring                             260 non-null    uint8
6   summer                             260 non-null    uint8
```

7	winter	260	non-null	uint8
8	station_avg_temp_c	260	non-null	float64
9	station_max_temp_c	260	non-null	float64
10	station_min_temp_c	260	non-null	float64
11	reanalysis_tdtr_k	260	non-null	float64
12	reanalysis_specific_humidity_g_per_kg	260	non-null	float64
13	reanalysis_precip_amt_kg_per_m2	260	non-null	float64
14	ndvi_ne	260	non-null	float64
15	ndvi_nw	260	non-null	float64
16	ndvi_se	260	non-null	float64
17	ndvi_sw	260	non-null	float64
18	ndvi_average	260	non-null	float64
19	ndvi_average_cat	260	non-null	object
20	grassy	260	non-null	uint8
21	soily	260	non-null	uint8
22	watery	260	non-null	uint8
23	station_avg_temp_c_shift	260	non-null	float64
24	station_max_temp_c_shift	260	non-null	float64
25	station_min_temp_c_shift	260	non-null	float64
26	reanalysis_precip_amt_kg_per_m2_shift	260	non-null	float64
27	reanalysis_specific_humidity_g_per_kg_shift	260	non-null	float64
28	reanalysis_tdtr_k_shift	260	non-null	float64
29	grassy_shift	260	non-null	float64
30	soily_shift	260	non-null	float64
31	watery_shift	260	non-null	float64
32	station_max_temp_c_shift_18	260	non-null	float64
33	station_min_temp_c_shift_18	260	non-null	float64
34	station_avg_temp_c_shift_18	260	non-null	float64
35	reanalysis_tdtr_k_shift_8	260	non-null	float64
36	reanalysis_specific_humidity_g_per_kg_shift_12	260	non-null	float64
37	reanalysis_precip_amt_kg_per_m2_shift_8	260	non-null	float64
38	grassy_shift_20	260	non-null	float64
39	soily_shift_20	260	non-null	float64
40	watery_shift_20	260	non-null	float64

dtypes: datetime64[ns](1), float64(31), int64(1), object(1), uint8(7)

memory usage: 71.0+ KB

```
[89]: # Export the final datasets as csv to be used for modeling
train_final.to_csv("train_final.csv")
test_final.to_csv("test_final.csv")
```

7.1.1 Export as PDF:

```
[90]: # Packages required for using nbconvert PDF
# ! apt-get install texlive texlive-xetex texlive-latex-extra pandoc
# ! pip install pypandoc
# ! pip install nbconvert
```

```
[91]: # First you need to download a copy of the ipynb notebook and upload it back to ↪ the drive, it is placed under /content/  
# ! jupyter nbconvert --to pdf /content/notebook_preprocessing.ipynb
```