

notebook_preprocessing

April 6, 2023

1 Predicting Dengue Cases

- Student name: *Aysu Erdemir*
- Student pace: *Flex*
- Scheduled project review date/time: *March, 2023*
- Instructor name: *Abhineet Kulkarni*

1.1 Overview:

Dengue fever is a mosquito-borne disease that occurs in tropical and sub-tropical parts of the world. In mild cases, symptoms are similar to the flu: fever, rash, and muscle and joint pain. In severe cases, dengue fever can cause severe bleeding, low blood pressure, and even death.

Because it is carried by mosquitoes, the transmission dynamics of dengue are related to climate variables such as temperature and precipitation; however the relationship to climate is known to be complex. The way the disease spreads and causes endemics has significant public health implications worldwide.

- CDC is interested in predicting local epidemics of dengue fever so that they can take necessary precautions and efforts before the next spike. They want to know if we can predict the number of dengue fever cases reported each week in San Juan, Puerto Rico.
- My goal is to build several machine learning models to forecast the upcoming weekly dengue cases as accurately as possible.

1.2 Business and Data Understanding

- The data was obtained from [DrivenData](#). The data set included weekly dengue case counts along with environmental data collected by various U.S. Federal Government agencies—from the Centers for Disease Control and Prevention to the National Oceanic and Atmospheric Administration in the U.S. Department of Commerce.
- The full dataset included cases from year 1990 to 2008. The data from 2008-2013 included only features without case counts.
- In this project I will be focusing on data on **Puerto Rico** only. The relevant variables/features included in the dataset are:

Target Feature: * total_cases - Weekly total dengue cases.

Predictive Features:

Date Indicators:

- week_start_date - Date given in yyyy-mm-dd format.

NOAA's GHCN daily climate data weather station measurements:

- station_max_temp_c - Maximum temperature
- station_min_temp_c - Minimum temperature
- station_avg_temp_c - Average temperature
- station_precip_mm - Total precipitation
- station_diur_temp_rng_c - Diurnal temperature range

PERSIANN satellite precipitation measurements (0.25x0.25 degree scale):

- precipitation_amt_mm - Total precipitation

NOAA's NCEP Climate Forecast System Reanalysis measurements (0.5x0.5 degree scale):

- reanalysis_sat_precip_amt_mm - Total precipitation
- reanalysis_dew_point_temp_k - Mean dew point temperature
- reanalysis_air_temp_k - Mean air temperature
- reanalysis_relative_humidity_percent - Mean relative humidity
- reanalysis_specific_humidity_g_per_kg - Mean specific humidity
- reanalysis_precip_amt_kg_per_m2 - Total precipitation
- reanalysis_max_air_temp_k - Maximum air temperature
- reanalysis_min_air_temp_k - Minimum air temperature
- reanalysis_avg_temp_k - Average air temperature
- reanalysis_tdtr_k - Diurnal temperature range

Satellite vegetation -greenness - Normalized difference vegetation index (NDVI) - NOAA's CDR Normalized Difference Vegetation Index (0.5x0.5 degree scale) measurements:

- ndvi_se - Pixel southeast of city centroid
- ndvi_sw - Pixel southwest of city centroid
- ndvi_ne - Pixel northeast of city centroid
- ndvi_nw - Pixel northwest of city centroid

For example, when you have negative values, it's highly likely that it's water. On the other hand, if you have an NDVI value close to +1, there's a high possibility that it's dense green leaves. But when NDVI is close to zero, there are likely no green leaves and it could even be an urbanized area.

1.3 Preprocessing:

1.3.1 Null Replacement:

- Null values for the climate features - except the four ndvi fatures - were imputed with **interpolation** since the missing data points are scarce.
- Null values for the four ndvi fatures were imputed using **k-Nearest Neighbors - KNN** since there were bigger chunks of missing values.

1.3.2 Feature Engineering:

- Create **month** and **seasons**: Created new variables representing the month and seasons.
- Create **average_ndvi** and its **categorical** version: Created a new feature representing the average NDVI values using the four different locations. Then created a categorical version of **average_ndvi** to represent watery, soily, sparse_grassy areas.
- Create **shifts** and **rolled averages** for the main climate variables: Research seems to indicate that past sustained heat, precipitation or humidity impacts dengue cases more profoundly than the climate situation right at the time of cases.
 - **Shifted** the variables by 2 weeks to account for the mosquito to reach adulthood and the incubation period of the virus until someone tests positive.
 - Engineered **rolled - cumulative** means over a period of time ranging from 1 weeks to 20 weeks to see the variable with the highest correlation. The lag with the highest corralation was kept in the final dataset. The final lags ranged from 2 months to 4 months.

Some initial thoughts based on past research:

Precipitation: - Mosquitos thrive wet climates, the wetter the better! - A rise in **accumulated rainfall** was shown to result in an increase in the number and quality of breeding sites. - Elevated relative risk of dengue was observed when the weekly average rainfall was more than 150 mm at **lagged weeks 12 to 20**. - However, **above a certain rainfall level**, suitable mosquito breeding sites can be exposed to flooding, so the population is likely to decrease in such cases.

Humidity: - Humidity generates conditions that are favorable to adult mosquitoes increasing the life of the mosquito.

- Humidity range of **60% - 90%** is the optimum moisture for growth and development of the *Aedes aegypti* mosquito.

Temperature: - Mosquitoes more likely to transmit dengue virus in hot weather. Mosquitoes function best at 80 degrees F. **Higher temperatures (73-85°F)** results in more rapid viral growth and higher levels of virus. - The relative risk of dengue fever increases when the weekly average temperature is high **at lagged weeks 5 to 18**. - Under **fluctuating temperatures**, the mosquitoes show lower levels of virus in their salivary glands. - **Extreme heat waves** can negatively impact mosquito life as very high temperatures reduce adult lifespan and egg survival, resulting in reduced *Aedes* population and lower risk of dengue transmission.

NDVI index: - The relationship between dengue epidemic and greenness indexes is not clear or consistent. - Some studies indicate increased density of vegetation to provide suitable habitat for the immature mosquitoes. - Others indicated that low vegetation cover areas with increased dengue incidence rates. This inconsistency may be explained by regional differences.

```
[93]: # Import required packages
```

```

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import missingno
from sklearn.preprocessing import StandardScaler
from sklearn.impute import KNNImputer

import warnings
warnings.filterwarnings("ignore")
# check package versions when necessary:
# pd.__version__

```

2 Data Exploration:

```

[2]: # load the files
from google.colab import files
uploaded = files.upload()

```

<IPython.core.display.HTML object>

Saving dengue_labels_train.csv to dengue_labels_train (1).csv
 Saving dengue_features_train.csv to dengue_features_train (1).csv
 Saving dengue_features_test.csv to dengue_features_test (1).csv

```

[3]: # Read the Data
train_features = pd.read_csv("dengue_features_train.csv")
train_labels = pd.read_csv("dengue_labels_train.csv")
test_features = pd.read_csv("dengue_features_test.csv")

```

```

[4]: train_features.head()

```

```

[4]:   year  weekofyear  week_start_date  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  1990         18      4/30/90  0.122600  0.103725  0.198483  0.177617
1  1990         19      5/7/90  0.169900  0.142175  0.162357  0.155486
2  1990         20      5/14/90  0.032250  0.172967  0.157200  0.170843
3  1990         21      5/21/90  0.128633  0.245067  0.227557  0.235886
4  1990         22      5/28/90  0.196200  0.262200  0.251200  0.247340

      precipitation_amt_mm  reanalysis_air_temp_k  reanalysis_avg_temp_k  ...  \
0                12.42        297.572857        297.742857  ...
1                22.82        298.211429        298.442857  ...

```

2	34.54	298.781429	298.878571	...
3	15.36	298.987143	299.228571	...
4	7.52	299.518571	299.664286	...

	reanalysis_precip_amt_kg_per_m2	reanalysis_relative_humidity_percent	\
0	32.00	73.365714	
1	17.94	77.368571	
2	26.10	82.052857	
3	13.90	80.337143	
4	12.20	80.460000	

	reanalysis_sat_precip_amt_mm	reanalysis_specific_humidity_g_per_kg	\
0	12.42	14.012857	
1	22.82	15.372857	
2	34.54	16.848571	
3	15.36	16.672857	
4	7.52	17.210000	

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c	\
0	2.628571	25.442857	6.900000	
1	2.371429	26.714286	6.371429	
2	2.300000	26.714286	6.485714	
3	2.428571	27.471429	6.771429	
4	3.014286	28.942857	9.371429	

	station_max_temp_c	station_min_temp_c	station_precip_mm
0	29.4	20.0	16.0
1	31.7	22.2	8.6
2	32.2	22.8	41.4
3	33.3	23.3	4.0
4	35.0	23.9	5.8

[5 rows x 23 columns]

```
[5]: train_labels.head()
```

```
[5]:   year  weekofyear  total_cases
0  1990         18           4
1  1990         19           5
2  1990         20           4
3  1990         21           3
4  1990         22           6
```

```
[6]: test_features.head()
```

```
[6]:   year  weekofyear  week_start_date  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  2008         18        4/29/08   -0.0189 -0.018900  0.102729  0.091200
```

1	2008	19	5/6/08	-0.0180	-0.012400	0.082043	0.072314
2	2008	20	5/13/08	-0.0015	NaN	0.151083	0.091529
3	2008	21	5/20/08	NaN	-0.019867	0.124329	0.125686
4	2008	22	5/27/08	0.0568	0.039833	0.062267	0.075914

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	...	\
0	78.60	298.492857	298.550000	...	
1	12.56	298.475714	298.557143	...	
2	3.66	299.455714	299.357143	...	
3	0.00	299.690000	299.728571	...	
4	0.76	299.780000	299.671429	...	

	reanalysis_precip_amt_kg_per_m2	reanalysis_relative_humidity_percent	\
0	25.37	78.781429	
1	21.83	78.230000	
2	4.12	78.270000	
3	2.20	73.015714	
4	4.36	74.084286	

	reanalysis_sat_precip_amt_mm	reanalysis_specific_humidity_g_per_kg	\
0	78.60	15.918571	
1	12.56	15.791429	
2	3.66	16.674286	
3	0.00	15.775714	
4	0.76	16.137143	

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c	\
0	3.128571	26.528571	7.057143	
1	2.571429	26.071429	5.557143	
2	4.428571	27.928571	7.785714	
3	4.342857	28.057143	6.271429	
4	3.542857	27.614286	7.085714	

	station_max_temp_c	station_min_temp_c	station_precip_mm
0	33.3	21.7	75.2
1	30.0	22.2	34.3
2	32.8	22.8	3.0
3	33.3	24.4	0.3
4	33.3	23.3	84.1

[5 rows x 23 columns]

```
[7]: # Merge the features dataset with the labels dataset (total cases) to form
      ↪ train.
train = pd.merge(train_features, train_labels, on=[ "year", "weekofyear"])
train.head()
```

```

[7]:   year  weekofyear week_start_date  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  1990           18      4/30/90  0.122600  0.103725  0.198483  0.177617
1  1990           19      5/7/90  0.169900  0.142175  0.162357  0.155486
2  1990           20      5/14/90  0.032250  0.172967  0.157200  0.170843
3  1990           21      5/21/90  0.128633  0.245067  0.227557  0.235886
4  1990           22      5/28/90  0.196200  0.262200  0.251200  0.247340

      precipitation_amt_mm  reanalysis_air_temp_k  reanalysis_avg_temp_k  ...  \
0                12.42          297.572857          297.742857  ...
1                22.82          298.211429          298.442857  ...
2                34.54          298.781429          298.878571  ...
3                15.36          298.987143          299.228571  ...
4                 7.52          299.518571          299.664286  ...

      reanalysis_relative_humidity_percent  reanalysis_sat_precip_amt_mm  \
0                73.365714                12.42
1                77.368571                22.82
2                82.052857                34.54
3                80.337143                15.36
4                80.460000                7.52

      reanalysis_specific_humidity_g_per_kg  reanalysis_tdtr_k  \
0                14.012857                2.628571
1                15.372857                2.371429
2                16.848571                2.300000
3                16.672857                2.428571
4                17.210000                3.014286

      station_avg_temp_c  station_diur_temp_rng_c  station_max_temp_c  \
0                25.442857                6.900000                29.4
1                26.714286                6.371429                31.7
2                26.714286                6.485714                32.2
3                27.471429                6.771429                33.3
4                28.942857                9.371429                35.0

      station_min_temp_c  station_precip_mm  total_cases
0                20.0          16.0          4
1                22.2           8.6          5
2                22.8          41.4          4
3                23.3           4.0          3
4                23.9           5.8          6

```

[5 rows x 24 columns]

```

[8]: # check row and column numbers
print(train_features.shape)
print(train_labels.shape)

```

```
print(train.shape)
```

```
(936, 23)
```

```
(936, 3)
```

```
(936, 24)
```

```
[9]: train.info()  
# all variables are numerical except week_start_date which is an object
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 936 entries, 0 to 935  
Data columns (total 24 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   year                                936 non-null    int64  
1   weekofyear                          936 non-null    int64  
2   week_start_date                     936 non-null    object  
3   ndvi_ne                             745 non-null    float64  
4   ndvi_nw                             887 non-null    float64  
5   ndvi_se                             917 non-null    float64  
6   ndvi_sw                             917 non-null    float64  
7   precipitation_amt_mm                927 non-null    float64  
8   reanalysis_air_temp_k               930 non-null    float64  
9   reanalysis_avg_temp_k               930 non-null    float64  
10  reanalysis_dew_point_temp_k         930 non-null    float64  
11  reanalysis_max_air_temp_k           930 non-null    float64  
12  reanalysis_min_air_temp_k           930 non-null    float64  
13  reanalysis_precip_amt_kg_per_m2     930 non-null    float64  
14  reanalysis_relative_humidity_percent 930 non-null    float64  
15  reanalysis_sat_precip_amt_mm        927 non-null    float64  
16  reanalysis_specific_humidity_g_per_kg 930 non-null    float64  
17  reanalysis_tdtr_k                   930 non-null    float64  
18  station_avg_temp_c                  930 non-null    float64  
19  station_diur_temp_rng_c             930 non-null    float64  
20  station_max_temp_c                  930 non-null    float64  
21  station_min_temp_c                  930 non-null    float64  
22  station_precip_mm                   930 non-null    float64  
23  total_cases                          936 non-null    int64  
dtypes: float64(20), int64(3), object(1)  
memory usage: 182.8+ KB
```

```
[10]: train.total_cases.describe()  
# Total weekly cases range from 0 to 461
```

```
[10]: count    936.000000  
      mean     34.180556  
      std     51.381372
```



```

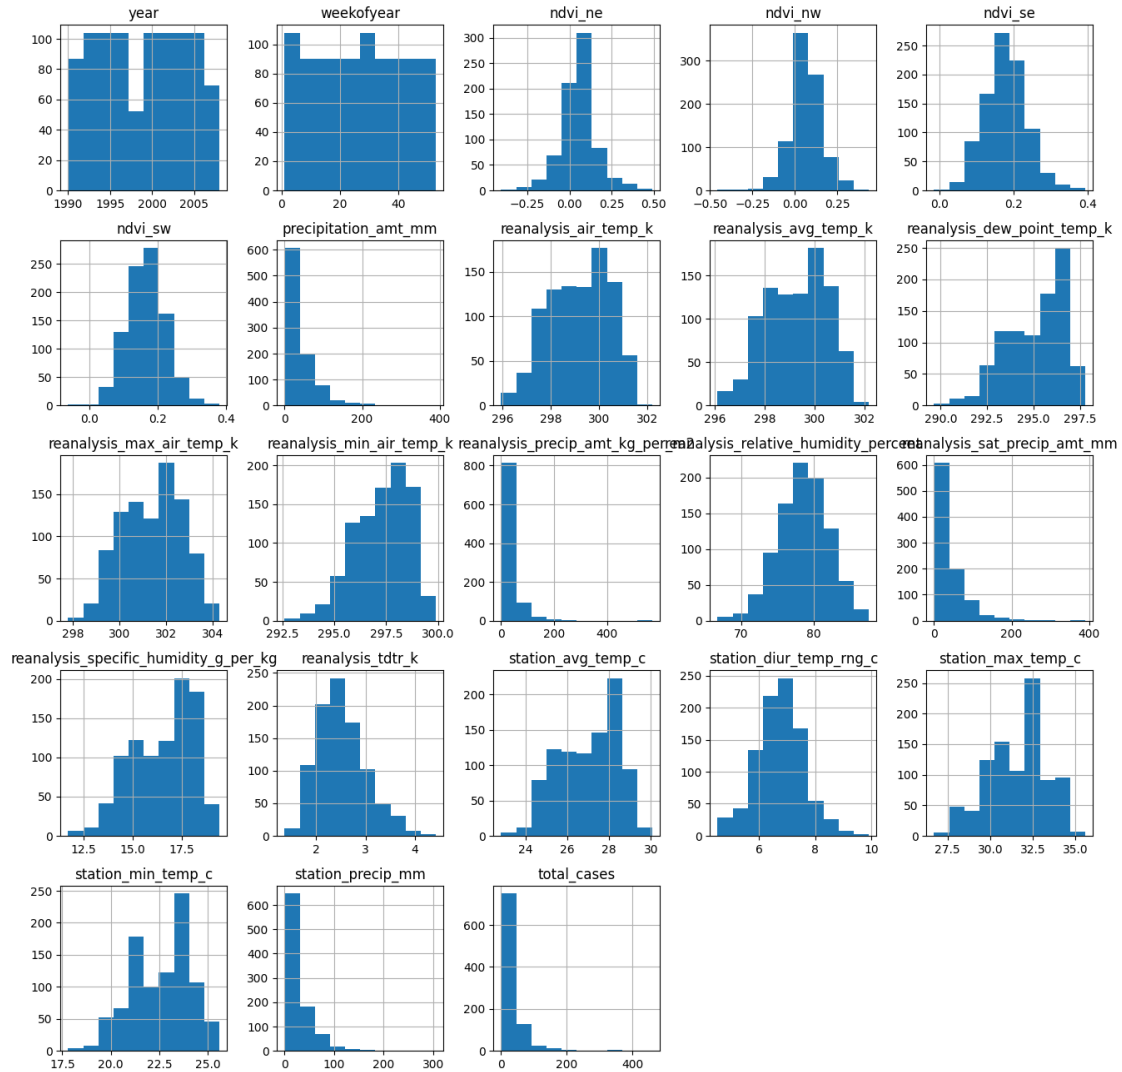
min          0.000000
25%          9.000000
50%         19.000000
75%         37.000000
max         461.000000
Name: total_cases, dtype: float64

```

```

[11]: # Let's see the distribution of the data for each one of the variables.
train.hist(figsize = (16,16));

```




- All feature variables seem to follow more or less a normal distribution except for precipitation features.
- Precipitation follows a positive skew, where there are rare but extremely high values.

```
[12]: # See the distribution of case counts per each successive year:
with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=(11,7))

    years = np.arange(1990,2008,1).astype(int)
    for year in years:
        sns.lineplot(data=train[train.year == year],
                      x="week_start_date", y='total_cases', ax = ax, label =_
↪year)
        ax.get_xaxis().set_ticks([]) # hide week_start_date
        ax.set_xlabel('Time')
        ax.set_ylabel('Total Cases')
        ax.set_title('Total Dengue Cases on a Yearly Basis')
    fig.patch.set_alpha(0) # make the figure background transparent
    plt.legend(bbox_to_anchor=(1.12, 1), loc="upper right");
    fig.savefig('total_cases_years.png', dpi=300)
    files.download("total_cases_years.png")
```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



2.0.1 Overall Summary:

- The data is from 1990 through 2008 with peak outbreaks at certain years.
- Total weekly cases range from 0 to 461.
- There are many null values, especially in ndvi index values.

- Many of the temperature features coded more than once in celcius and fehrenheit using different data sources.
- All variables appear as numerical except for week_start_date which is an object / string.
- All variables appear as normally distributed except target variable and precipitation measures.

2.1 Check for null values:

```
[13]: train.isnull().sum()
      # There are many null values
```

```
[13]: year                0
      weekofyear          0
      week_start_date     0
      ndvi_ne             191
      ndvi_nw             49
      ndvi_se             19
      ndvi_sw             19
      precipitation_amt_mm 9
      reanalysis_air_temp_k 6
      reanalysis_avg_temp_k 6
      reanalysis_dew_point_temp_k 6
      reanalysis_max_air_temp_k 6
      reanalysis_min_air_temp_k 6
      reanalysis_precip_amt_kg_per_m2 6
      reanalysis_relative_humidity_percent 6
      reanalysis_sat_precip_amt_mm 9
      reanalysis_specific_humidity_g_per_kg 6
      reanalysis_tdtr_k    6
      station_avg_temp_c    6
      station_diur_temp_rng_c 6
      station_max_temp_c    6
      station_min_temp_c    6
      station_precip_mm     6
      total_cases           0
      dtype: int64
```

```
[14]: # Proportion of null values for each variable:
      nulls = ((train.isnull().sum()*100) / len(train_features)).
      ↪sort_values(ascending=False)
      nulls[nulls > 0]
```

```
[14]: ndvi_ne            20.405983
      ndvi_nw            5.235043
      ndvi_se            2.029915
      ndvi_sw            2.029915
      precipitation_amt_mm 0.961538
      reanalysis_sat_precip_amt_mm 0.961538
```

```

reanalysis_min_air_temp_k      0.641026
reanalysis_precip_amt_kg_per_m2 0.641026
station_avg_temp_c            0.641026
reanalysis_tdtr_k             0.641026
reanalysis_specific_humidity_g_per_kg 0.641026
reanalysis_relative_humidity_percent 0.641026
reanalysis_dew_point_temp_k    0.641026
station_min_temp_c            0.641026
reanalysis_max_air_temp_k      0.641026
station_max_temp_c            0.641026
reanalysis_avg_temp_k          0.641026
reanalysis_air_temp_k          0.641026
station_precip_mm             0.641026
station_diur_temp_rng_c        0.641026
dtype: float64

```

2.1.1 Display missing values:


- **Missingno** library offers a very nice way to visualize the distribution of Null values.

```


[15]: # Display null values across all rows/columns to check for specific patterns
      ↪ for the absence of data:
missingno.matrix(train_features , figsize=(14, 7), fontsize = 10)

```


```
[15]: <Axes: >
```



```
[16]: # let's see the same data with a bar chart
missingno.bar(train_features, figsize = (14,7), fontsize = 12);
```




```
[17]: # Check the null matrix for the four variables with most null values to see if
      ↪ there is a pattern
missingno.matrix(train_features[['ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw']],
      ↪ figsize=(14, 7), fontsize = 12);
```



- Null values for most climate measures are scarce but ndvi indexes have null values in bigger chunks.

```
[18]: # Heat Map showing the correlation between all variables including the target
corr = train_features.corr().abs()
fig, ax = plt.subplots(figsize=(14,14))
matrix = np.triu(corr) # Getting the Lower Triangle of the correlation matrix
cbar_kws={"label": "Correlation", "shrink":0.8}
heatmap = sns.heatmap(data = corr, linewidths = 1, square= True, 
    cmap='Blues', ax=ax, annot=True, mask=matrix, fmt= ".2f", cbar_kws=cbar_kws)
fig.suptitle('Heatmap of Correlation Between All Features', fontsize=18, y=.84, 
    x = .43);
```

Heatmap of Correlation Between All Features



- There are strong correlations among the majority of the variables.

3 Feature Engineering: Null Replacement

3.0.1 Null replacement using interpolation and predictive modeling:

- We will replace the null values of all climate features except ndvi variables using **interpolation** since the missing data points are scarce.
- We will replace the null values of features for ndvi variables using **k-Nearest Neighbors** since there are bigger chunks of missing values.

[19]: *# Null replacement with interpolation for the below variables:*

```
train_features_interpolated = train_features

vars_to_interpolate = ['precipitation_amt_mm', 'reanalysis_air_temp_k',
```

```

'reanalysis_avg_temp_k', 'reanalysis_dew_point_temp_k',
'reanalysis_max_air_temp_k', 'reanalysis_min_air_temp_k',
'reanalysis_precip_amt_kg_per_m2',
'reanalysis_relative_humidity_percent', 'reanalysis_sat_precip_amt_mm',
'reanalysis_specific_humidity_g_per_kg', 'reanalysis_tdtr_k',
'station_avg_temp_c', 'station_diur_temp_rng_c', 'station_max_temp_c',
'station_min_temp_c', 'station_precip_mm']

for var in vars_to_interpolate:
    train_features_interpolated[var].interpolate(method='linear',
    ↪limit_direction='forward', inplace=True)

```

Null replacement using KNN neighbours for the remaining ndvi variables:

- The default distance measure is a Euclidean distance measure that is NaN aware.
- The number of neighbors is set to five by default and can be configured by the “n_neighbors” argument.
- week_start_date column was dropped and the data was scaled before imputing using KNN neighbours.
- After imputation, reversed the scaling to bring the values back to original scaling.

```

[20]: # Drop `week_start_date` since we cannot work with this variable.
train_features_interpolated = train_features_interpolated.
    ↪drop("week_start_date", axis = 1)

```

```

[21]: # Scale the data first using StandardScaler
scaler = StandardScaler()
train_features_scaled = pd.DataFrame(scaler.
    ↪fit_transform(train_features_interpolated), columns =
    ↪train_features_interpolated.columns)
# Define imputer
imputer = KNNImputer(n_neighbors=5)
# The fit imputer is applied to the dataset to create a copy of the dataset
    ↪with all missing values for each column replaced with an estimated value.
train_features_imputed= pd.DataFrame(imputer.
    ↪fit_transform(train_features_scaled), columns = train_features_scaled.
    ↪columns)
# inverse the Standard Scaling
train_features_full = pd.DataFrame(scaler.
    ↪inverse_transform(train_features_imputed), columns = train_features_imputed.
    ↪columns)
train_features_full.head()

```

```

[21]:      year  weekofyear  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  1990.0         18.0  0.122600  0.103725  0.198483  0.177617
1  1990.0         19.0  0.169900  0.142175  0.162357  0.155486
2  1990.0         20.0  0.032250  0.172967  0.157200  0.170843

```


3	1990.0	21.0	0.128633	0.245067	0.227557	0.235886
4	1990.0	22.0	0.196200	0.262200	0.251200	0.247340

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	\
0	12.42	297.572857	297.742857	
1	22.82	298.211429	298.442857	
2	34.54	298.781429	298.878571	
3	15.36	298.987143	299.228571	
4	7.52	299.518571	299.664286	

	reanalysis_dew_point_temp_k	...	reanalysis_precip_amt_kg_per_m2	\
0	292.414286	...	32.00	
1	293.951429	...	17.94	
2	295.434286	...	26.10	
3	295.310000	...	13.90	
4	295.821429	...	12.20	

	reanalysis_relative_humidity_percent	reanalysis_sat_precip_amt_mm	\
0	73.365714	12.42	
1	77.368571	22.82	
2	82.052857	34.54	
3	80.337143	15.36	
4	80.460000	7.52	

	reanalysis_specific_humidity_g_per_kg	reanalysis_tdtr_k	\
0	14.012857	2.628571	
1	15.372857	2.371429	
2	16.848571	2.300000	
3	16.672857	2.428571	
4	17.210000	3.014286	

	station_avg_temp_c	station_diur_temp_rng_c	station_max_temp_c	\
0	25.442857	6.900000	29.4	
1	26.714286	6.371429	31.7	
2	26.714286	6.485714	32.2	
3	27.471429	6.771429	33.3	
4	28.942857	9.371429	35.0	

	station_min_temp_c	station_precip_mm
0	20.0	16.0
1	22.2	8.6
2	22.8	41.4
3	23.3	4.0
4	23.9	5.8

[5 rows x 22 columns]

```
[22]: # Making sure no null values remained
train_features_imputed.isna().sum().any() == 0
```

[22]: True

```
[23]: # Display the dataset for ndvi values before and after knn imputation
with plt.style.context('seaborn-talk'):
    fig, (ax1, ax2) = plt.subplots(nrows=2, ncols=1, figsize=(14,10))
    for var in ["ndvi_ne", "ndvi_nw", "ndvi_se", "ndvi_sw"]:
        train_features[150:300][var].plot.line(lw=1.2, ax = ax1)
    ax1.set_title('Vegetation Index over Time - Original Data')
    ax1.set_xlabel('Time')
    ax1.legend(loc='lower right')
    for var in ["ndvi_ne", "ndvi_nw", "ndvi_se", "ndvi_sw"]:
        train_features_full[150:300][var].plot.line(lw=1.2, ax = ax2)
    ax2.set_title('Vegetation Index over Time - Imputed Data')
    ax2.set_xlabel('Time')
    ax2.legend(loc='lower right')
    fig.tight_layout();
    fig.patch.set_alpha(0) # make the figure background transparent
    fig.savefig('KNN_ndvi.png', dpi=300, bbox_inches='tight')
    files.download("KNN_ndvi.png")
```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>



```
[24]: # Merge the imputed dataset with the labels
train_full = pd.merge(train_features_full, train_labels, on=[ "year",
↪ "weekofyear" ])
train_full.head()
```

```
[24]:
```

	year	weekofyear	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	\
0	1990.0	18.0	0.122600	0.103725	0.198483	0.177617	
1	1990.0	19.0	0.169900	0.142175	0.162357	0.155486	
2	1990.0	20.0	0.032250	0.172967	0.157200	0.170843	
3	1990.0	21.0	0.128633	0.245067	0.227557	0.235886	
4	1990.0	22.0	0.196200	0.262200	0.251200	0.247340	

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	\
0	12.42	297.572857	297.742857	
1	22.82	298.211429	298.442857	
2	34.54	298.781429	298.878571	
3	15.36	298.987143	299.228571	
4	7.52	299.518571	299.664286	

	reanalysis_dew_point_temp_k	...	reanalysis_relative_humidity_percent	\
0	292.414286	...	73.365714	
1	293.951429	...	77.368571	

2	295.434286	...	82.052857
3	295.310000	...	80.337143
4	295.821429	...	80.460000

	reanalysis_sat_precip_amt_mm	reanalysis_specific_humidity_g_per_kg	\
0	12.42	14.012857	
1	22.82	15.372857	
2	34.54	16.848571	
3	15.36	16.672857	
4	7.52	17.210000	

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c	\
0	2.628571	25.442857	6.900000	
1	2.371429	26.714286	6.371429	
2	2.300000	26.714286	6.485714	
3	2.428571	27.471429	6.771429	
4	3.014286	28.942857	9.371429	

	station_max_temp_c	station_min_temp_c	station_precip_mm	total_cases
0	29.4	20.0	16.0	4
1	31.7	22.2	8.6	5
2	32.2	22.8	41.4	4
3	33.3	23.3	4.0	3
4	35.0	23.9	5.8	6

[5 rows x 23 columns]

3.1 Repeat all imputation steps for the test_features dataset:

```
[25]: # See how many null values present in the test dataset
test_features.isnull().sum()
```

```
[25]: year                0
weekofyear              0
week_start_date         0
ndvi_ne                43
ndvi_nw                11
ndvi_se                1
ndvi_sw                1
precipitation_amt_mm    2
reanalysis_air_temp_k   2
reanalysis_avg_temp_k   2
reanalysis_dew_point_temp_k  2
reanalysis_max_air_temp_k  2
reanalysis_min_air_temp_k  2
reanalysis_precip_amt_kg_per_m2  2
reanalysis_relative_humidity_percent  2
```

```

reanalysis_sat_precip_amt_mm          2
reanalysis_specific_humidity_g_per_kg  2
reanalysis_tdtr_k                      2
station_avg_temp_c                     2
station_diur_temp_rng_c                2
station_max_temp_c                     2
station_min_temp_c                     2
station_precip_mm                       2
dtype: int64

```

```

[26]: # Interpolation
test_features_interpolated = test_features
for var in vars_to_interpolate:
    test_features_interpolated[var].interpolate(method='linear',
    ↪limit_direction='forward', inplace=True)

# Drop week_start_date:
test_features_interpolated = test_features_interpolated.drop("week_start_date",
    ↪axis=1)

# Scale, imputer using KNN imputer, inverse scale
test_features_scaled = pd.DataFrame(scaler.
    ↪fit_transform(test_features_interpolated), columns=
    ↪test_features_interpolated.columns)
test_features_imputed = pd.DataFrame(imputer.
    ↪fit_transform(test_features_scaled), columns=test_features_scaled.columns)
test_features_full = pd.DataFrame(scaler.
    ↪inverse_transform(test_features_imputed), columns=test_features_imputed.
    ↪columns)

# Making sure no null values remained
test_features_full.isna().sum().any() == 0

```

[26]: True

Full Imputed datasets are: * train_full * test_features_full

4 Feature Engineering: Feature Selection / Creation

```

[27]: train_featured = train_full.copy()

```

Let's create a new month variable and dummy coded season variables:

```

[28]: # Add `the week_start_date` column from the original dataset to the new dataset
train_featured['week_start_date'] = train_features['week_start_date']

```

```
[29]: # create a new month variable:
train_featured["week_start_date"] = pd.
    ↳to_datetime(train_featured["week_start_date"])
train_featured['month'] = train_featured['week_start_date'].dt.month
```

```
[30]: # create a new season variable:
seasons = ["winter", "winter", "spring", "spring", "spring",
           "summer", "summer", "summer", "fall", "fall", "fall", "winter"]

month_to_season = dict(zip(range(1,13), seasons))
month_to_season

train_featured['season'] = train_featured['month'].map(month_to_season)
```

```
[31]: # See the new variables:
train_featured[['week_start_date', 'month', 'season']].sample(5)
```

```
[31]:      week_start_date  month  season
485      1999-08-27      8  summer
327      1996-08-12      8  summer
566      2001-03-19      3  spring
892      2007-06-25      6  summer
671      2003-03-26      3  spring
```

```
[32]: # Get the season dummy coded
season_features = pd.get_dummies(train_featured['season'])
train_featured = pd.concat([train_featured, season_features], axis = 1)
train_featured.head()
```

```
[32]:      year  weekofyear  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0  1990.0      18.0  0.122600  0.103725  0.198483  0.177617
1  1990.0      19.0  0.169900  0.142175  0.162357  0.155486
2  1990.0      20.0  0.032250  0.172967  0.157200  0.170843
3  1990.0      21.0  0.128633  0.245067  0.227557  0.235886
4  1990.0      22.0  0.196200  0.262200  0.251200  0.247340

      precipitation_amt_mm  reanalysis_air_temp_k  reanalysis_avg_temp_k  \
0              12.42      297.572857      297.742857
1              22.82      298.211429      298.442857
2              34.54      298.781429      298.878571
3              15.36      298.987143      299.228571
4               7.52      299.518571      299.664286


      reanalysis_dew_point_temp_k  ...  station_min_temp_c  station_precip_mm  \
0              292.414286  ...      20.0              16.0
1              293.951429  ...      22.2               8.6
2              295.434286  ...      22.8             41.4
```

3	295.310000	...			23.3		4.0
4	295.821429	...			23.9		5.8

	total_cases	week_start_date	month	season	fall	spring	summer	winter
0	4	1990-04-30	4	spring	0	1	0	0
1	5	1990-05-07	5	spring	0	1	0	0
2	4	1990-05-14	5	spring	0	1	0	0
3	3	1990-05-21	5	spring	0	1	0	0
4	6	1990-05-28	5	spring	0	1	0	0

[5 rows x 30 columns]

```
[33]: # Seasonality: See the distribution of case counts on a monthly basis:
with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=(10,4))
    sns.lineplot(data=train_featured, x="month", y='total_cases', ax = ax)
    # Maximum number of cases are usually are seen in the fall.
```




5 Feature elimination / selection:

```
[34]: # Show the correlated pairs starting with strongest correlations:
dataCorr = train_featured.drop('total_cases',axis =1).corr().abs()
dataCorr = dataCorr.mask(np.triu(np.ones(dataCorr.shape)).astype(bool)) #
    ↳convert upper triangle of values to NaN to remove repeated values from the
    ↳table
dataCorr = dataCorr.stack().reset_index().sort_values(0, ascending=False) #0 is
    ↳the column automatically generated by the stacking
dataCorr = dataCorr[(dataCorr[0]>.8) & (dataCorr[0]<1)]
```


207	station_avg_temp_c	0.898506
113	reanalysis_avg_temp_k	0.896420
116	reanalysis_min_air_temp_k	0.896376
44	reanalysis_avg_temp_k	0.895373
143	reanalysis_air_temp_k	0.880871
144	reanalysis_avg_temp_k	0.879118
151	reanalysis_specific_humidity_g_per_kg	0.869982
145	reanalysis_dew_point_temp_k	0.868837
188	station_avg_temp_c	0.865240
115	reanalysis_max_air_temp_k	0.853629
146	reanalysis_max_air_temp_k	0.852831
199	reanalysis_dew_point_temp_k	0.850479
205	reanalysis_specific_humidity_g_per_kg	0.849573
54	reanalysis_dew_point_temp_k	0.847654
147	reanalysis_min_air_temp_k	0.841300
197	reanalysis_air_temp_k	0.833158
201	reanalysis_min_air_temp_k	0.829792
65	reanalysis_max_air_temp_k	0.828665
198	reanalysis_avg_temp_k	0.827497
14	ndvi_se	0.820109

[35]: *# Show how strongly the features are correlated with the target variable - total cases:*

```
fig, ax = plt.subplots(figsize=(14,7))
train_featured.corr()['total_cases'].drop('total_cases').
    sort_values(ascending=False).plot.barh(ax=ax);
```



- Many of the temperature data are strongly correlated with one another especially because the same feature was coded multiple times from different resources.
- However, none of the features seem to have a strong relationship with the target variable - total_cases. Total_cases seems to only have weak correlations with other variables.

- We need to engineer some new features hoping they would have stronger relationship with `total_cases`.

5.0.1 Select the best average temperature variable:

- `station_avg_temp_c` has the strongest correlation


```
[36]: # see the correlations among all average temperature variables:
train_featured[['total_cases', 'station_avg_temp_c', 'reanalysis_air_temp_k', 'reanalysis_avg_tem
↪corr()
```

```
[36]:
```

	total_cases	station_avg_temp_c	reanalysis_air_temp_k	\
total_cases	1.000000	0.194823	0.180311	
station_avg_temp_c	0.194823	1.000000	0.880871	
reanalysis_air_temp_k	0.180311	0.880871	1.000000	
reanalysis_avg_temp_k	0.173670	0.879118	0.997507	

	reanalysis_avg_temp_k
total_cases	0.173670
station_avg_temp_c	0.879118
reanalysis_air_temp_k	0.997507
reanalysis_avg_temp_k	1.000000

```
[37]: with plt.style.context('seaborn-talk'):
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=(14,4))
sns.regplot(data=train_featured, x="station_avg_temp_c", y='total_cases',
↪ax = ax1, label = 'station_avg_temp_c')
sns.regplot(data=train_featured, x="reanalysis_air_temp_k",
↪y='total_cases', ax = ax2, label = 'reanalysis_air_temp_k')
```



5.0.2 Select the best daily temperature change variable:

- `reanalysis_tdtr_k` has the strongest correlation
- Let's also replace the single outlier with a better value.


```
[38]: train_featured[['total_cases', 'station_diur_temp_rng_c', 'reanalysis_tdtr_k']].
      ↪corr()
```

```
[38]:
```

	total_cases	station_diur_temp_rng_c	\
total_cases	1.000000	0.035303	
station_diur_temp_rng_c	0.035303	1.000000	
reanalysis_tdtr_k	-0.067623	0.372414	

	reanalysis_tdtr_k
total_cases	-0.067623
station_diur_temp_rng_c	0.372414
reanalysis_tdtr_k	1.000000

```
[39]: with plt.style.context('seaborn-talk'):
      fig, ax = plt.subplots(figsize=(10,4))
      sns.regplot(data=train_featured, x="station_diur_temp_rng_c",
      ↪y='total_cases', ax = ax, label = 'station_diur_temp_rng_c')
      sns.regplot(data=train_featured, x="reanalysis_tdtr_k", y='total_cases', ax=
      ↪ax, label = 'reanalysis_tdtr_k')
```



```
[40]: # check out the outlier for 'reanalysis_tdtr_k':
      train_featured[train_featured['reanalysis_tdtr_k'] ==
      ↪train_featured['reanalysis_tdtr_k'].max()]
```

```
[40]:
```

	year	weekofyear	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	\
799	2005.0	36.0	0.0022	-0.0271	0.205029	0.220233	

	precipitation_amt_mm	reanalysis_air_temp_k	reanalysis_avg_temp_k	\
799	23.3	301.465714	301.514286	

```

reanalysis_dew_point_temp_k ... station_min_temp_c station_precip_mm \
799          296.642857 ...          24.4          8.9

total_cases week_start_date month season fall spring summer winter
799          131      2005-09-10     9   fall     1     0     0     0

[1 rows x 30 columns]

```

```

[41]: # replace the outlier with the previous value in the series
train_featured = train_featured.replace(train_featured['reanalysis_tdtr_k'].
    ↳max(), method='ffill')
train_featured[799:800]

```

```

[41]:      year weekofyear ndvi_ne ndvi_nw ndvi_se ndvi_sw \
799  2005.0         36.0  0.0022 -0.0271  0.205029  0.220233

precipitation_amt_mm reanalysis_air_temp_k reanalysis_avg_temp_k \
799          23.3          301.465714          301.514286

reanalysis_dew_point_temp_k ... station_min_temp_c station_precip_mm \
799          296.642857 ...          24.4          8.9

total_cases week_start_date month season fall spring summer winter
799          131      2005-09-10     9   fall     1     0     0     0

[1 rows x 30 columns]

```

```

[42]: # check the correlations again:
train_featured[['total_cases', 'station_diur_temp_rng_c', 'reanalysis_tdtr_k']].
    ↳corr()

```

```

[42]:      total_cases station_diur_temp_rng_c \
total_cases      1.000000      0.035303
station_diur_temp_rng_c      0.035303      1.000000
reanalysis_tdtr_k      -0.073160      0.374047

      reanalysis_tdtr_k
total_cases      -0.073160
station_diur_temp_rng_c      0.374047
reanalysis_tdtr_k      1.000000


```

```

[43]: with plt.style.context('seaborn-talk'):
fig, ax = plt.subplots(figsize=(10,4))
sns.regplot(data=train_featured, x="station_diur_temp_rng_c",
    ↳y='total_cases', ax = ax, label = 'station_diur_temp_rng_c')

```

```
sns.regplot(data=train_featured, x="reanalysis_tdtr_k", y='total_cases', ax=
↪ ax, label = 'reanalysis_tdtr_k')
```



5.0.3 Select the best humidity variable:

- reanalysis_specific_humidity_g_per_kg has the strongest correlation

```
[44]: train_featured[['total_cases', 'reanalysis_relative_humidity_percent',
↪ 'reanalysis_specific_humidity_g_per_kg', 'reanalysis_dew_point_temp_k']].
↪ corr()
```

```
[44]:
```


	total_cases	\
total_cases	1.000000	
reanalysis_relative_humidity_percent	0.144404	
reanalysis_specific_humidity_g_per_kg	0.206942	
reanalysis_dew_point_temp_k	0.202807	

	reanalysis_relative_humidity_percent	\
total_cases	0.144404	
reanalysis_relative_humidity_percent	1.000000	
reanalysis_specific_humidity_g_per_kg	0.673010	
reanalysis_dew_point_temp_k	0.678116	

	reanalysis_specific_humidity_g_per_kg	\
total_cases	0.206942	
reanalysis_relative_humidity_percent	0.673010	
reanalysis_specific_humidity_g_per_kg	1.000000	
reanalysis_dew_point_temp_k	0.998533	

	reanalysis_dew_point_temp_k
total_cases	0.202807
reanalysis_relative_humidity_percent	0.678116
reanalysis_specific_humidity_g_per_kg	0.998533
reanalysis_dew_point_temp_k	1.000000

```
[45]: with plt.style.context('seaborn-talk'):
fig, (ax1, ax2, ax3) = plt.subplots(ncols = 3, nrows = 1, figsize=(16,4))
sns.regplot(data=train_featured, x="reanalysis_relative_humidity_percent",
↪y='total_cases', ax = ax1, label = 'reanalysis_relative_humidity_percent')
sns.regplot(data=train_featured, x="reanalysis_specific_humidity_g_per_kg",
↪y='total_cases', ax = ax2, label = 'reanalysis_specific_humidity_g_per_kg')
sns.regplot(data=train_featured, x="reanalysis_dew_point_temp_k",
↪y='total_cases', ax = ax3, label = 'reanalysis_dew_point_temp_k')
```



5.0.4 Select the best precipitation variable:

- reanalysis_precip_amt_kg_per_m2 has the strongest correlation

```
[46]: train_featured[['total_cases', 'reanalysis_sat_precip_amt_mm', 'station_precip_mm',
↪'reanalysis_precip_amt_kg_per_m2', 'precipitation_amt_mm']].
↪corr()
```

```
[46]:
```

	total_cases	reanalysis_sat_precip_amt_mm	\
total_cases	1.000000	0.060296	
reanalysis_sat_precip_amt_mm	0.060296	1.000000	
station_precip_mm	0.051883	0.566660	
reanalysis_precip_amt_kg_per_m2	0.106939	0.508274	
precipitation_amt_mm	0.060296	1.000000	


	station_precip_mm	\
total_cases	0.051883	
reanalysis_sat_precip_amt_mm	0.566660	

station_precip_mm	1.000000
reanalysis_precip_amt_kg_per_m2	0.477984
precipitation_amt_mm	0.566660

	reanalysis_precip_amt_kg_per_m2 \
total_cases	0.106939
reanalysis_sat_precip_amt_mm	0.508274
station_precip_mm	0.477984
reanalysis_precip_amt_kg_per_m2	1.000000
precipitation_amt_mm	0.508274

	precipitation_amt_mm
total_cases	0.060296
reanalysis_sat_precip_amt_mm	1.000000
station_precip_mm	0.566660
reanalysis_precip_amt_kg_per_m2	0.508274
precipitation_amt_mm	1.000000

```
[47]: with plt.style.context('seaborn-talk'):
fig, ax = plt.subplots(figsize=(12,4))
for var in ["reanalysis_sat_precip_amt_mm", "station_precip_mm",
↪ "reanalysis_precip_amt_kg_per_m2", "precipitation_amt_mm"]:
    sns.regplot(data=train_featured, x=var, y='total_cases', ax = ax, label =
↪ var)
plt.legend()
```



5.1 Summary - feature selection:

Let's focus on these variables below since they provide the highest correlations. Let's keep all the temperature variables same scale (celcius) for interpretability except the diurnal which gave better correlation in kelvin.

Using NOAA's NCEP Climate Forecast System Reanalysis measurements (0.5x0.5 degree scale):

Temperature variables: station_avg_temp_c, station_min_temp_c, station_max_temp_c, reanalysis_tdtr_k (Diurnal temperature range)


Humidity: reanalysis_specific_humidity_g_per_kg

Precipitation: reanalysis_precip_amt_kg_per_m2 (Total precipitation)

Vegetation: ndvi_ne, ndvi_nw, ndvi_se, ndvi_sw

```
[48]: # Keeping the below variables as primary:
train_featured = train_featured[['total_cases', 'year', 'weekofyear',
    ↪ 'week_start_date',
    'month', 'fall', 'spring', 'summer', 'winter',
    'station_avg_temp_c', 'station_max_temp_c',
    'station_min_temp_c', 'reanalysis_tdtr_k',
    'reanalysis_specific_humidity_g_per_kg',
    'reanalysis_precip_amt_kg_per_m2',
    'ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw']]
```

```
[49]: # See how these hand picked variables correlate with total cases
fig, ax = plt.subplots(figsize=(12,6))
train_featured.corr()['total_cases'].drop('total_cases').
    ↪ sort_values(ascending=False).plot.barh(ax=ax);
```



```
[50]: # Plot how the total case numbers differ based on each climate feature:
columns = ['station_avg_temp_c', 'station_max_temp_c',
    'station_min_temp_c', 'reanalysis_precip_amt_kg_per_m2',
    'reanalysis_tdtr_k', 'reanalysis_specific_humidity_g_per_kg',
    'ndvi_ne', 'ndvi_nw',
    'ndvi_se', 'ndvi_sw']

labels = ["Avg Temp", "Max Temp",
    "Min Temp", "Precipitation",
```



```

        "Daily Temp Range1", "Specific_Humidity",
        "Ndvi NE", "Ndvi_NW",
        "Ndvi_SE", "Ndvi_SW"]

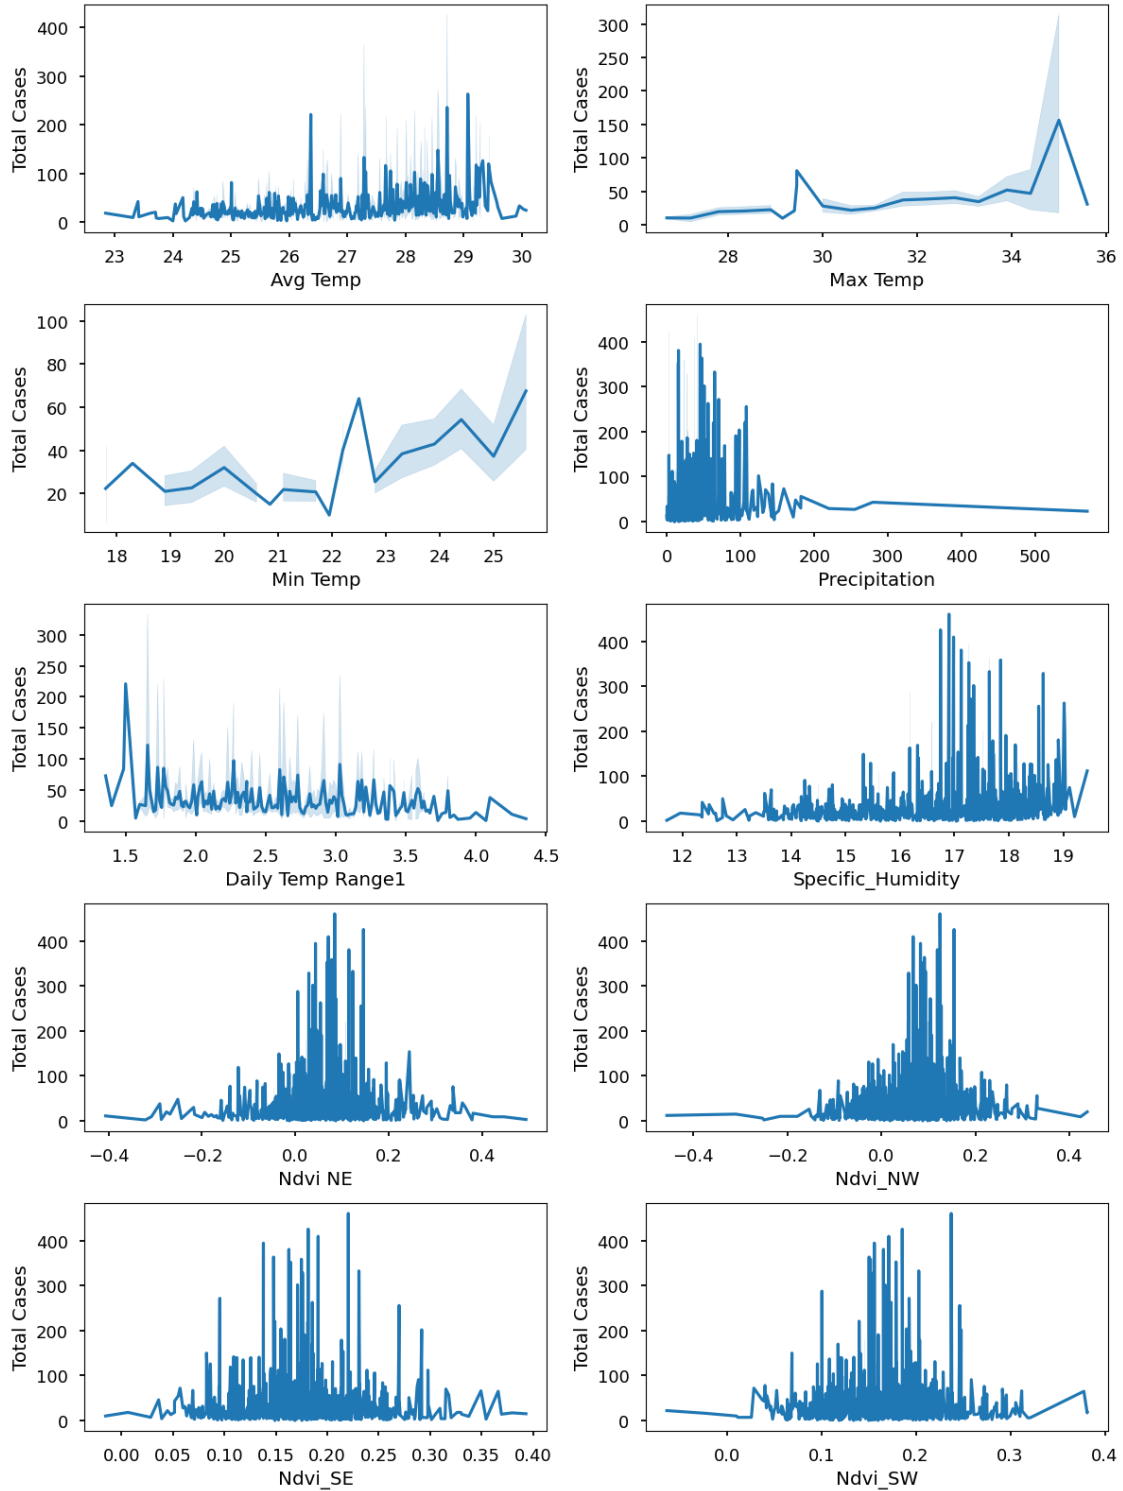
nrows =5
ncols =2
with plt.style.context('seaborn-talk'):

    fig, ax_list = plt.subplots(nrows = nrows, ncols = ncols, figsize=(12,16))

    j=0
    for i in range(nrows):
        for u in range(ncols):
            sns.lineplot(data =train_featured, x=columns[j], y="total_cases",
↪ax = ax_list[i,u]) # need to use index for column because otherwise it does
↪not iterate.
            ax_list[i,u].set_xlabel(labels[j])
            ax_list[i,u].set_ylabel("Total Cases")
            j = j+1

    fig.tight_layout();

```



- For all the variables, the relationship seems stronger until case number 100 reaches, possible because there are fewer extra high dengue cases.
- Extremely high **average temp**, **maximum temp**, **temp range**, or **precipitation** seem to

impact dengue cases negatively.

- There is no clear linear relationship between ndvi variables and total cases as lower and higher values tend to result in lower total cases, but moderate values tend to result in higher total cases.

5.2 Convert NDVI into Categorical variables:

- NDVI calculation range from -1 to 1. Negative values correspond to areas with water surfaces, manmade structures, rocks, clouds, snow. Bare soil usually falls within 0.1- 0.2 range. Plants will always have positive values between 0.2 and 1. Healthy, dense vegetation canopy should be above 0.5. Sparse vegetation will most likely fall within 0.2 to 0.5.
- Since there is no clear linear relationship between ndvi and total cases, let's create a categorical version of the variables. |
 - Let's first create a new feature representing the average NDVI values from the four different locations.
 - Then let's create a categorical version of the variable to represent **watery**, **soily**, **sparse_grassy** areas.

```
[51]: # Create `ndvi_average` by taking the mean of all 4 coordinates.
```

```
train_featured['ndvi_average'] =   
    (train_featured['ndvi_ne']+train_featured['ndvi_nw']+  
       
     train_featured['ndvi_se']+train_featured['ndvi_sw'])/4
```

```
[52]: # Let's check the distribution of values:
```


```
train_featured[['ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw', 'ndvi_average']].  
describe()
```

```
[52]:
```

	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	ndvi_average
count	936.000000	936.000000	936.000000	936.000000	936.000000
mean	0.057266	0.067853	0.177654	0.165855	0.117157
std	0.100001	0.090603	0.056694	0.055681	0.056231
min	-0.406250	-0.456100	-0.015533	-0.063457	-0.092565
25%	0.008050	0.018706	0.139862	0.129778	0.079570
50%	0.057667	0.068750	0.177171	0.165906	0.112724
75%	0.108288	0.115017	0.212336	0.202549	0.147122
max	0.493400	0.437100	0.393129	0.381420	0.342338

```
[53]: # Let's see the bottom and top 15th percentile for `ndvi_average`:
```

```
fig, ax = plt.subplots(figsize=(12,6))  
sns.regplot(data = train_featured, x='ndvi_average', y="total_cases", ax = ax)  
ax.axvline(x=train_featured['ndvi_average'].quantile(0.15), ymin=0, ymax=1,  
           color='y', linestyle='--')  
ax.axvline(x=train_featured['ndvi_average'].quantile(0.85), ymin=0, ymax=1,  
           color='y', linestyle='--');
```




- Total cases seem to be low for low and high ndvi, and high for moderate ndvi

```
[54]: # Create a function to transform `ndvi_average` into a categorical variables.
# watery (dead plants) --> soily (unhealthy plants) --> grassy (healthy plants)
def get_ndvi_cat(x):
    if x < train_featured['ndvi_average'].quantile(0.15):
        return "watery"
    if x >= train_featured['ndvi_average'].quantile(0.15) and x <=
    ↪ train_featured['ndvi_average'].quantile(0.85):
        return "soily"
    else:
        return "grassy"
```

```
[55]: # Apply the transformation
train_featured["ndvi_average_cat"] = train_featured["ndvi_average"].
    ↪ apply(get_ndvi_cat)
```

```
[56]: # See the count values after transformation
train_featured['ndvi_average_cat'].value_counts().plot(kind='bar');
```



```
[57]: # Get dummy codes for 'ndvi_average_cat' and merge with the dataset:
ndvi_features = pd.get_dummies(train_featured['ndvi_average_cat'])
train_featured = pd.concat([train_featured, ndvi_features], axis = 1)
train_featured.head()
```

```
[57]:
```

	total_cases	year	weekofyear	week_start_date	month	fall	spring	\
0	4	1990.0	18.0	1990-04-30	4	0	1	
1	5	1990.0	19.0	1990-05-07	5	0	1	
2	4	1990.0	20.0	1990-05-14	5	0	1	
3	3	1990.0	21.0	1990-05-21	5	0	1	
4	6	1990.0	22.0	1990-05-28	5	0	1	

	summer	winter	station_avg_temp_c	...	reanalysis_precip_amt_kg_per_m2	\
0	0	0	25.442857	...	32.00	
1	0	0	26.714286	...	17.94	
2	0	0	26.714286	...	26.10	
3	0	0	27.471429	...	13.90	
4	0	0	28.942857	...	12.20	

	ndvi_ne	ndvi_nw	ndvi_se	ndvi_sw	ndvi_average	ndvi_average_cat	\
--	---------	---------	---------	---------	--------------	------------------	---

0	0.122600	0.103725	0.198483	0.177617	0.150606	soily
1	0.169900	0.142175	0.162357	0.155486	0.157479	soily
2	0.032250	0.172967	0.157200	0.170843	0.133315	soily
3	0.128633	0.245067	0.227557	0.235886	0.209286	grassy
4	0.196200	0.262200	0.251200	0.247340	0.239235	grassy

	grassy	soily	watery
0	0	1	0
1	0	1	0
2	0	1	0
3	1	0	0
4	1	0	0

[5 rows x 24 columns]

```
[58]: # Let's see if correlations improved:
ndvi_data =
↳ train_featured[['total_cases', 'ndvi_average', 'grassy', 'soily', 'watery']]
ndvi_data.corr()['total_cases']
```

```
[58]: total_cases      1.000000
ndvi_average      0.052466
grassy           -0.043124
soily             0.102880
watery           -0.088839
Name: total_cases, dtype: float64
```

5.3 Create new shifted variables with rolled means:

- Research seems to indicated that past **sustained** heat, precipitation or humidity impacts dengue cases more profoundly than the climate situation right at the time of cases.
- I will be creating a series of rolled means for each of the chosen variable.
 - First shift the variables by 2 weeks to account for the growth of mosquito and the incubation period of the virus until testing positive.
 - Then create rolled - cumulative means with a range of lags to find the variable with the highest correlation.

```
[59]: train_shifted = train_featured.copy()
```

```
[60]: # The variables to shift and roll:
shifted_varbls = ['station_avg_temp_c', 'station_max_temp_c',
↳ 'station_min_temp_c', 'reanalysis_tdtr_k',
                  'reanalysis_specific_humidity_g_per_kg',
↳ 'reanalysis_precip_amt_kg_per_m2',
                  'grassy', 'soily', 'watery']
```

```
[61]: # shift the variables two weeks ahead so that total cases would correspond with
      ↪ climate variables from two weeks ago:
      for var in shifted_varbls:
          train_shifted[f"{var}_shift"] = train_shifted[var].shift(2)
      # drop the two rows with NA
      train_shifted.dropna(axis=0, inplace=True)
      # reset the index
      train_shifted = train_shifted.reset_index(drop=True)
```

```
[62]: # Making sure it shifted corretly
      train_shifted[['year', 'weekofyear', 'week_start_date', 'station_avg_temp_c',
      ↪ 'station_avg_temp_c_shift', 'grassy', 'grassy_shift']].head()
```

```
[62]:
```

	year	weekofyear	week_start_date	station_avg_temp_c \
0	1990.0	20.0	1990-05-14	26.714286
1	1990.0	21.0	1990-05-21	27.471429
2	1990.0	22.0	1990-05-28	28.942857
3	1990.0	23.0	1990-06-04	28.114286
4	1990.0	24.0	1990-06-11	27.414286

	station_avg_temp_c_shift	grassy	grassy_shift
0	25.442857	0	0.0
1	26.714286	1	0.0
2	26.714286	1	0.0
3	27.471429	1	1.0
4	28.942857	0	1.0

```
[63]: # create another copy to get the rolled means
      train_rolled = train_shifted.copy()
```

```
[64]: # We will check the correlations between these variables and its shifted
      ↪ versions
      varbls_to_see_lags = ['total_cases',
                           'reanalysis_precip_amt_kg_per_m2',
                           'reanalysis_specific_humidity_g_per_kg',
                           'reanalysis_tdtr_k',
                           'station_avg_temp_c',
                           'station_max_temp_c',
                           'station_min_temp_c',
                           'grassy', 'soily', 'watery',
                           'reanalysis_precip_amt_kg_per_m2_shift',
                           'reanalysis_specific_humidity_g_per_kg_shift',
                           'reanalysis_tdtr_k_shift',
                           'station_avg_temp_c_shift',
                           'station_max_temp_c_shift',
                           'station_min_temp_c_shift',
                           'grassy_shift', 'soily_shift', 'watery_shift']
```

```
[65]: train_rolled = train_rolled[varbls_to_see_lags]
```

```
[66]: rolled_varbls = ['reanalysis_precip_amt_kg_per_m2_shift',
                      'reanalysis_specific_humidity_g_per_kg_shift',
                      'reanalysis_tdtr_k_shift',
                      'station_avg_temp_c_shift',
                      'station_max_temp_c_shift',
                      'station_min_temp_c_shift',
                      'grassy_shift', 'soily_shift', 'watery_shift']
```

```
[67]: # Create cumulative means for lags of 2 through 20 weeks (about 3-4 months):
# Use a min period of 10 for a lag of 20 so we do not lose all the first 20
# weeks of data.
window = np.linspace(2,20,10).astype(int)
min_periods = np.linspace(1,10,10).astype(int)

for var in rolled_varbls:
    for num,min in zip(window,min_periods):
        train_rolled[f"{var}_{num}"] = train_rolled[var].rolling(window = num,
#min_periods = min).mean()
```

```
[68]: # Create 4 separate datasets for temp, humid, prec, ndvi variables with shifted
# and rolled versions:
temp_cols = [col for col in train_rolled.columns if 'temp' in col or 'tdtr' in
#col]
hum_cols = [col for col in train_rolled.columns if 'hum' in col]
prec_cols = [col for col in train_rolled.columns if 'prec' in col]
ndvi_cols = [col for col in train_rolled.columns if 'grassy' in col or 'soily'
#in col or 'watery' in col]
```

```
[94]: # Add total_cases
temp = train_rolled[temp_cols]
temp['total_cases'] = train_rolled['total_cases']

hum = train_rolled[hum_cols]
hum['total_cases'] = train_rolled['total_cases']


prec = train_rolled[prec_cols]
prec['total_cases'] = train_rolled['total_cases']

ndvi = train_rolled[ndvi_cols]
ndvi['total_cases'] = train_rolled['total_cases']
```

```
[70]: # see the correlations between all temperature variables along with their
# lagged versions and total case counts
fig, ax = plt.subplots(figsize=(12,12))
```



```
temp.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).
    plot.barh(ax=ax);
```



```
[71]: # create a function to see how lagged versions of a variable is related to
    total cases.
def lag_graph(df, var):

    columns = [f"{var}", f"{var}_shift",
               f"{var}_shift_2", f"{var}_shift_4",
               f"{var}_shift_6", f"{var}_shift_8",
               f"{var}_shift_10", f"{var}_shift_12",
               f"{var}_shift_14", f"{var}_shift_16",
               f"{var}_shift_18", f"{var}_shift_20"]

    labels = ["original", "shift_2", "shift_2 + roll_2", "shift_2 + roll_4",
    "shift_2 + roll_6",
               "shift_2 + roll_8", "shift_2 + roll_10", "shift_2 + roll_12",
    "shift_2 + roll_14",
```

```

        "shift_2 + roll_16", "shift_2 + roll_18", "shift_2 + roll_20"]

ncols = 2
nrows = 6
with plt.style.context('seaborn-talk'):

    fig, ax_list = plt.subplots(nrows = nrows, ncols = ncols,
↪figsize=(12,16))

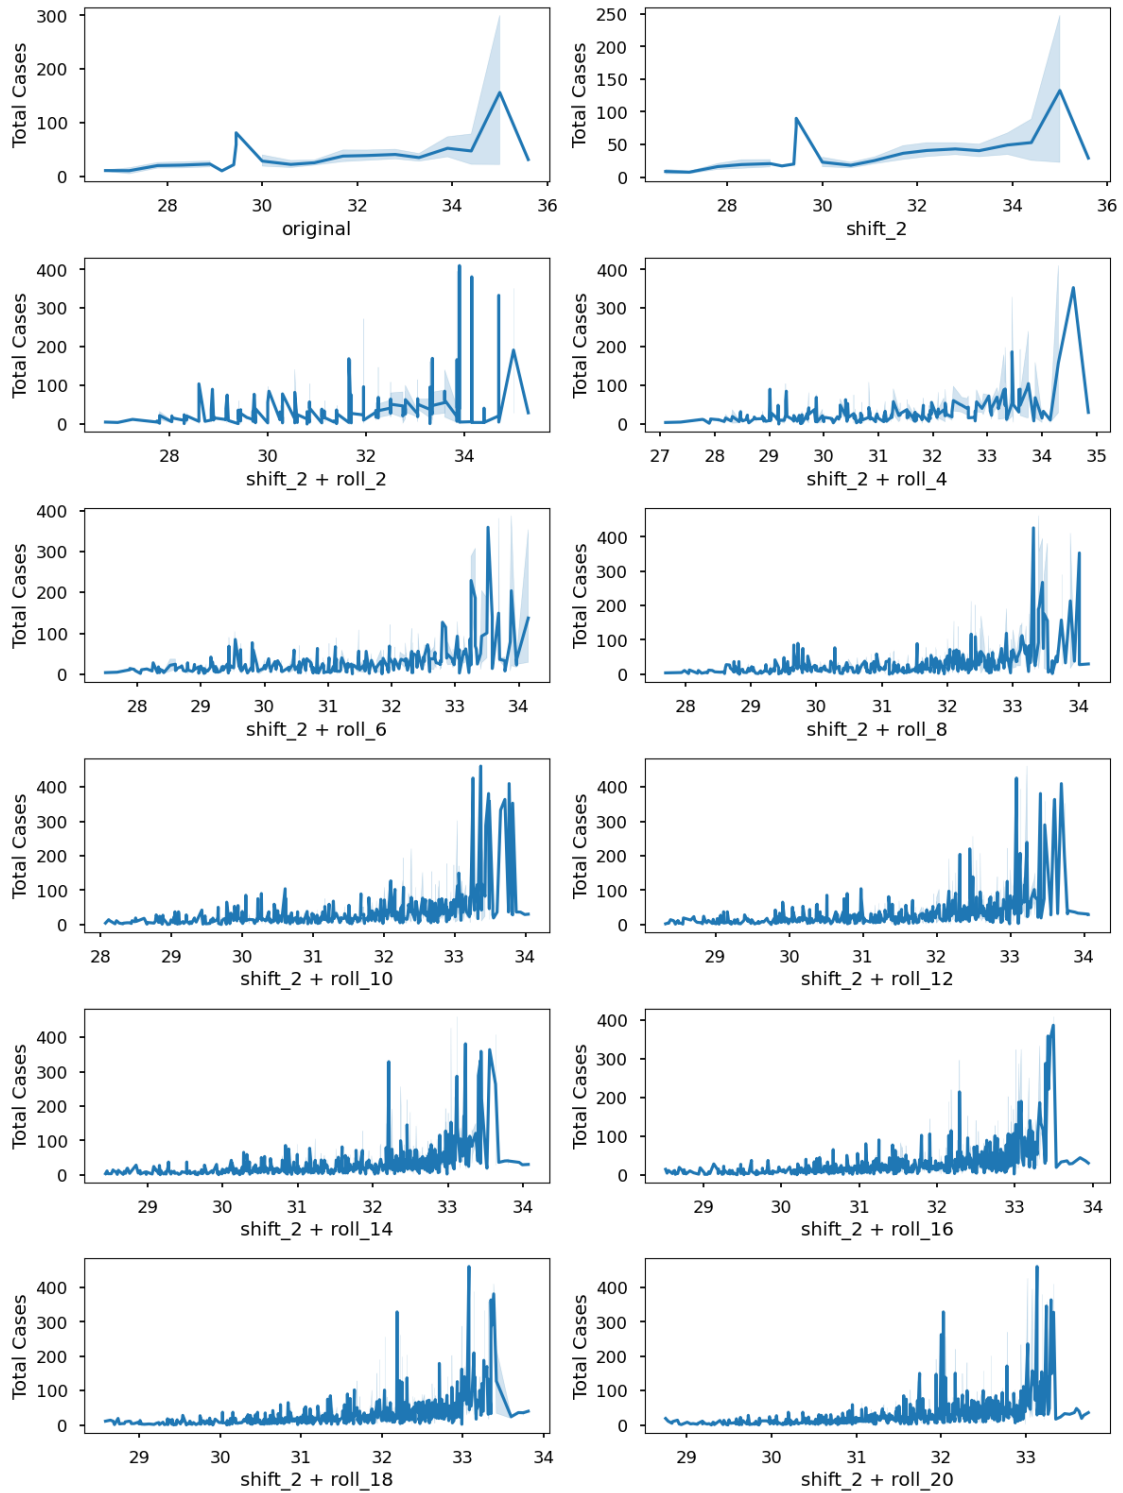
    j=0
    for i in range(nrows):
        for u in range(ncols):
            sns.lineplot(data = df, x=columns[j], y="total_cases", ax =
↪ax_list[i,u]) # need to use index for column because otherwise it does not
↪iterate.

            ax_list[i,u].set_xlabel(labels[j])
            ax_list[i,u].set_ylabel("Total Cases")
            j = j+1

    fig.tight_layout();


```

```
[72]: lag_graph(train_rolled, "station_max_temp_c")
```

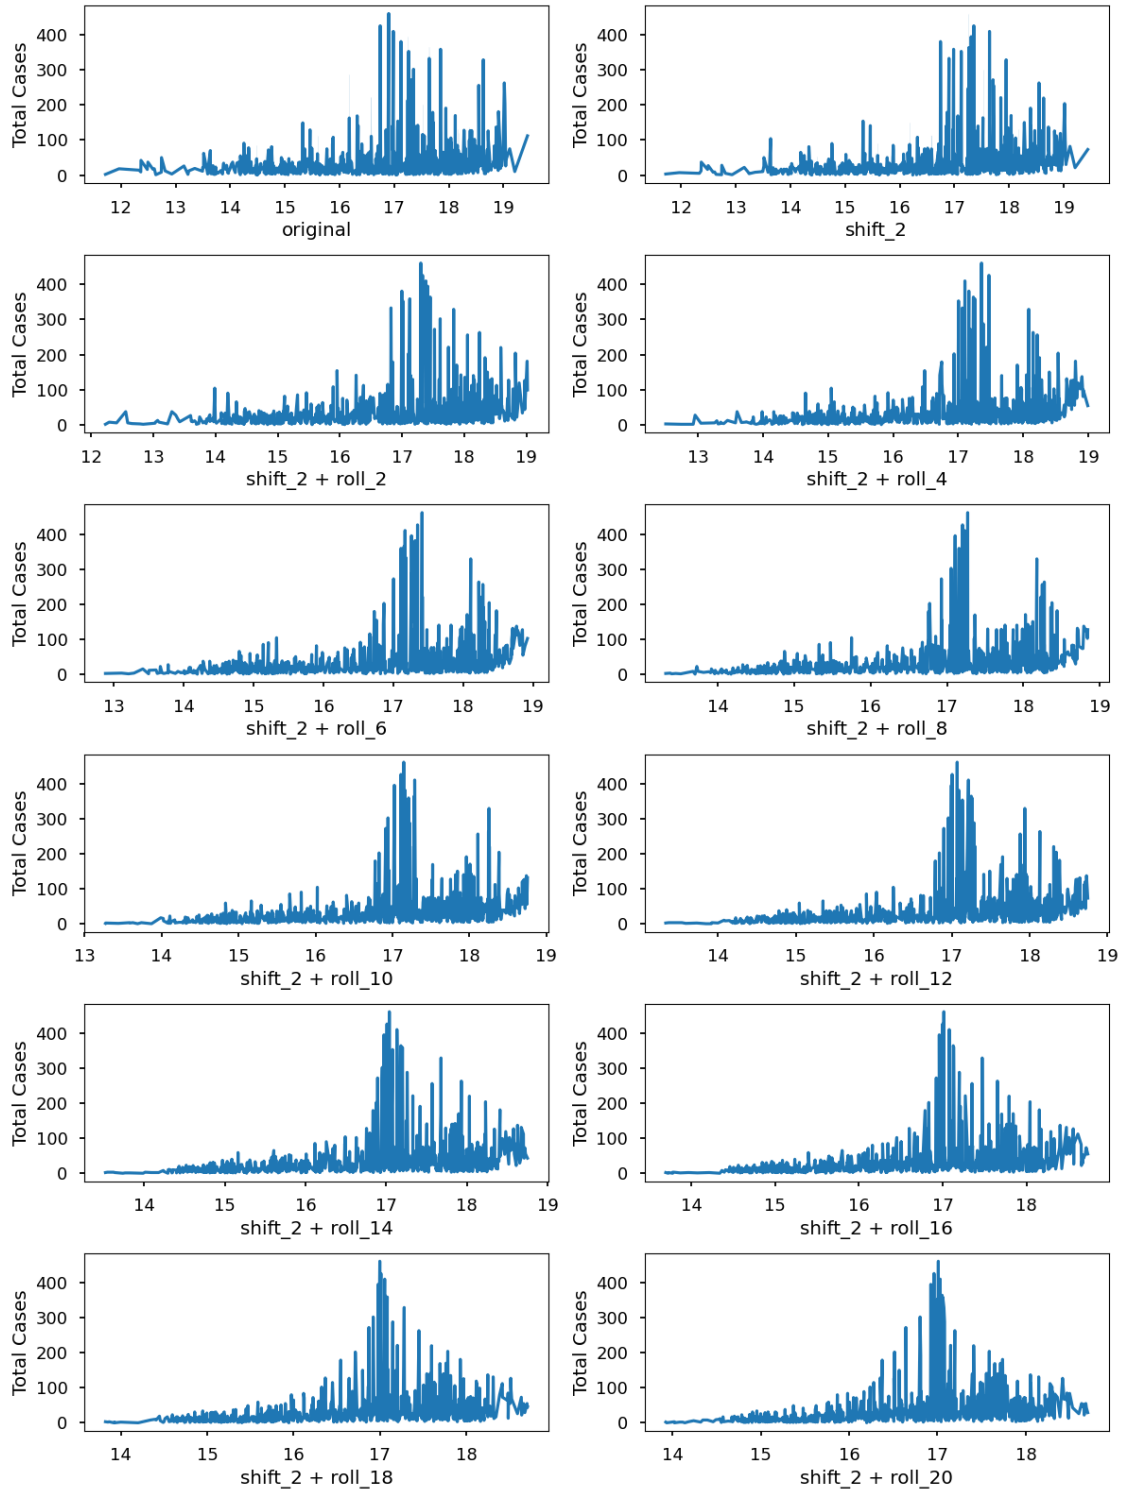


```
[73]: # see the correlations between the humidity variable along with its lagged
      ↪ versions and total case counts
```

```
fig, ax = plt.subplots(figsize=(10,6))
hum.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).plot.
↳ barh(ax=ax);
```




```
[74]: lag_graph(train_rolled, "reanalysis_specific_humidity_g_per_kg")
```




```
[75]: # see the correlations between the precipitation variable along with its lagged
      ↪ versions and total case counts
```

```
fig, ax = plt.subplots(figsize=(10,6))
prec.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).
    ↪plot.barh(ax=ax);
```



[76]: *# see the correlations between the vegetation variables along with their lagged versions and total case counts*

```
fig, ax = plt.subplots(figsize=(10,12))
ndvi.corr()['total_cases'].drop('total_cases').sort_values(ascending=False).
    ↪plot.barh(ax=ax);
```



6 Based on above graphs I will be using these variables which provided the highest correlations to total cases:

- station_max_temp_c_shift_18,
- station_min_temp_c_shift_18,
- station_avg_temp_c_shift_18,
- reanalysis_tdtr_k_shift_8,
- reanalysis_specific_humidity_g_per_kg_shift_12,
- reanalysis_precip_amt_kg_per_m2_shift_8,
- grassy_shift_20,

- soily_shift_20,
- watery_shift_20

```
[77]: rolled_varblds_to_use = ['station_max_temp_c_shift_18',
                              'station_min_temp_c_shift_18',
                              'station_avg_temp_c_shift_18',
                              'reanalysis_tdtr_k_shift_8',
                              'reanalysis_specific_humidity_g_per_kg_shift_12',
                              'reanalysis_precip_amt_kg_per_m2_shift_8',
                              'grassy_shift_20',
                              'soily_shift_20', 'watery_shift_20']
```

```
[79]: # Add the rolled variables to the dataset
train_final = train_shifted.join(train_rolled[rolled_varblds_to_use])
train_final.head(10)
```

```
[79]:   total_cases   year  weekofyear week_start_date  month  fall  spring \
0             4  1990.0         20.0   1990-05-14     5     0      1
1             3  1990.0         21.0   1990-05-21     5     0      1
2             6  1990.0         22.0   1990-05-28     5     0      1
3             2  1990.0         23.0   1990-06-04     6     0      0
4             4  1990.0         24.0   1990-06-11     6     0      0
5             5  1990.0         25.0   1990-06-18     6     0      0
6            10  1990.0         26.0   1990-06-25     6     0      0
7             6  1990.0         27.0   1990-07-02     7     0      0
8             8  1990.0         28.0   1990-07-09     7     0      0
9             2  1990.0         29.0   1990-07-16     7     0      0
```

```
   summer  winter  station_avg_temp_c  ...  watery_shift  \
0         0       0          26.714286  ...          0.0
1         0       0          27.471429  ...          0.0
2         0       0          28.942857  ...          0.0
3         1       0          28.114286  ...          0.0
4         1       0          27.414286  ...          0.0
5         1       0          28.371429  ...          0.0
6         1       0          28.328571  ...          0.0
7         1       0          28.328571  ...          0.0
8         1       0          27.557143  ...          0.0
9         1       0          28.128571  ...          0.0
```

```
   station_max_temp_c_shift_18  station_min_temp_c_shift_18  \
0                          NaN                          NaN
1                          NaN                          NaN
2                          NaN                          NaN
3                          NaN                          NaN
4                          NaN                          NaN
5                          NaN                          NaN
```


6	NaN	NaN
7	NaN	NaN
8	32.888889	22.777778
9	32.990000	22.940000

	station_avg_temp_c_shift_18	reanalysis_tdtr_k_shift_8 \
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	2.432143
4	NaN	2.548571
5	NaN	2.473810
6	NaN	2.412245
7	NaN	2.307143
8	27.501587	2.214286
9	27.584286	2.169643

	reanalysis_specific_humidity_g_per_kg_shift_12 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
5	16.221667
6	16.366327
7	16.567679
8	16.703492
9	16.840286

	reanalysis_precip_amt_kg_per_m2_shift_8	grassy_shift_20	soily_shift_20 \
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	22.485000	NaN	NaN
4	20.428000	NaN	NaN
5	21.438333	NaN	NaN
6	23.890000	NaN	NaN
7	24.653750	NaN	NaN
8	25.342500	NaN	NaN
9	26.650000	0.3	0.7

	watery_shift_20
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

```

5          NaN
6          NaN
7          NaN
8          NaN
9          0.0

```

[10 rows x 42 columns]

```

[80]: # We are losing the first 9 rows
train_final = train_final.dropna().reset_index(drop=True)

```

```

[81]: train_final.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 925 entries, 0 to 924
Data columns (total 42 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   total_cases                               925 non-null    int64
1   year                                       925 non-null    float64
2   weekofyear                               925 non-null    float64
3   week_start_date                           925 non-null
datetime64[ns]
4   month                                     925 non-null    int64
5   fall                                       925 non-null    uint8
6   spring                                    925 non-null    uint8
7   summer                                    925 non-null    uint8
8   winter                                    925 non-null    uint8
9   station_avg_temp_c                        925 non-null    float64
10  station_max_temp_c                        925 non-null    float64
11  station_min_temp_c                        925 non-null    float64
12  reanalysis_tdtr_k                         925 non-null    float64
13  reanalysis_specific_humidity_g_per_kg    925 non-null    float64
14  reanalysis_precip_amt_kg_per_m2          925 non-null    float64
15  ndvi_ne                                   925 non-null    float64
16  ndvi_nw                                   925 non-null    float64
17  ndvi_se                                   925 non-null    float64
18  ndvi_sw                                   925 non-null    float64
19  ndvi_average                             925 non-null    float64
20  ndvi_average_cat                          925 non-null    object
21  grassy                                    925 non-null    uint8
22  soily                                     925 non-null    uint8
23  watery                                    925 non-null    uint8
24  station_avg_temp_c_shift                  925 non-null    float64
25  station_max_temp_c_shift                  925 non-null    float64
26  station_min_temp_c_shift                  925 non-null    float64
27  reanalysis_tdtr_k_shift                   925 non-null    float64

```

```

28 reanalysis_specific_humidity_g_per_kg_shift      925 non-null    float64
29 reanalysis_precip_amt_kg_per_m2_shift           925 non-null    float64
30 grassy_shift                                     925 non-null    float64
31 soily_shift                                       925 non-null    float64
32 watery_shift                                     925 non-null    float64
33 station_max_temp_c_shift_18                     925 non-null    float64
34 station_min_temp_c_shift_18                     925 non-null    float64
35 station_avg_temp_c_shift_18                     925 non-null    float64
36 reanalysis_tdtr_k_shift_8                       925 non-null    float64
37 reanalysis_specific_humidity_g_per_kg_shift_12  925 non-null    float64
38 reanalysis_precip_amt_kg_per_m2_shift_8         925 non-null    float64
39 grassy_shift_20                                  925 non-null    float64
40 soily_shift_20                                    925 non-null    float64
41 watery_shift_20                                  925 non-null    float64
dtypes: datetime64[ns](1), float64(31), int64(2), object(1), uint8(7)
memory usage: 259.4+ KB

```

6.1 Repeat all steps for the final test set:

- First add the last 21 (19 +2 for shifting) rows of the train_full to test_full to not to lose the first part of the dataset while transforming

```

[85]: # Add week_start_date to both datasets
test_features_full['week_start_date'] = test_features['week_start_date']
train_features_full['week_start_date'] = train_features['week_start_date']

```

```

[86]: test_features_long = pd.concat([train_features_full.tail(21),
↳ test_features_full], ignore_index=True)
test_features_long

```

```

[86]:      year  weekofyear  ndvi_ne  ndvi_nw  ndvi_se  ndvi_sw  \
0    2007.0         49.0 -0.03976 -0.042350  0.095600  0.089000
1    2007.0         50.0 -0.13305 -0.045550  0.151440  0.143171
2    2007.0         51.0  0.02945 -0.039000  0.173417  0.150171
3    2007.0         52.0  0.01480  0.016300  0.207267  0.144578
4    2008.0          1.0  0.00060 -0.309600  0.239814  0.195557
..     ...         ...     ...     ...     ...     ...
276  2013.0         13.0 -0.08740 -0.016183  0.156343  0.105186
277  2013.0         14.0 -0.20325 -0.077833  0.204171  0.178914
278  2013.0         15.0 -0.11760 -0.008200  0.192700  0.170429
279  2013.0         16.0  0.08275  0.031200  0.135014  0.074857
280  2013.0         17.0 -0.08730 -0.048667  0.129814  0.117671

      precipitation_amt_mm  reanalysis_air_temp_k  reanalysis_avg_temp_k  \
0                17.85          299.020000          299.021429
1                31.30          298.900000          298.971429
2                62.11          298.668571          298.757143
3                 0.00          298.602857          298.750000

```

4	0.00	298.038571	298.121429
..
276	30.34	298.670000	298.885714
277	6.55	298.035714	298.157143
278	0.00	299.057143	299.328571
279	0.00	298.912857	299.064286
280	45.47	298.067143	298.042857

	reanalysis_dew_point_temp_k	...	reanalysis_relative_humidity_percent	\
0	294.288571	...	75.368571	
1	294.774286	...	78.015714	
2	294.977143	...	80.178571	
3	293.928571	...	75.448571	
4	293.514286	...	76.148571	
..	
276	294.675714	...	78.780000	
277	294.628571	...	81.650000	
278	294.948571	...	78.285714	
279	294.678571	...	77.674286	
280	294.132857	...	79.045714	

	reanalysis_sat_precip_amt_mm	reanalysis_specific_humidity_g_per_kg	\
0	17.85	15.675714	
1	31.30	16.130000	
2	62.11	16.344286	
3	0.00	15.318571	
4	0.00	14.911429	
..	
276	30.34	15.985714	
277	6.55	15.881429	
278	0.00	16.212857	
279	0.00	15.965714	
280	45.47	15.451429	

	reanalysis_tdtr_k	station_avg_temp_c	station_diur_temp_rng_c	\
0	2.100000	25.842857	5.400000	
1	2.485714	25.771429	5.085714	
2	2.371429	25.071429	4.914286	
3	2.985714	25.085714	6.242857	
4	1.842857	25.400000	5.300000	
..	
276	3.314286	27.542857	7.942857	
277	2.828571	26.642857	6.642857	
278	3.171429	27.914286	8.114286	
279	3.042857	27.728571	6.942857	
280	2.342857	26.442857	6.742857	

	station_max_temp_c	station_min_temp_c	station_precip_mm	\
0	29.4	22.8	34.5	
1	28.9	22.2	30.2	
2	28.9	21.7	108.2	
3	28.3	21.1	16.8	
4	29.4	22.2	55.5	
..	
276	33.9	22.8	3.5	
277	33.3	22.8	17.6	
278	32.8	23.3	9.4	
279	31.7	23.9	22.9	
280	31.1	21.7	47.5	

	week_start_date
0	12/3/07
1	12/10/07
2	12/17/07
3	12/24/07
4	1/1/08
..	...
276	3/26/13
277	4/2/13
278	4/9/13
279	4/16/13
280	4/23/13

[281 rows x 23 columns]

[87]: *# create a new month variable:*

```
test_featured = test_features_long.copy()
test_featured["week_start_date"] = pd.
    ↪to_datetime(test_featured["week_start_date"])
test_featured['month'] = test_featured['week_start_date'].dt.month

# create a new season variable:
test_featured['season'] = test_featured['month'].map(month_to_season)
season_features = pd.get_dummies(test_featured['season'])
test_featured = pd.concat([test_featured, season_features], axis = 1)

test_featured = test_featured[['year', 'weekofyear', 'week_start_date',
                                'month', 'fall', 'spring', 'summer', 'winter',
                                'station_avg_temp_c', 'station_max_temp_c',
                                'station_min_temp_c', 'reanalysis_tdtr_k',
                                'reanalysis_specific_humidity_g_per_kg',
                                'reanalysis_precip_amt_kg_per_m2',
                                'ndvi_ne', 'ndvi_nw', 'ndvi_se', 'ndvi_sw']]
```

```

test_featured['ndvi_average'] = (
    (test_featured['ndvi_ne']+test_featured['ndvi_nw']+
     test_featured['ndvi_se']+test_featured['ndvi_sw']))/4

test_featured["ndvi_average_cat"] = test_featured["ndvi_average"].
    apply(get_ndvi_cat)

ndvi_features = pd.get_dummies(test_featured['ndvi_average_cat'])
test_featured = pd.concat([test_featured, ndvi_features], axis = 1)

# shift by 2 weeks
test_shifted = test_featured.copy()

for var in shifted_varbls:
    test_shifted[f"{var}_shift"] = test_shifted[var].shift(2)

test_shifted.dropna(axis=0, inplace=True)
# test_shifted = test_shifted.reset_index(drop=True)

# get rolled means
test_rolled = test_shifted.copy()

varbls_to_see_lags = ['reanalysis_precip_amt_kg_per_m2',
                      'reanalysis_specific_humidity_g_per_kg',
                      'reanalysis_tdtr_k',
                      'station_avg_temp_c',
                      'station_max_temp_c',
                      'station_min_temp_c',
                      'grassy', 'soily', 'watery',
                      'reanalysis_precip_amt_kg_per_m2_shift',
                      'reanalysis_specific_humidity_g_per_kg_shift',
                      'reanalysis_tdtr_k_shift',
                      'station_avg_temp_c_shift',
                      'station_max_temp_c_shift',
                      'station_min_temp_c_shift',
                      'grassy_shift', 'soily_shift', 'watery_shift']

test_rolled = test_rolled[varbls_to_see_lags]

for var in rolled_varbls:
    for num in window:
        test_rolled[f"{var}_{num}"] = test_rolled[var].rolling(num).mean()

```

```

rolled_varbls_to_use = ['station_avg_temp_c_shift',
                        'station_max_temp_c_shift',
                        'station_min_temp_c_shift',
                        'reanalysis_precip_amt_kg_per_m2_shift',
                        'reanalysis_specific_humidity_g_per_kg_shift',
                        'reanalysis_tdtr_k_shift',
                        'grassy_shift', 'soily_shift', 'watery_shift',
                        'station_max_temp_c_shift_18',
                        'station_min_temp_c_shift_18',
                        'station_avg_temp_c_shift_18',
                        'reanalysis_tdtr_k_shift_8',
                        'reanalysis_specific_humidity_g_per_kg_shift_12',
                        'reanalysis_precip_amt_kg_per_m2_shift_8',
                        'grassy_shift_20',
                        'soily_shift_20', 'watery_shift_20']

test_final = test_featured.join(test_rolled[rolled_varbls_to_use])
test_final = test_final.dropna().reset_index(drop=True)
test_final

```

```

[87]:
   year  weekofyear week_start_date  month  fall  spring  summer  winter  \
0  2008.0         18.0    2008-04-29     4     0     1     0     0
1  2008.0         19.0    2008-05-06     5     0     1     0     0
2  2008.0         20.0    2008-05-13     5     0     1     0     0
3  2008.0         21.0    2008-05-20     5     0     1     0     0
4  2008.0         22.0    2008-05-27     5     0     1     0     0
..    ...         ...            ...    ...    ...    ...    ...
255  2013.0         13.0    2013-03-26     3     0     1     0     0
256  2013.0         14.0    2013-04-02     4     0     1     0     0
257  2013.0         15.0    2013-04-09     4     0     1     0     0
258  2013.0         16.0    2013-04-16     4     0     1     0     0
259  2013.0         17.0    2013-04-23     4     0     1     0     0

   station_avg_temp_c  station_max_temp_c  ...  watery_shift  \
0          26.528571          33.3  ...          0.0
1          26.071429          30.0  ...          1.0
2          27.928571          32.8  ...          1.0
3          28.057143          33.3  ...          1.0
4          27.614286          33.3  ...          0.0
..            ...            ...    ...            ...
255         27.542857          33.9  ...          0.0
256         26.642857          33.3  ...          1.0
257         27.914286          32.8  ...          1.0
258         27.728571          31.7  ...          1.0
259         26.442857          31.1  ...          1.0

```

	station_max_temp_c_shift_18	station_min_temp_c_shift_18	\
0	28.816667	21.300000	
1	28.972222	21.388889	
2	29.250000	21.422222	
3	29.283333	21.422222	
4	29.533333	21.516667	
..	
255	30.766667	22.377778	
256	30.550000	22.194444	
257	30.611111	22.105556	
258	30.733333	22.016667	
259	30.827778	21.983333	

	station_avg_temp_c_shift_18	reanalysis_tdtr_k_shift_8	\
0	24.997619	2.678571	
1	25.107143	2.887500	
2	25.187302	2.951786	
3	25.224603	3.012500	
4	25.388889	3.082143	
..	
255	26.509524	2.723214	
256	26.352381	2.775000	
257	26.298413	2.891071	
258	26.227778	2.910714	
259	26.244444	2.971429	

	reanalysis_specific_humidity_g_per_kg_shift_12	\
0	14.065833	
1	14.171071	
2	14.395476	
3	14.514167	
4	14.692143	
..	...	
255	14.662024	
256	14.439048	
257	14.439524	
258	14.490000	
259	14.648452	

	reanalysis_precip_amt_kg_per_m2_shift_8	grassy_shift_20	soily_shift_20	\
0	9.28500	0.0	0.50	
1	8.74875	0.0	0.50	
2	11.10750	0.0	0.50	
3	12.66375	0.0	0.45	
4	13.06625	0.0	0.45	
..	
255	7.32500	0.0	0.60	

256	6.88750	0.0	0.55
257	6.24375	0.0	0.55
258	12.39375	0.0	0.50
259	11.53125	0.0	0.45

	watery_shift_20
0	0.50
1	0.50
2	0.50
3	0.55
4	0.55
..	...
255	0.40
256	0.45
257	0.45
258	0.50
259	0.55

[260 rows x 41 columns]

```
[88]: # Making sure test_final has the same length with older version
len(test_features_full) == len(test_final)
```

[88]: True

```
[89]: # Making sure test_final follows train_final corrrctly in terms of date
print(train_final['week_start_date'])
print('-----')
print(test_final['week_start_date'])
```

0	1990-07-16
1	1990-07-23
2	1990-07-30
3	1990-08-06
4	1990-08-13

	...
920	2008-03-25
921	2008-04-01
922	2008-04-08
923	2008-04-15
924	2008-04-22

Name: week_start_date, Length: 925, dtype: datetime64[ns]

0	2008-04-29
1	2008-05-06
2	2008-05-13
3	2008-05-20

```

4      2008-05-27
...
255    2013-03-26
256    2013-04-02
257    2013-04-09
258    2013-04-16
259    2013-04-23
Name: week_start_date, Length: 260, dtype: datetime64[ns]

```

```
[90]: train_final.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 925 entries, 0 to 924
Data columns (total 42 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   total_cases                             925 non-null    int64
1   year                                    925 non-null    float64
2   weekofyear                             925 non-null    float64
3   week_start_date                         925 non-null    datetime64[ns]
4   month                                  925 non-null    int64
5   fall                                   925 non-null    uint8
6   spring                                 925 non-null    uint8
7   summer                                 925 non-null    uint8
8   winter                                 925 non-null    uint8
9   station_avg_temp_c                     925 non-null    float64
10  station_max_temp_c                     925 non-null    float64
11  station_min_temp_c                     925 non-null    float64
12  reanalysis_tdtr_k                      925 non-null    float64
13  reanalysis_specific_humidity_g_per_kg  925 non-null    float64
14  reanalysis_precip_amt_kg_per_m2        925 non-null    float64
15  ndvi_ne                                925 non-null    float64
16  ndvi_nw                                925 non-null    float64
17  ndvi_se                                925 non-null    float64
18  ndvi_sw                                925 non-null    float64
19  ndvi_average                           925 non-null    float64
20  ndvi_average_cat                       925 non-null    object
21  grassy                                 925 non-null    uint8
22  soily                                  925 non-null    uint8
23  watery                                 925 non-null    uint8
24  station_avg_temp_c_shift               925 non-null    float64
25  station_max_temp_c_shift               925 non-null    float64
26  station_min_temp_c_shift               925 non-null    float64
27  reanalysis_tdtr_k_shift                925 non-null    float64
28  reanalysis_specific_humidity_g_per_kg_shift  925 non-null    float64
29  reanalysis_precip_amt_kg_per_m2_shift  925 non-null    float64
30  grassy_shift                           925 non-null    float64

```

```

31  soily_shift                925 non-null    float64
32  watery_shift              925 non-null    float64
33  station_max_temp_c_shift_18 925 non-null    float64
34  station_min_temp_c_shift_18 925 non-null    float64
35  station_avg_temp_c_shift_18 925 non-null    float64
36  reanalysis_tdtr_k_shift_8   925 non-null    float64
37  reanalysis_specific_humidity_g_per_kg_shift_12 925 non-null    float64
38  reanalysis_precip_amt_kg_per_m2_shift_8         925 non-null    float64
39  grassy_shift_20             925 non-null    float64
40  soily_shift_20              925 non-null    float64
41  watery_shift_20             925 non-null    float64
dtypes: datetime64[ns](1), float64(31), int64(2), object(1), uint8(7)
memory usage: 259.4+ KB

```

```
[91]: test_final.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 260 entries, 0 to 259
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   year                                260 non-null    float64
1   weekofyear                          260 non-null    float64
2   week_start_date                    260 non-null    datetime64[ns]
3   month                              260 non-null    int64
4   fall                               260 non-null    uint8
5   spring                             260 non-null    uint8
6   summer                             260 non-null    uint8
7   winter                             260 non-null    uint8
8   station_avg_temp_c                 260 non-null    float64
9   station_max_temp_c                 260 non-null    float64
10  station_min_temp_c                 260 non-null    float64
11  reanalysis_tdtr_k                   260 non-null    float64
12  reanalysis_specific_humidity_g_per_kg 260 non-null    float64
13  reanalysis_precip_amt_kg_per_m2       260 non-null    float64
14  ndvi_ne                             260 non-null    float64
15  ndvi_nw                             260 non-null    float64
16  ndvi_se                             260 non-null    float64
17  ndvi_sw                             260 non-null    float64
18  ndvi_average                        260 non-null    float64
19  ndvi_average_cat                    260 non-null    object
20  grassy                              260 non-null    uint8
21  soily                               260 non-null    uint8
22  watery                              260 non-null    uint8
23  station_avg_temp_c_shift            260 non-null    float64
24  station_max_temp_c_shift            260 non-null    float64
25  station_min_temp_c_shift            260 non-null    float64

```

```

26 reanalysis_precip_amt_kg_per_m2_shift      260 non-null    float64
27 reanalysis_specific_humidity_g_per_kg_shift 260 non-null    float64
28 reanalysis_tdtr_k_shift                    260 non-null    float64
29 grassy_shift                               260 non-null    float64
30 soily_shift                                260 non-null    float64
31 watery_shift                               260 non-null    float64
32 station_max_temp_c_shift_18                260 non-null    float64
33 station_min_temp_c_shift_18                260 non-null    float64
34 station_avg_temp_c_shift_18                260 non-null    float64
35 reanalysis_tdtr_k_shift_8                  260 non-null    float64
36 reanalysis_specific_humidity_g_per_kg_shift_12 260 non-null    float64
37 reanalysis_precip_amt_kg_per_m2_shift_8     260 non-null    float64
38 grassy_shift_20                            260 non-null    float64
39 soily_shift_20                              260 non-null    float64
40 watery_shift_20                            260 non-null    float64
dtypes: datetime64[ns](1), float64(31), int64(1), object(1), uint8(7)
memory usage: 71.0+ KB

```

```

[92]: # Export the final datasets as csv to be used for modeling
train_final.to_csv("train_final.csv")
test_final.to_csv("test_final.csv")

```

6.1.1 Export as PDF:

```

[ ]: # Packages required for using nbconvert PDF
# ! apt-get install texlive texlive-xetex texlive-latex-extra pandoc
# ! pip install py pandoc
# ! pip install nbconvert

```

```

[ ]: # First you need to download a copy of the ipynb notebook and upload it back to
↳ the drive, it is placed under /content/
# ! jupyter nbconvert --to pdf /content/notebook_preprocessing.ipynb

```